

# Kubernetes Overview

# 4

## 1. The Project



`github.com/openshift/origin`

`github.com/kubernetes`

## 2. Community Deployment

`kops`



`kubeadm`

`minikube/minishift/crc` (codeready containers)

## 3. Product

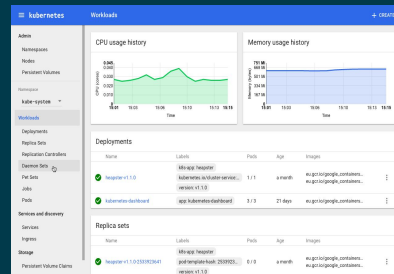
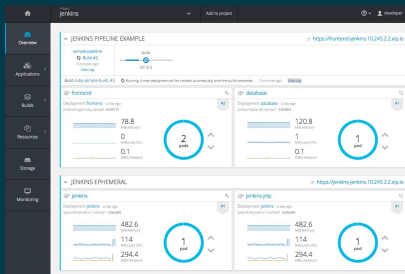


## 4. Service



# 4

## 1. Dashboard



## 2. Command Line Interface

```
$ oc/kubectl apply -f my-new-deployment.yaml
```

## 3. SDK/Client Libraries

```
pod, err :=  
c.Pods(v1.NamespaceDefault).Get("my-pod")  
  
if err != nil {  
    fmt.Println(err)  
    return  
}
```

## 4. Helm

```
$ helm install stable/mariadb
```

# API

**oc/kubectl**

A command line interface for running commands  
against Kubernetes clusters.

# Installing kubectl

Homebrew on macOS

```
$ brew install kubectl
```

Chocolatey on Windows

```
$ choco install kubectl
```

Linux/macOS/Windows

```
$ curl -LO https://storage.googleapis.com/kubernetes-release/release
```

# How to Provide Your K8s Auth Info

3

## 1 - argument

```
oc --kubeconfig config cluster-info
```

## 2 - variable

```
export KUBECONFIG=config  
oc cluster-info
```

## 3 - .kube/config

```
oc cluster-info
```

```
cat .kube/config
```



# Anatomy of a kubeconfig

## Clusters

```
apiVersion: v1
clusters:
- cluster:
  certificate-authority-data: <cacert>
  server: https://192.168.56.60:6443
  name: my-first-cluster
```

## Context

```
contexts:
- context:
  cluster: my-first-cluster
  user: kubelet
  name: default
```

## Current

```
current-context: default
kind: Config
preferences: {}
```

## Users

```
users:
- name: kubelet
  user:
    client-certificate-data: <yourclientcert>
    client-key-data: <yourclientkey>
```

# About Client-Go

A collection of tools/frameworks (in the form of Go packages) for all your Kubernetes programming needs.

# Contents of Client-Go

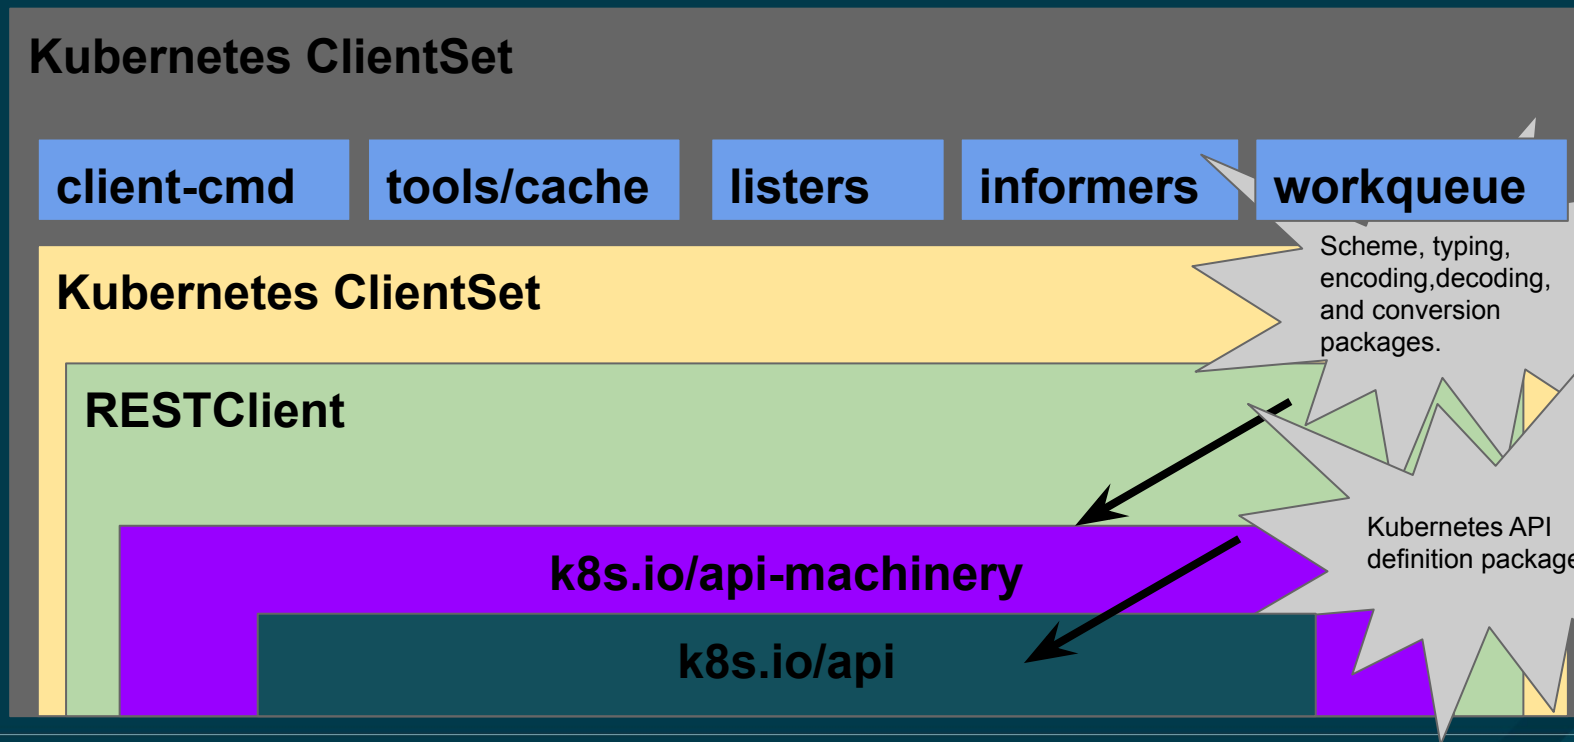
## Clients

- Clientset
- Dynamic Client
- RESTclient


## Utilities for Writing Controllers

- Workqueue
- Informers/Shared Informers

# Client-Go Implementation



# Out-of-Cluster Interaction With Kubernetes API



1

client-cmd

# Fetch the kube-config file and use current-context

```
func main() {  
    var kubeconfig *string  
    if home := homeDir(); home != "" {  
        kubeconfig = flag.String("kubeconfig", filepath.Join(home, ".kube", "config"), "(optional) absolute  
path to the kubeconfig file")  
    } else {  
        kubeconfig = flag.String("kubeconfig", "", "absolute path to the kubeconfig file")  
    }  
    flag.Parse()  
  
    config, err := clientcmd.BuildConfigFromFlags("", *kubeconfig)  
    if err != nil {  
        panic(err.Error())  
    }  
}
```

["k8s.io/client-go/tools/clientcmd"](https://k8s.io/client-go/tools/clientcmd)

# Create the client-set

```
clientset, err := kubernetes.NewForConfig(config)
if err != nil {
    panic(err.Error())
}
```

**"k8s.io/client-go/kubernetes"**

Retrieve the Corev1 Client via clientset and list all pods in the cluster (across all namespaces)

```
for {  
    pods, err := clientset.CoreV1().Pods("").List(metav1.ListOptions{})  
    if err != nil {  
        panic(err.Error())  
    }  
    fmt.Printf("There are %d pods in the cluster\n", len(pods.Items))  
}
```

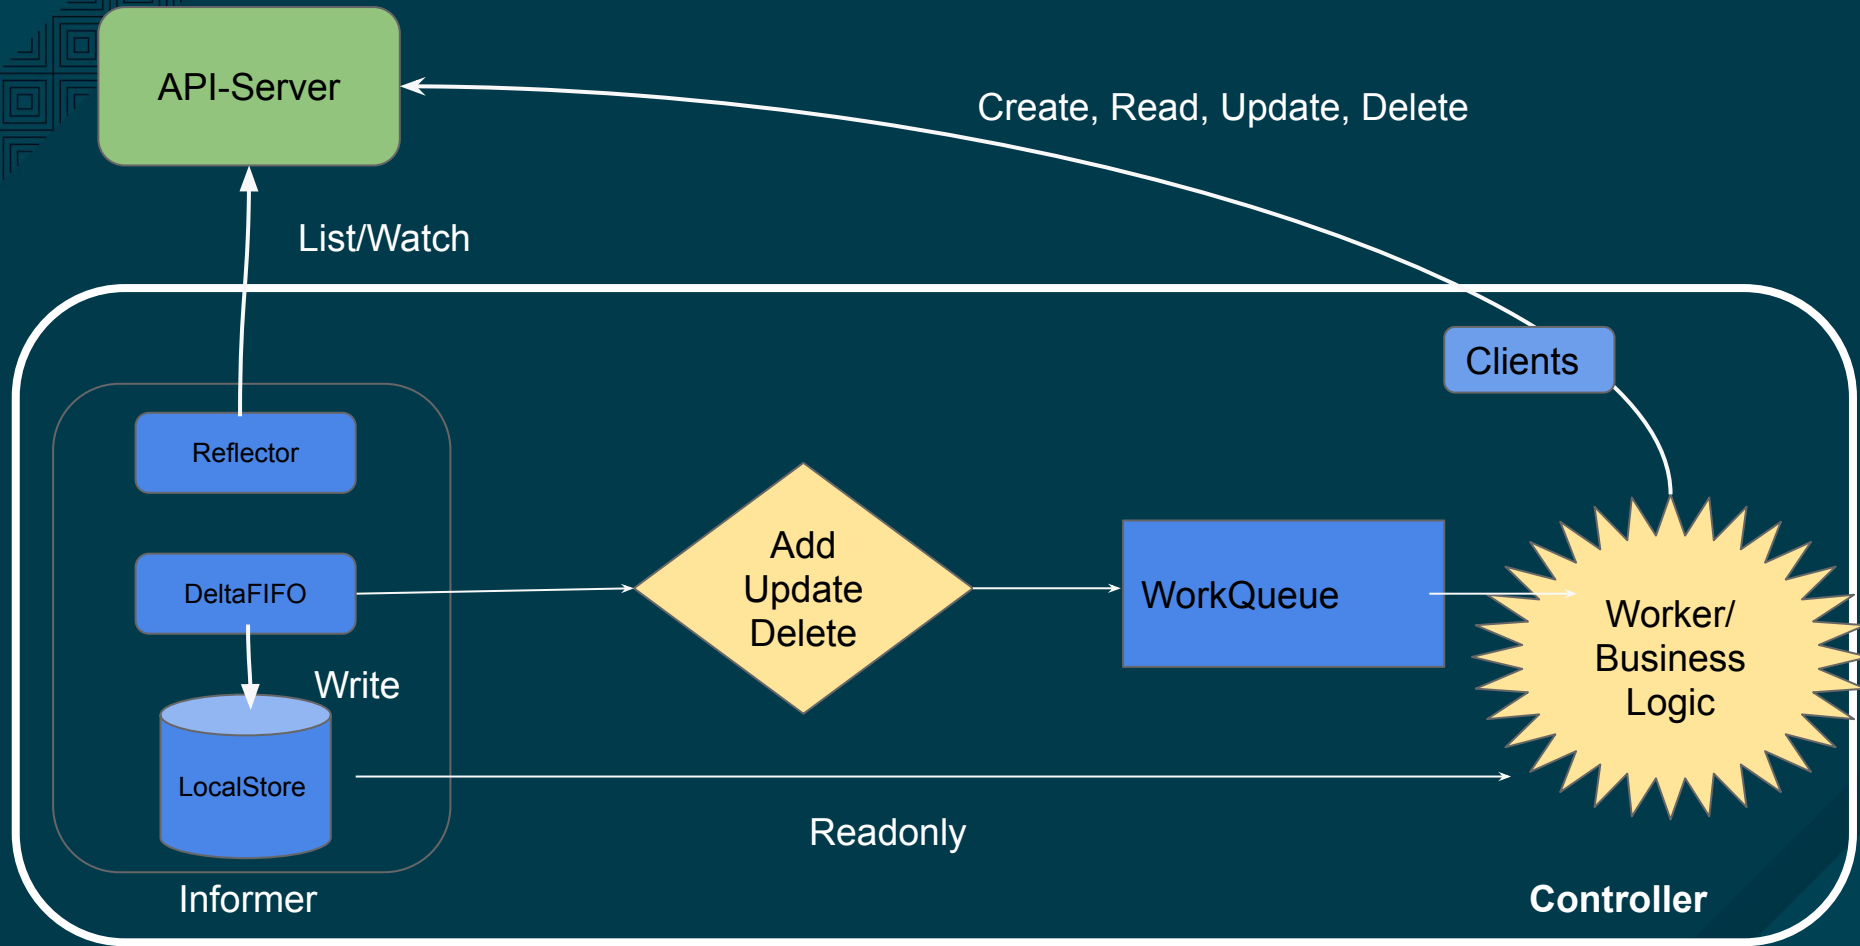
```
Verb(verb string) *Request  
Post() *Request  
Put() *Request  
Patch(pt types.PatchType) *Request  
Get() *Request  
Delete() *Request  
APIVersion() schema.GroupVersion
```

```
type ListOptions struct {  
    TypeMeta      `json:",inline"`  
    LabelSelector string  
    `json:"labelSelector,omitempty"`  
    FieldSelector string  
    `json:"fieldSelector,omitempty"`  
    protobuf:"bytes,2,opt,name=fieldSele
```



Retrieve the Corev1 Client via clientset and get **spec** for **individual pod** in the **default** namespace.

```
for {  
    pod, err := clientset.CoreV1().Pods("default").Get("my-pod", metav1.GetOptions{})  
    if err != nil {  
        panic(err.Error())  
    }  
    fmt.Printf("%v\n\n\n", pod.spec)
```



```
1  while true {  
2    receiveInfoAboutAPIObjects()  
3    synchronizeRealStateToMatchFetchedInfo()  
4  }
```

# Kubernetes Concepts

# What is a Kubernetes Resource?

# Most Common Definition...

Any individual Kubernetes item such as a deployment, pod, service, or secret, etc.

# Kubernetes Resources

- Nodes
- Namespaces
- Pods
- Endpoints
- Services
- Deployments
- ReplicaSets
- Persistent Volumes
- PersistentVolumeClaims
- ConfigMaps
- DaemonSets
- StatefulSets
- Events
- PodDisruptionBudgets
- PodSecurityPolicies
- ResourceQuotas
- Service Accounts
- HorizontalPodAutoScalers



# A Better Definition...

A Kubernetes Resource is a  
**declarative API** with well defined  
Schema structure and endpoints.\*

\*Because the structure of the Schema and Endpoints are predictable and structured, most Kubernetes tools work with any Kubernetes API even if they are not part of the core (e.g. extensions through CRDs).

```
oc proxy
```

```
curl localhost:8001
```

# What is a Declarative API?

# Declarative vs. Imperative API

- Declarative expresses a fixed state that the cluster must continually work towards.
- “What to Do”
  - Example: `$ replicas 3`
- Imperative API expresses an operation that may change state but does not define an absolute state that must be maintained.
- “How to Do It”
  - Example: `$ add-pods 2`

# ReplicaSet Manifest

# Resource Schema Components

GVK aka TypeMeta

Metadata aka ObjectMeta

Spec

Status

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
```

```
metadata:
  name: my-first-replica-set
  namespace: myproject
```

```
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 5
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
```

```
status:
  availableReplicas: 1
  fullyLabeledReplicas: 1
  observedGeneration: 1
  readyReplicas: 1
  replicas: 1
```

# Resource Schema: Group, Version, Kind (GVK)

```
apiVersion: extensions/v1beta1  
kind: ReplicaSet
```

- The resource **Group** is similar to package in a language. It disambiguates different APIs that may happen to have identically named Kinds. Groups often contain a domain name, such as redhat.com.
- The resource **Version** defines the stability of the API and backward compatibility guarantees - such as v1beta1 or v1.
- The resource **Kind** is the name of the API - such as Deployment or Service.



# A Note About API Versions

```
└─ apis
  └─ workloads
    └─ v1
      ├── containerset_types_test.go
      ├── containerset_types.go
      ├── doc.go
      ├── register.go
      └── v1_suite_test.go
    └─ v1beta1
      ├── containerset_types_test.go
      ├── containerset_types.go
      ├── doc.go
      ├── register.go
      ├── v1beta1_suite_test.go
      └── zz_generated.deepcopy.go
    └─ v1beta2
      ├── containerset_types_test.go
      ├── containerset_types.go
      ├── doc.go
      ├── register.go
      └── v1beta2_suite_test.go
  ├── group.go
  ├── addtoscheme_workloads_v1.go
  ├── addtoscheme_workloads_v1beta1.go
  ├── addtoscheme_workloads_v1beta2.go
  └── apis.go
```

# Difference between API Version Numbers

i.e. apps/v1beta1, apps/v1beta2

- Unspecified fields may have different defaults.
- The same logical fields may have different names.

```
kubectl explain deployments.spec --api-version="apps/v1beta1"
```

```
revisionHistoryLimit  <integer>
```

The number of old ReplicaSets to retain to allow rollback. This is a pointer to distinguish between explicit zero and not specified.  
**Defaults to 2.**

```
kubectl explain deployments.spec --api-version="apps/v1beta2"
```

```
revisionHistoryLimit  <integer>
```

The number of old ReplicaSets to retain to allow rollback. This is a pointer to distinguish between explicit zero and not specified.  
**Defaults to 10.**

# API Versions

## Alpha (i.e. v1alpha1)

- Disabled by default. Must be enabled via API.
- May contain bugs. Features may be changed or removed. Field names may also be changed and not supported in the future.
- Only use for short-lived testing clusters.

## Beta (i.e. v1beta1)

- Enabled by default.
- Considered safe. Backwards compatibility on field names.
- Support for the feature will not be dropped, though details may change.

## Stable (i.e. v1,v2)

- Stable versions of features will appear in many subsequent versions.

# Not Flexible

`http://kubernetes:6443/api/v1/pods`

`http://kubernetes:6443/api/v1/replicaset`

`http://kubernetes:6443/api/v1/services`

`http://kubernetes:6443/api/v1/deployments`

# Flexible

```
curl -k https://kubernetes:6443
```

```
"/api/v1"  
"/apis/authentication.k8s.io/v1"  
"/apis/authentication.k8s.io/v1beta1"  
"/apis/authorization.k8s.io/v1"  
"/apis/authorization.k8s.io/v1beta1"  
"/apis/certificates.k8s.io/v1beta1"  
"/apis/certificates.k8s.io"  
"/apis/extensions/v1beta1"  
"/apis/policy/v1beta1"  
"/apis/rbac.authorization.k8s.io/v1beta1"  
"/apis/rbac.authorization.k8s.io/v1alpha1"  
"/apis/storage.k8s.io/v1"  
"/apis/storage.k8s.io/v1beta1"
```

Allows the program to move, change, and grow over time.

Engineers can advertise to support older API versions, and offer backward-compatibility guarantees.



# See Current API-Versions

```
oc api-versions
```

```
"/api/v1"  
"/apis/authentication.k8s.io/v1"  
"/apis/authentication.k8s.io/v1beta1"  
"/apis/authorization.k8s.io/v1"  
"/apis/authorization.k8s.io/v1beta1"  
"/apis/certificates.k8s.io/v1beta1"  
"/apis/certificates.k8s.io"  
"/apis/extensions/v1beta1"  
"/apis/policy/v1beta1"  
"/apis/rbac.authorization.k8s.io/v1beta1"  
"/apis/rbac.authorization.k8s.io/v1alpha1"  
"/apis/storage.k8s.io/v1"  
"/apis/storage.k8s.io/v1beta1"
```

# News Snippet About Introduction of v1 NetworkPolicy

## Two of the changes you need to be aware of are:

### » The v1beta1 NetworkPolicy API Has Been Deprecated

The v1beta1 version of the NetworkPolicy API has been deprecated in favor of moving forward with the new behaviors and updating the behavior of the *extensions* API to allow for future expansion and development. Keep in mind that while the v1 NetworkPolicy API eclipses the existing beta, the new API endpoint will only be available on Kubernetes 1.7+ (as older versions do not include the v1 API code). As such, as you work towards upgrading, you'll want to ensure that you are using the correct version of Project Calico for the NetworkPolicy behavior you want.

### » The DefaultDeny Annotation Has Been Removed

One of the bigger changes in Kubernetes 1.7 is the removal of the DefaultDeny annotation. This means that when upgrading, you should **first delete any existing NetworkPolicy** objects in namespaces that previously **did not have** the "DefaultDeny" annotation (as this may cause Kubernetes to unintentionally block traffic now).



# Kubernetes API Actions and HTTP Method

<u>Verb (API)</u>	<u>HTTP Method</u>
Get	GET
List	GET
Watch	GET
Create	POST
Update	PUT
Patch	PATCH
Delete	DELETE

## Kubernetes Subcommand & HTTP Method

<u>Subcommand</u>	<u>Object does not exist</u>	<u>Object exists</u>
apply	POST	PATCH/DELETE
create	POST	error!
replace	error!	PUT
delete	error!	DELETE
patch	error!	PATCH

# Interacting with the API

Create!

```
kubectl/oc create -f podmanifest.json
```

```
curl -X POST http://localhost:8001/api/v1/namespaces/myproject/pods/ -H  
"Content-type: application/json" -d @podmanifest.json
```

Update!

```
kubectl/oc replace -f podmanifest.json
```

```
curl -X PUT http://localhost:8001/api/v1/namespaces/myproject/pods/mypod -H  
"Content-type: application/json" -d @newpodmanifest.json
```

Patch!

```
kubectl/oc patch -f patch.json
```


```
kubectl patch etcdcluster example-etcd-cluster --type='json' -p '[{"op":  
"replace", "path": "/spec/size", "value":5}]'
```

# Labels/Selectors


Key/value pairs attached to resources.

Used for grouping, viewing, and  
operating.

# Labels: Grouping



```
labels:  
  name: apache  
  app: mynewapp  
  role: frontend
```



```
labels:  
  name: mysql  
  app: mynewapp  
  role: db
```

# Labels: Viewing

```
kubectl get pods --show-labels
```

db-dev	1/1	Running	0	6s	app=my-app,environment=dev,tier=backend
www-dev	1/1	Running	0	6s	app=my-app,environment=dev,tier=frontend
www-prod	1/1	Running	0	6s	app=my-app,environment=production,tier=frontend

```
kubectl get pods -L app,environment,tier -l environment!=dev
```

www-prod	1/1	Running	0	4m	my-app	production	frontend
----------	-----	---------	---	----	--------	------------	----------

```
kubectl get pods -l "tier notin (backend,cache),environment in (dev)"
```

www-dev	1/1	Running	0	6m
---------	-----	---------	---	----

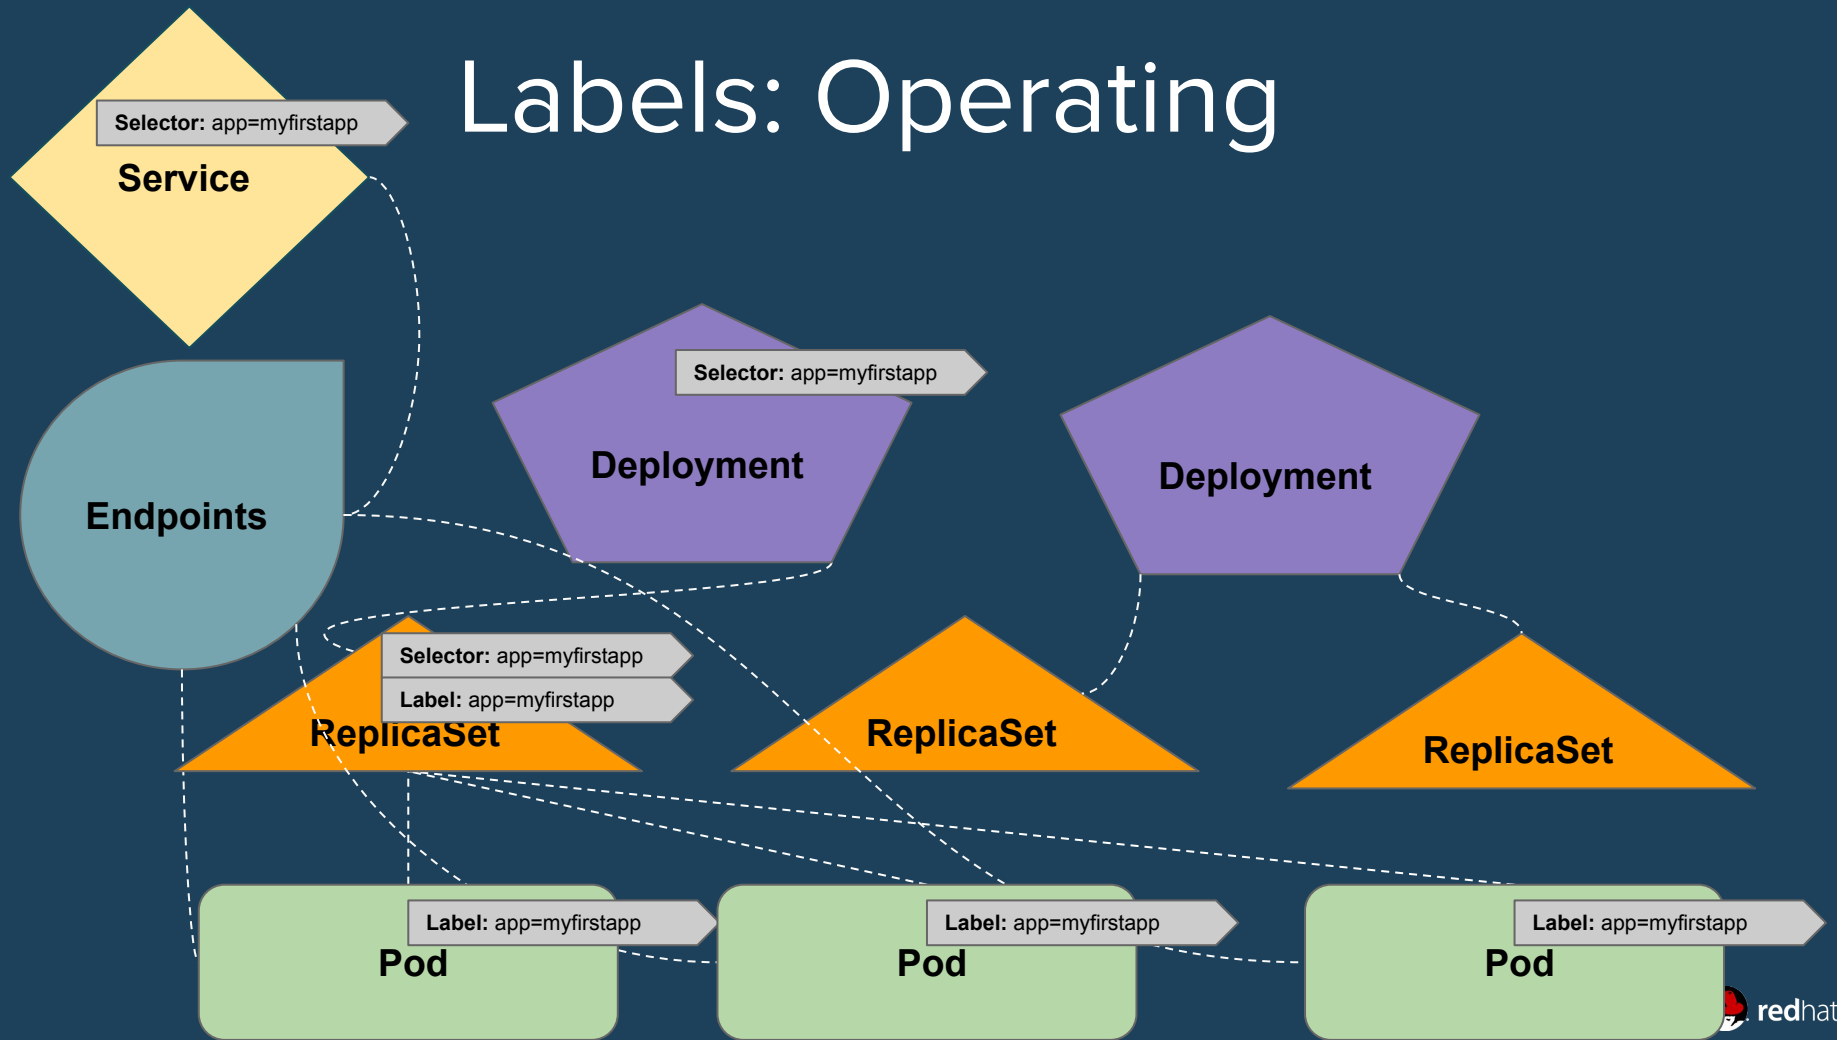
# We will Return at

12:00pm CT

11:00am UTC



# Labels: Operating



# Using Labels

Labels can be displayed as a column in output with `-L` option:

```
$ kubectl get pods -L tier
```

NAME	READY	STATUS	RESTARTS	AGE	TIER
my-nginx-3800858182-1v53o	1/1	Running	0	46s	backend
my-nginx-3800858182-2ds1q	1/1	Running	0	46s	backend

# Updating Labels

Sometimes existing pods and other resources need to be relabeled before creating new resources

```
$ kubectl label pods -l app=nginx,tier=fe # select by label app=nginx, apply tier=fe
pod "my-nginx-v4-9gw19" labeled
pod "my-nginx-v4-hayza" labeled
pod "my-nginx-v4-mde6m" labeled
pod "my-nginx-v4-sh6m8" labeled
pod "my-nginx-v4-wfof4" labeled
```

```
$ kubectl get pods -l app=nginx -L tier
```

NAME	READY	STATUS	RESTARTS	AGE	TIER
my-nginx-v4-9gw19	1/1	Running	0	15m	fe
my-nginx-v4-hayza	1/1	Running	0	14m	fe
my-nginx-v4-mde6m	1/1	Running	0	18m	fe
my-nginx-v4-sh6m8	1/1	Running	0	19m	fe
my-nginx-v4-wfof4	1/1	Running	0	16m	fe

# Using Labels Effectively

Examples of multiple labels for app, tier and role:

```
labels:  
  app: guestbook  
  tier: frontend
```

```
labels:  
  app: guestbook  
  tier: backend  
  role: master
```

```
labels:  
  app: guestbook  
  tier: backend  
  role: slave
```

Other example labels:

- "release" : "stable" or "canary"
- "partition" : "customerA" or "customerB"
- "track" : "daily" or "weekly"

# What is a Declarative API?

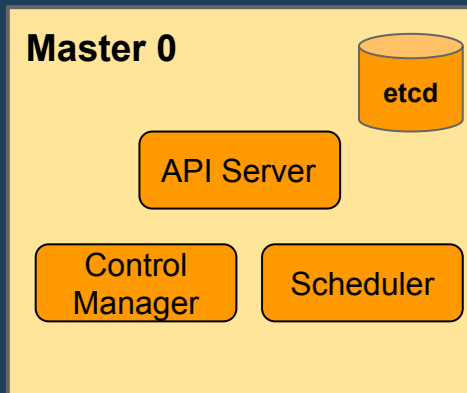
# Cluster

A group of servers (or virtual machines)  
configured to run a functioning  
Kubernetes system.

# Kubernetes Cluster



# Kubernetes Master

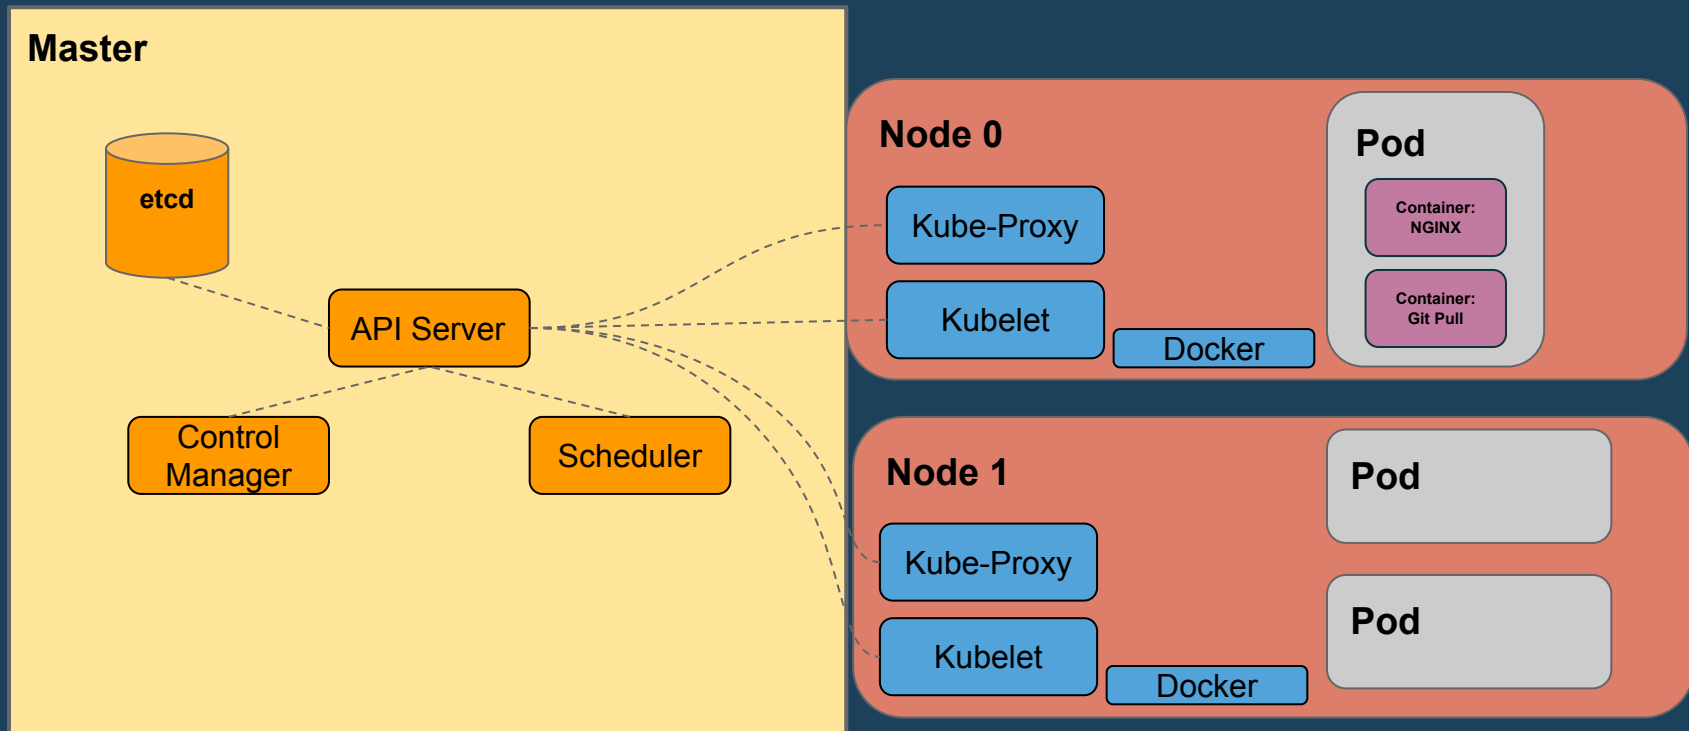


Also known as a  
*controller* or  
*head* node.





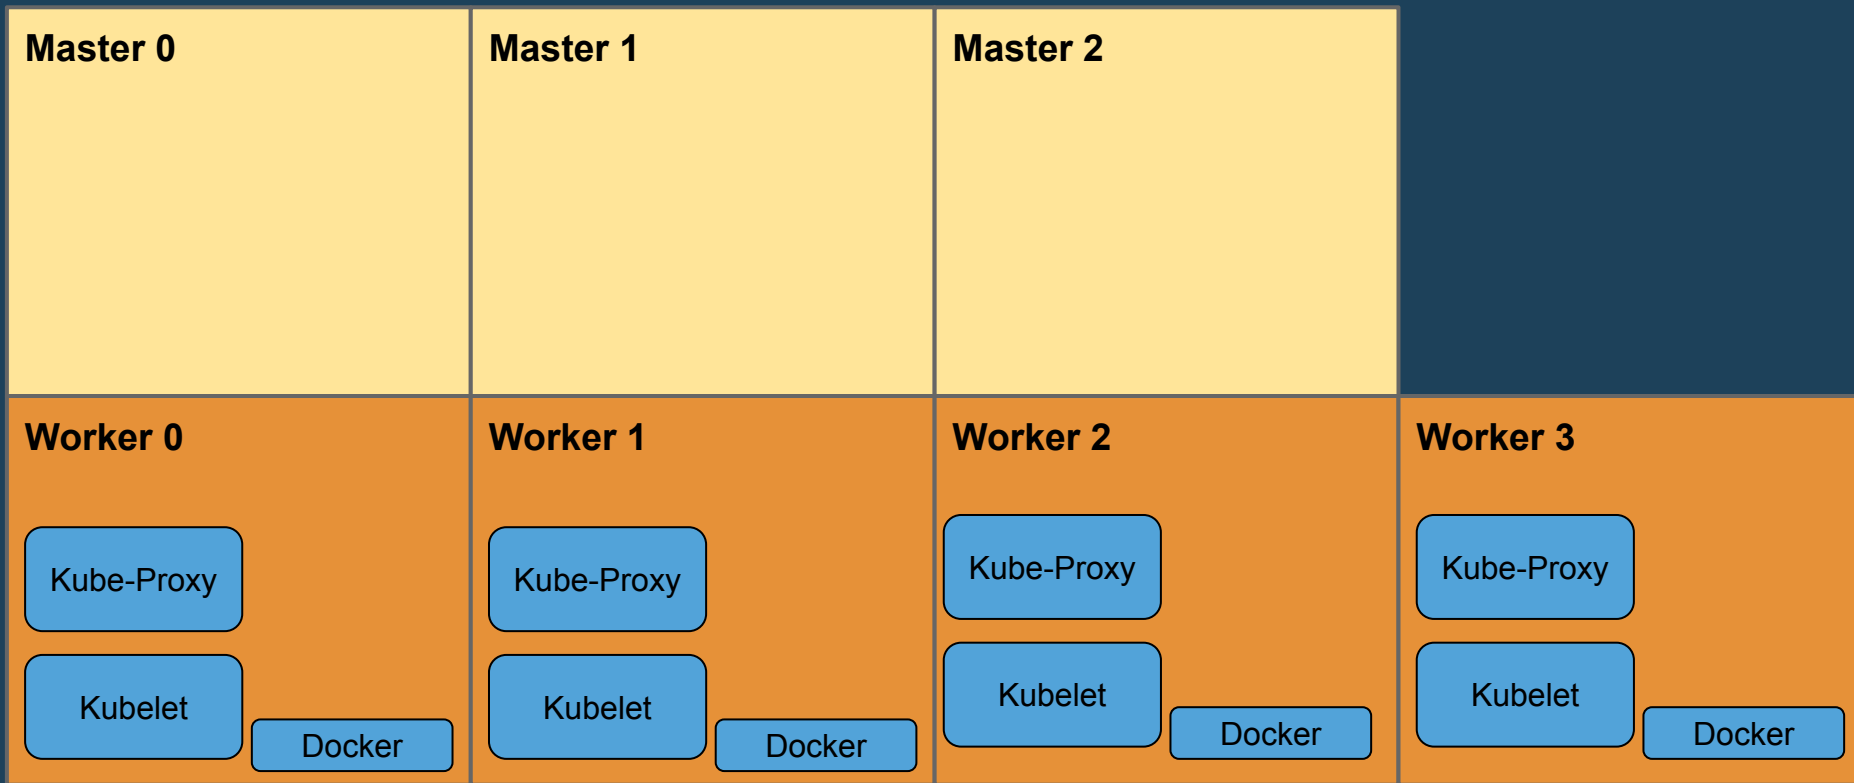
# Kubernetes Cluster



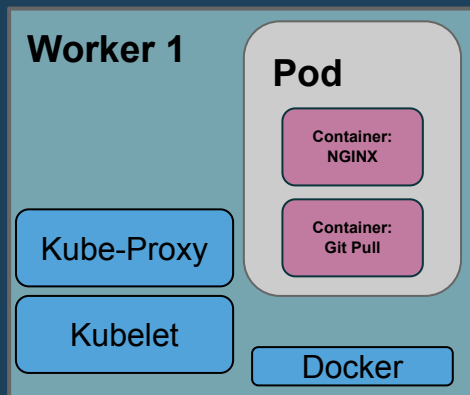
# Node

A worker machine (previously known as a *minion*). Contains the services necessary to run pods.

# Kubernetes Node/Worker



# Kubernetes Node/Worker



Previously  
known as  
minion.

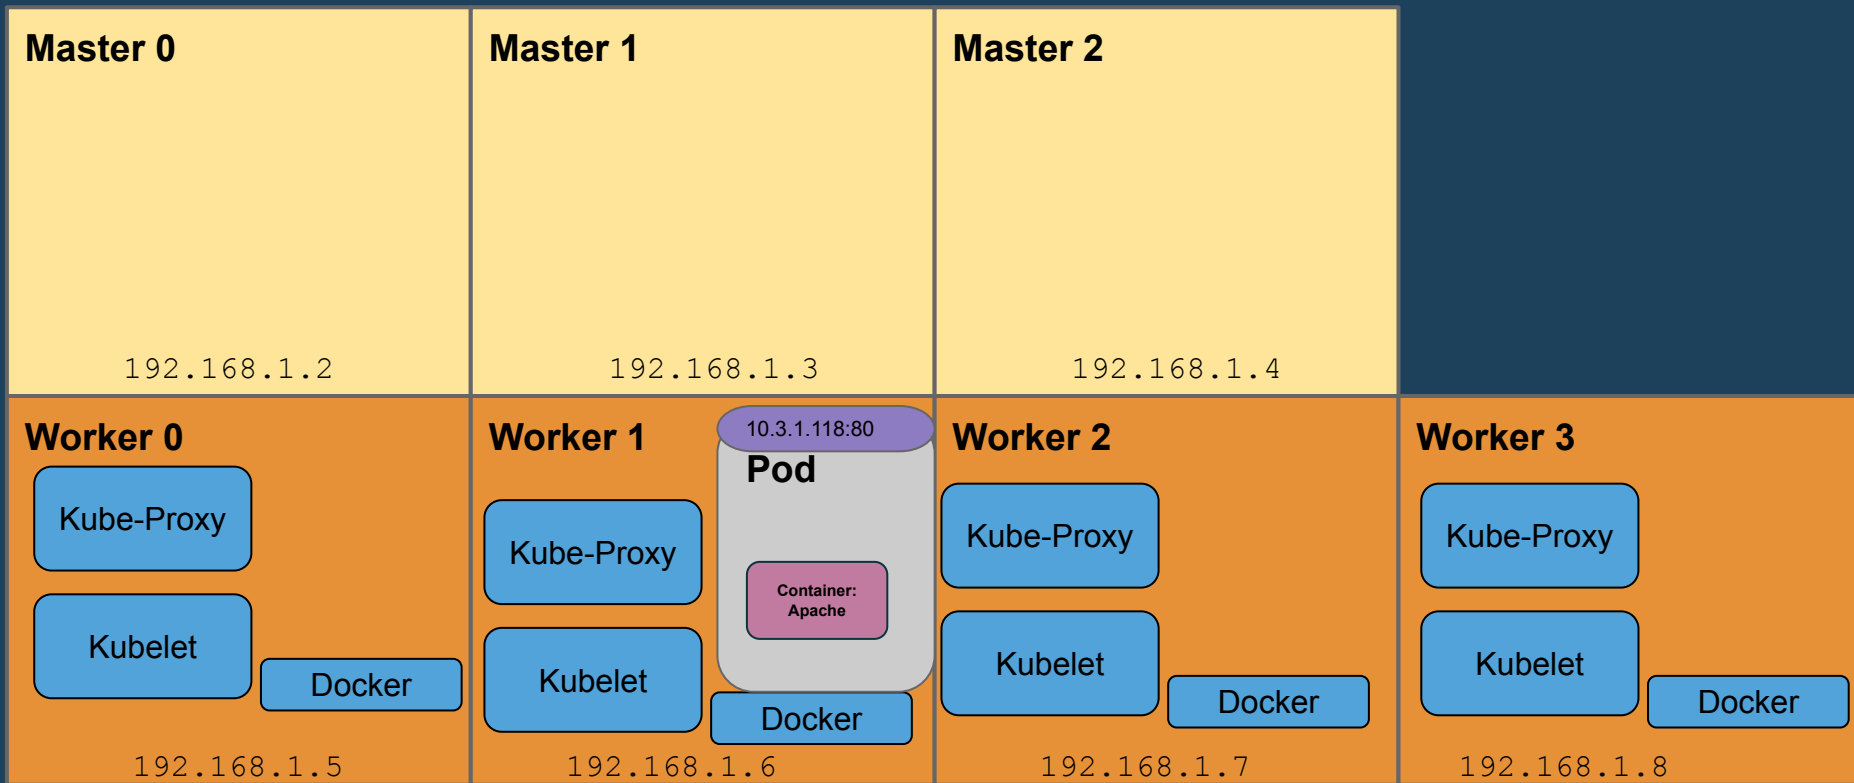
# Pod

A group of one or more containers  
running on a single node.

Pods are the smallest deployable units of computing that can be created and managed in Kubernetes.

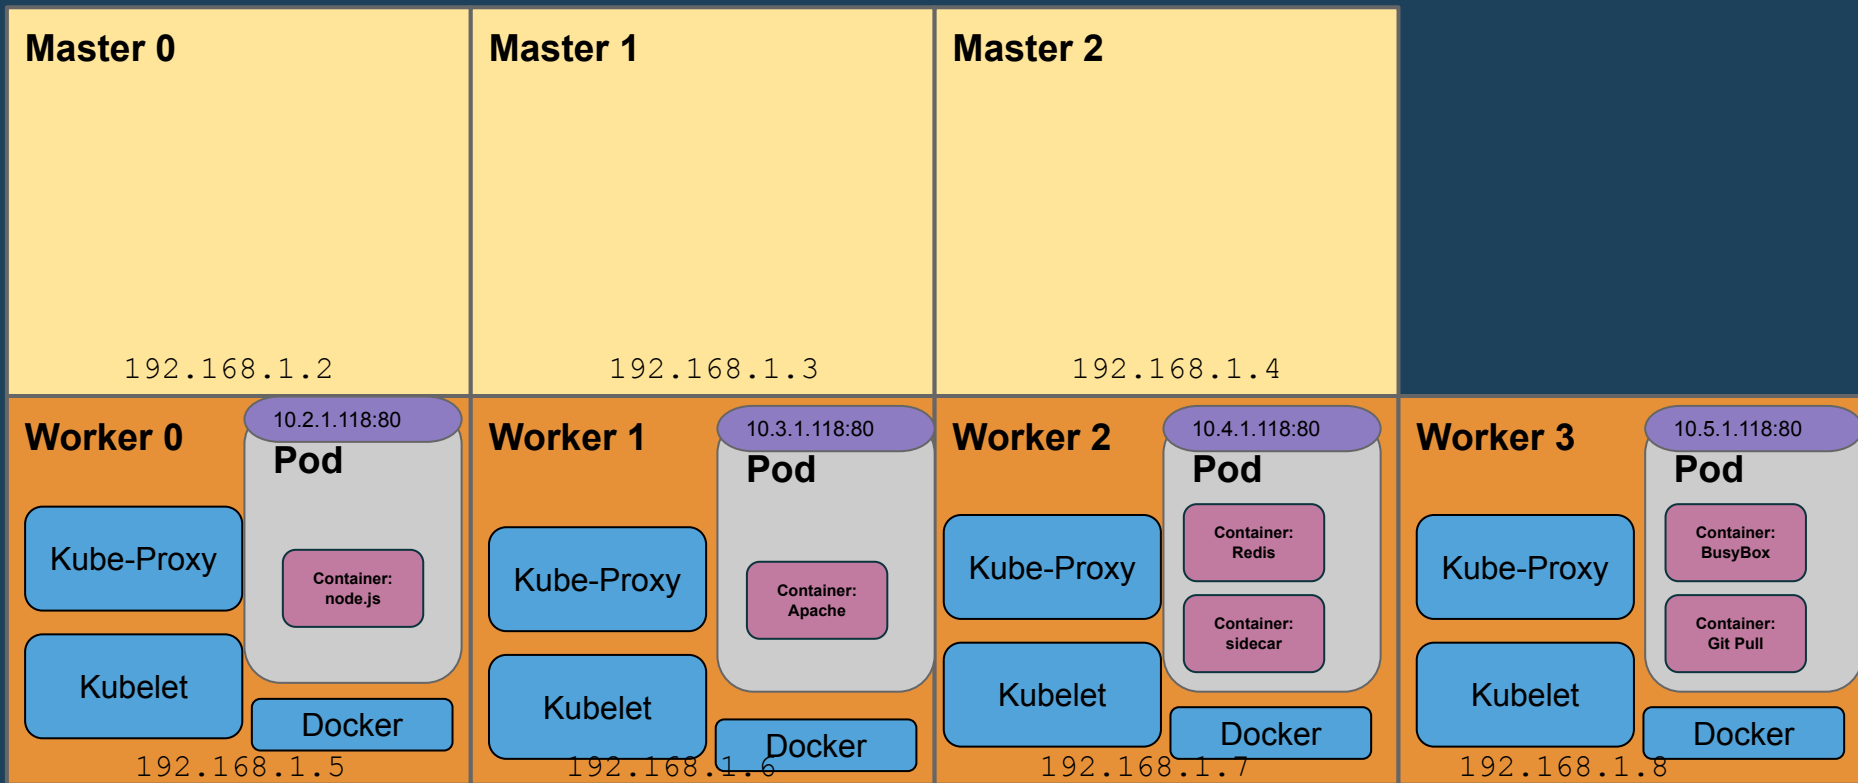
```
kubectl run my-first-pod --image=apache --restart=Never --port=80
```

# Kubernetes Pod



```
kubectl run my-first-pod --image=apache --restart=Always --port=80
```

# Kubernetes Pods






Let's take a closer look at a pod.

# About Pods

Contain one or more containers.

These “extra” containers are commonly called “side-cars”.



# What can a sidecar container do?

Data puller.

Data watcher.

Log watcher.

Reverse Proxy (Log/Modify Requests)

# What is separated/isolated?

# Mnt Namespaces

Separate filesystems.

## Pod

apache

```
$ ls /
```

bin	etc	lib	mnt	root	sbin	sys	usr
dev	home	media	proc	run	srv	tmp	var

git-sync

```
$ ls /
```

bin	etc	home	lib	mnt	root	sbin	sys
dev	home	media	proc	run	srv	tmp	var

# PID Namespaces

Separate processes.

Pod

apache

\$ ps

PID	USER	COMMAND
1	root	apache2

git-sync

\$ ps

PID	USER	COMMAND
1	root	pull.py

# User Namespaces

Separate user namespaces.

Pod

apache

\$ whoami

root

git-sync


\$ whoami

root



# Cgroups


Each container has its own  
cpu/memory cgroup.



# What is shared?

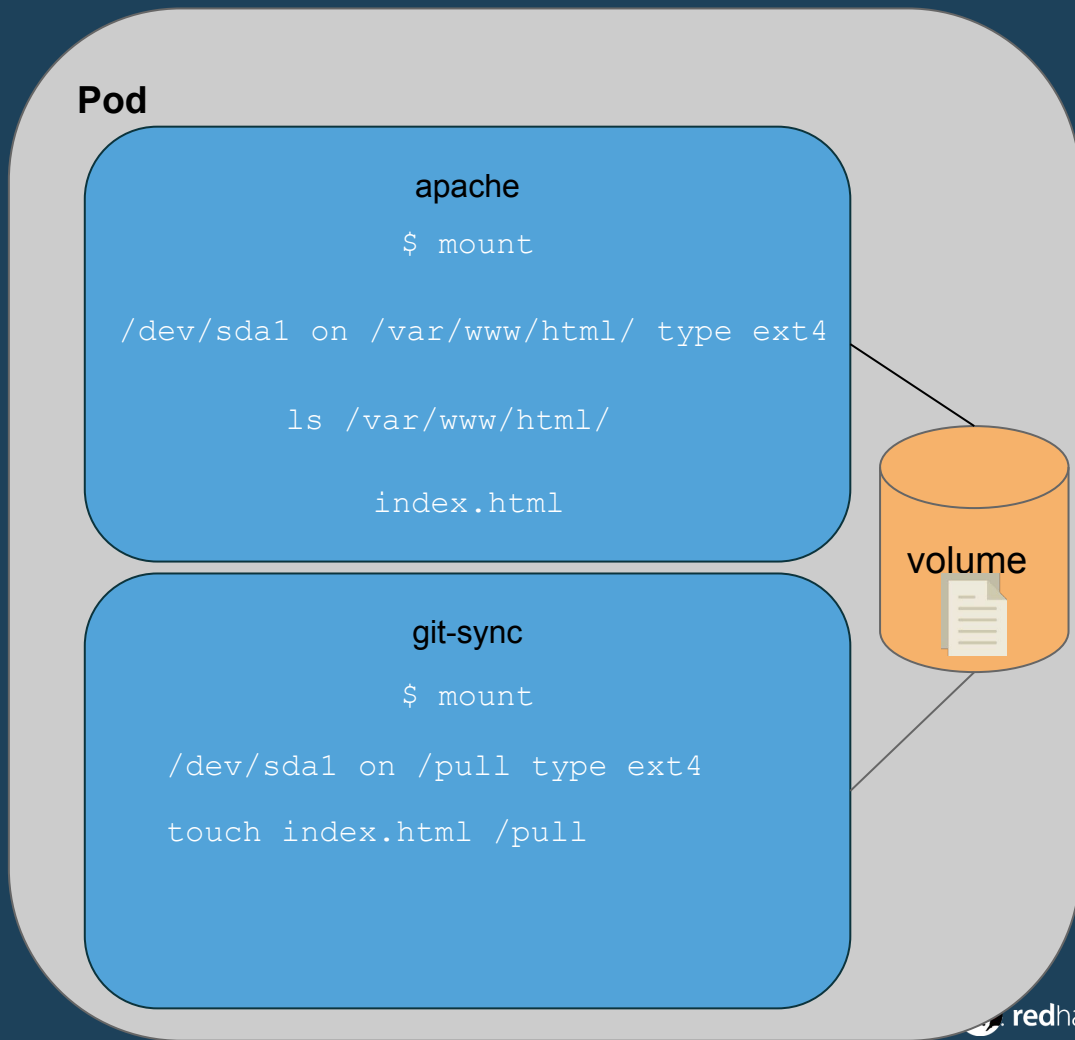
# UTS Namespaces

Same hostname.



# Volumes

Share the same volume but mounted as different directories.



# Net Namespaces

Same IP address.

## Pod

apache

\$ ip a

```
31: eth0@if32: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu
1500 qdisc noqueue state UP
```

```
link/ether 02:42:ac:11:00:04 brd ff:ff:ff:ff:ff:ff
```

```
inet 10.2.0.4/16 scope global eth0
```

```
valid_lft forever preferred_lft forever
```

```
inet6 fe80::42:acff:fe11:4/64 scope link
```

```
valid_lft forever preferred_lft forever
```

git-sync

\$ ip a

```
31: eth0@if32: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu
1500 qdisc noqueue state UP
```

```
link/ether 02:42:ac:11:00:04 brd ff:ff:ff:ff:ff:ff
```

```
inet 10.2.0.4/16 scope global eth0
```

```
valid_lft forever preferred_lft forever
```

```
inet6 fe80::42:acff:fe11:4/64 scope link
```

```
valid_lft forever preferred_lft forever
```

# Net Namespaces

Same network sockets.

## Pod

apache

```
$ netstat -ntlp
```

Active Internet connections (only servers)

Proto	Recv-Q	Send-Q	Local Address	State	Program
tcp	0	0	0.0.0.0:80	LISTEN	apache2

git-sync


```
$ netstat -ntlp
```

Active Internet connections (only servers)

Proto	Recv-Q	Send-Q	Local Address	State	Program
tcp	0	0	0.0.0.0:80	LISTEN	apache2

# IPC Namespace

Containers can communicate via SystemV or POSIX shared memory, semaphores, or messages.



# kubectl run

- ‘Off the cuff’ command to create **pods, deployments, or jobs**.
- Great for testing and troubleshooting.
- Similar to “docker run” in usage.




# Labels/Selectors


Key/value pairs attached to resources.

Used for grouping, viewing, and  
operating.

# Labels: Grouping



```
labels:  
  name: apache  
  app: mynewapp  
  role: frontend
```



```
labels:  
  name: mysql  
  app: mynewapp  
  role: db
```

# Labels: Viewing

```
kubectl get pods --show-labels
```

db-dev	1/1	Running	0	6s	app=my-app,environment=dev,tier=backend
www-dev	1/1	Running	0	6s	app=my-app,environment=dev,tier=frontend
www-prod	1/1	Running	0	6s	app=my-app,environment=production,tier=frontend

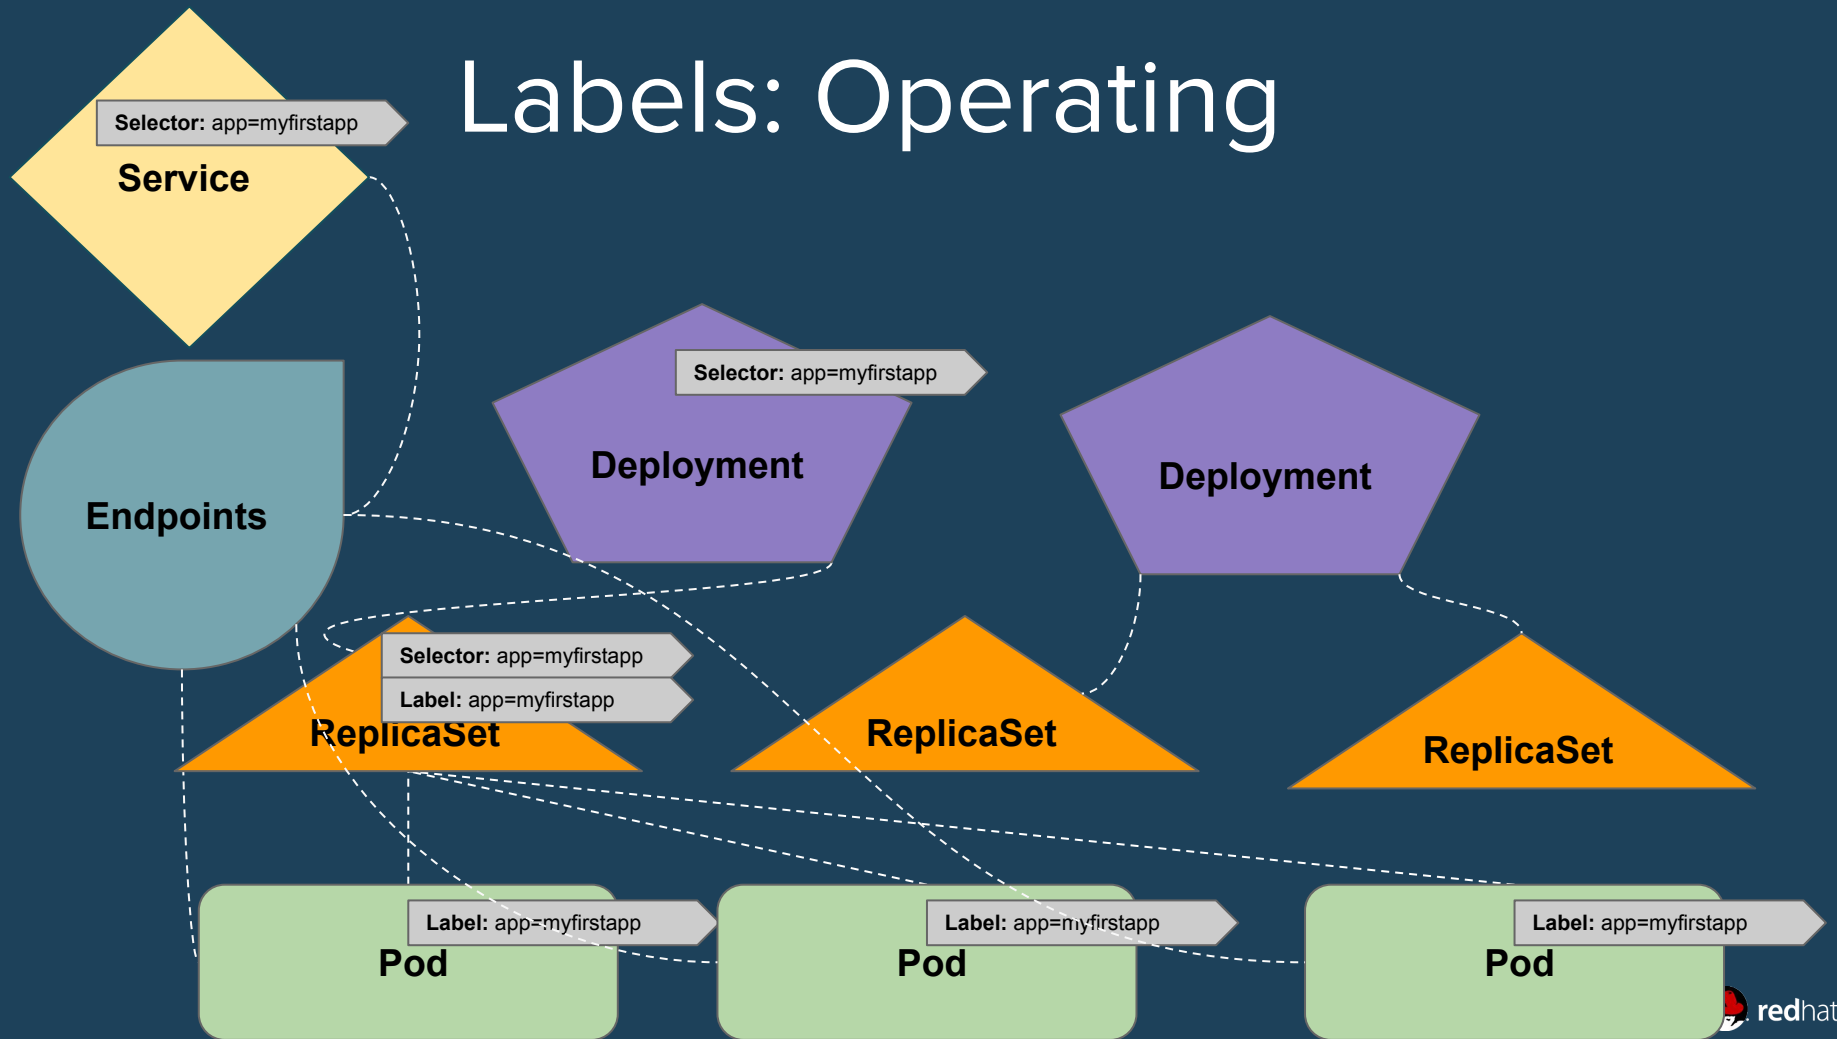
```
kubectl get pods -L app,environment,tier -l environment!=dev
```

www-prod	1/1	Running	0	4m	my-app	production	frontend
----------	-----	---------	---	----	--------	------------	----------

```
kubectl get pods -l "tier notin (backend,cache),environment in (dev)"
```

www-dev	1/1	Running	0	6m
---------	-----	---------	---	----

# Labels: Operating



# Using Labels

Labels can be displayed as a column in output with `-L` option:

```
$ kubectl get pods -L tier
```

NAME	READY	STATUS	RESTARTS	AGE	TIER
my-nginx-3800858182-1v53o	1/1	Running	0	46s	backend
my-nginx-3800858182-2ds1q	1/1	Running	0	46s	backend

# Updating Labels

Sometimes existing pods and other resources need to be relabeled before creating new resources

```
$ kubectl label pods -l app=nginx,tier=fe # select by label app=nginx, apply tier=fe
pod "my-nginx-v4-9gw19" labeled
pod "my-nginx-v4-hayza" labeled
pod "my-nginx-v4-mde6m" labeled
pod "my-nginx-v4-sh6m8" labeled
pod "my-nginx-v4-wfof4" labeled
```

```
$ kubectl get pods -l app=nginx -L tier
```

NAME	READY	STATUS	RESTARTS	AGE	TIER
my-nginx-v4-9gw19	1/1	Running	0	15m	fe
my-nginx-v4-hayza	1/1	Running	0	14m	fe
my-nginx-v4-mde6m	1/1	Running	0	18m	fe
my-nginx-v4-sh6m8	1/1	Running	0	19m	fe
my-nginx-v4-wfof4	1/1	Running	0	16m	fe

# Using Labels Effectively

Examples of multiple labels for app, tier and role:

```
labels:  
  app: guestbook  
  tier: frontend
```

```
labels:  
  app: guestbook  
  tier: backend  
  role: master
```

```
labels:  
  app: guestbook  
  tier: backend  
  role: slave
```

Other example labels:

- "release" : "stable" or "canary"
- "partition" : "customerA" or "customerB"
- "track" : "daily" or "weekly"



# Pods using a Configuration File

- Instead of using the `kubectl run` command to create a deployment object, you can specify a deployment YAML file
  - `kubectl` converts the YAML to JSON before sending it to the API Server
- Advantages over `kubectl run`:
  - Declarative (what to do) instead of imperative (how to do it)
  - Can track your changes in Git
  - Can have multiple containers in a pod
- Other Kubernetes objects (services, pods, etc.) can also be configured with YAML or JSON files

# Kubernetes API

# Pods using a Configuration File

- Instead of using the `kubectl run` command to create a deployment object, you can specify a deployment YAML file
  - `kubectl` converts the YAML to JSON before sending it to the API Server
- Advantages over `kubectl run`:
  - Declarative (what to do) instead of imperative (how to do it)
  - Can track your changes in Git
  - Can have multiple containers in a pod
- Other Kubernetes objects (services, pods, etc.) can also be configured with YAML or JSON files

# Not Flexible

`http://kubernetes:6443/api/v1/pods`

`http://kubernetes:6443/api/v1/replicasets`

`http://kubernetes:6443/api/v1/services`

`http://kubernetes:6443/api/v1/deployments`

# More Flexible

```
curl kubernetes:6443
```

```
"/api/v1"  
"/apis/authentication.k8s.io/v1"  
"/apis/authentication.k8s.io/v1beta1"  
"/apis/authorization.k8s.io/v1"  
"/apis/authorization.k8s.io/v1beta1"  
"/apis/certificates.k8s.io/v1beta1"  
"/apis/certificates.k8s.io"  
"/apis/extensions/v1beta1"  
"/apis/policy/v1beta1"  
"/apis/rbac.authorization.k8s.io/v1beta1"  
"/apis/rbac.authorization.k8s.io/v1alpha1"  
"/apis/storage.k8s.io/v1"  
"/apis/storage.k8s.io/v1beta1"
```

# API Versions

## Alpha (i.e. v1alpha1)

- Disabled by default. Must be enabled via API.
- May contain bugs. Features may be changed or removed.
- Only use for short-lived testing clusters.

## Beta (i.e. v1beta1)

- Enabled by default.
- Considered safe.
- Support for the feature will not be dropped, though details may change.

## Stable (i.e. v1,v2)

- Stable versions of features will appear in many subsequent versions.

# See current Api Versions available.

```
kubectl api-versions
```

```
"/api/v1"  
"/apis/authentication.k8s.io/v1"  
"/apis/authentication.k8s.io/v1beta1"  
"/apis/authorization.k8s.io/v1"  
"/apis/authorization.k8s.io/v1beta1"  
"/apis/certificates.k8s.io/v1beta1"  
"/apis/certificates.k8s.io"  
"/apis/extensions/v1beta1"  
"/apis/policy/v1beta1"  
"/apis/rbac.authorization.k8s.io/v1beta1"  
"/apis/rbac.authorization.k8s.io/v1alpha1"  
"/apis/storage.k8s.io/v1"  
"/apis/storage.k8s.io/v1beta1"
```

# TypeMeta, ObjectMeta, Spec

TypeMeta

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
```

ObjectMeta

```
metadata:
  name: my-first-replica-set
  namespace: myproject
```

Spec

```
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 5
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
```



# News Snippet About Introduction of v1 NetworkPolicy

## Two of the changes you need to be aware of are:

### » The v1beta1 NetworkPolicy API Has Been Deprecated

The v1beta1 version of the NetworkPolicy API has been deprecated in favor of moving forward with the new behaviors and updating the behavior of the *extensions* API to allow for future expansion and development. Keep in mind that while the v1 NetworkPolicy API eclipses the existing beta, the new API endpoint will only be available on Kubernetes 1.7+ (as older versions do not include the v1 API code). As such, as you work towards upgrading, you'll want to ensure that you are using the correct version of Project Calico for the NetworkPolicy behavior you want.

### » The DefaultDeny Annotation Has Been Removed

One of the bigger changes in Kubernetes 1.7 is the removal of the DefaultDeny annotation. This means that when upgrading, you should **first delete any existing NetworkPolicy** objects in namespaces that previously **did not have** the "DefaultDeny" annotation (as this may cause Kubernetes to unintentionally block traffic now).

# Kubernetes API Actions and HTTP Method

<u>Verb</u>	<u>HTTP Method</u>
Get	GET
List	GET
Watch	GET
Create	POST
Update	PUT
Patch	PATCH
Delete	DELETE

# Kubernetes Subcommand & HTTP Method

<u>Subcommand</u>	<u>Object does not exist</u>	<u>Object exists</u>
apply	POST	PATCH/DELETE
create	POST	error!
replace	error!	PUT
delete	error!	DELETE
Patch	PATCH	PATCH
Delete	DELETE	DELETE

# Interacting with the API

Create!

```
kubectl/oc create -f podmanifest.json
```

```
curl -X POST http://localhost:8001/api/v1/namespaces/myproject/pods/ -H  
"Content-type: application/json" -d @podmanifest.json
```

Update!

```
kubectl/oc replace -f podmanifest.json
```

```
curl -X PUT http://localhost:8001/api/v1/namespaces/myproject/pods/mypod -H  
"Content-type: application/json" -d @newpodmanifest.json
```

Patch!

```
kubectl/oc patch -f patch.json
```

```
kubectl patch etcdcluster example-etcd-cluster --type='json' -p '[{"op":  
"replace", "path": "/spec/size", "value":5}]'
```

# ReplicaSets

# Redundancy

Multiple running instances means failure can be tolerated.



# Scale

Multiple running instances mean more requests can be handled.



000  
cpu:80%

cpu:20%

001  
cpu:75%

cpu:25%

002  
cpu:70%

cpu:23%


003

004



# Sharding


Multiple running instances can handle different parts of a computation in parallel.





ReplicaSet is just a  
Reconciliation Loop.

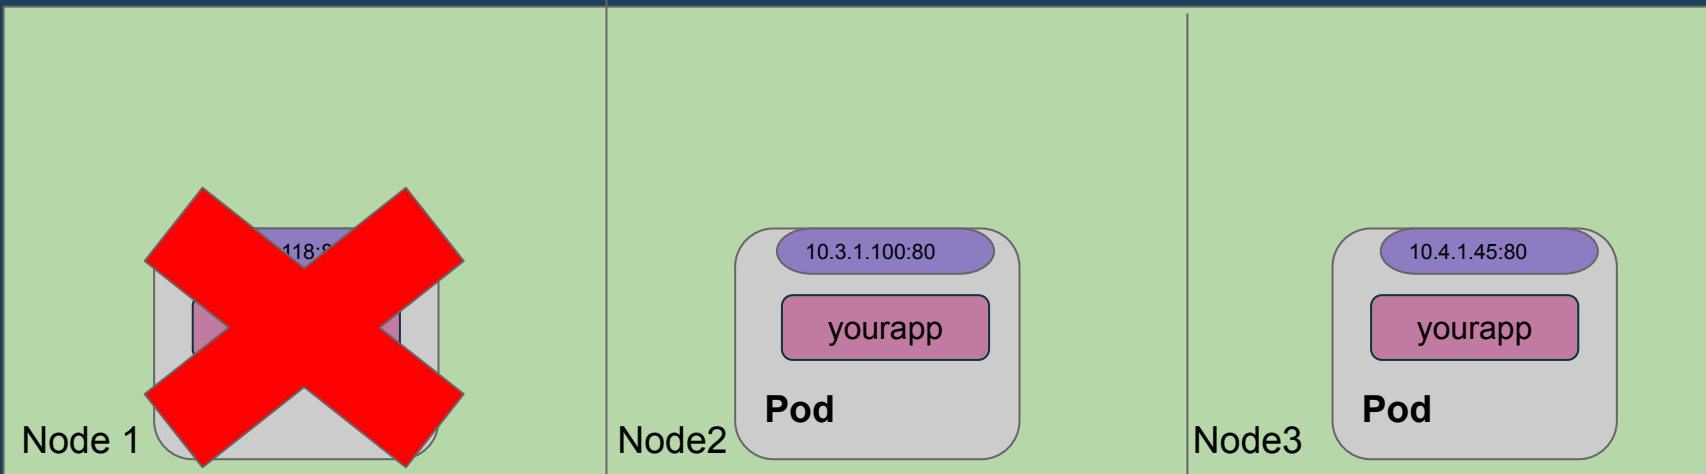
# Edge Driven vs. Level Driven.



```
1  while true {  
2  receiveInfoAboutAPIObjects()  
3  synchronizeRealStateToMatchFetchedInfo()  
4  }
```

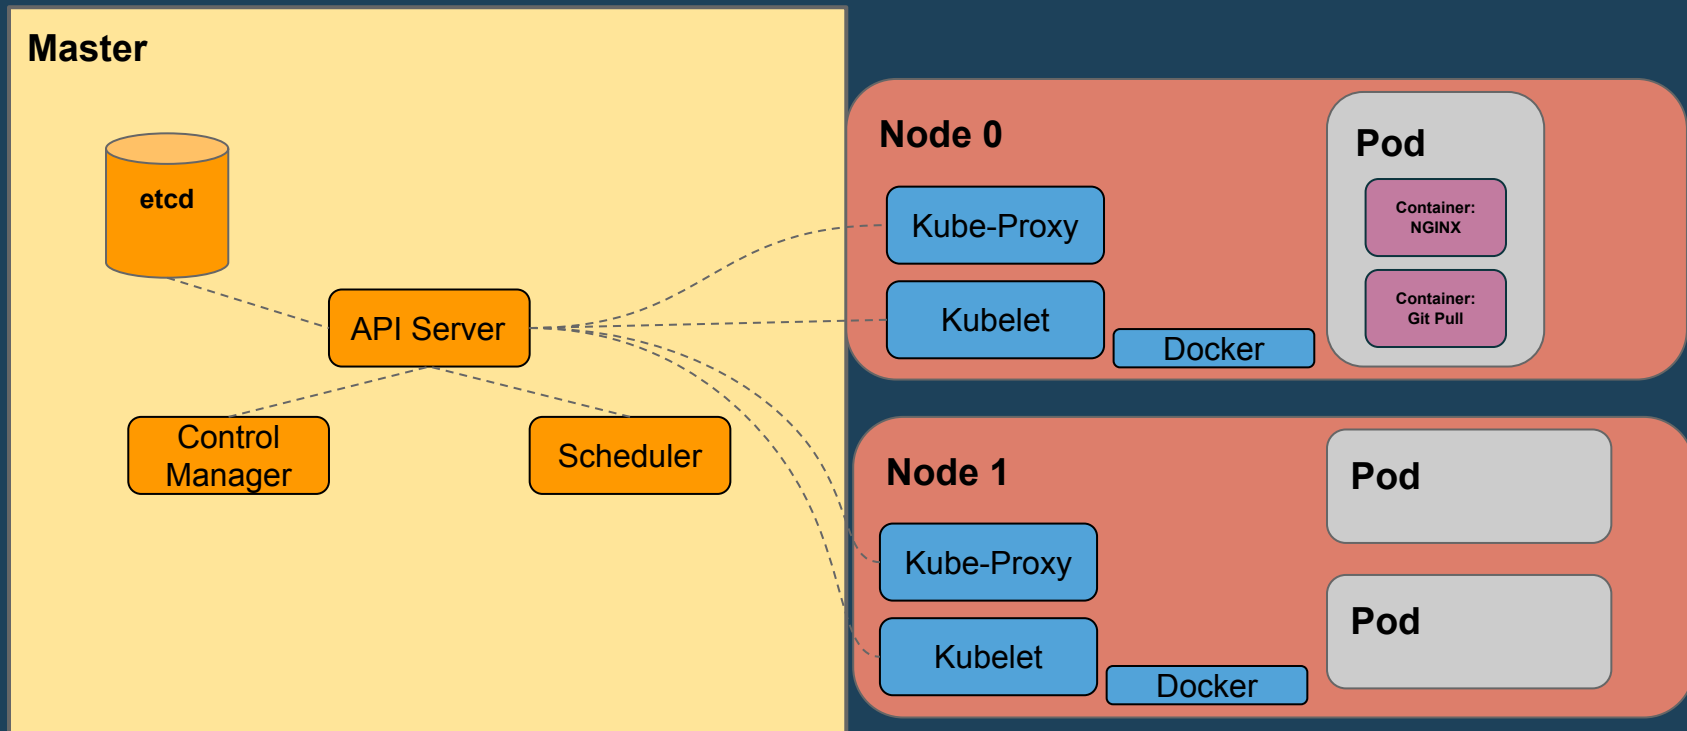
```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: my-first-replica-set
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
```

```
status:
  availableReplicas: 2
  fullyLabeledReplicas: 5
  observedGeneration: 724
  readyReplicas: 2
  replicas: 5
```



```
1  while true {  
2  receiveInfoAboutAPIObjects()  
3  synchronizeRealStateToMatchFetchedInfo()  
4  }
```

# Kubernetes Cluster



# Desired State and Actual State

- The replicaset object contains a desired state object (spec) from the user and the actual state from the ReplicaSet controller (status)



# Desired State vs. Current State

## Desired State:

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: my-first-replica-set
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 5
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
```

## Current State:

```
status:
  availableReplicas: 2
  fullyLabeledReplicas: 5
  observedGeneration: 724
  readyReplicas: 2
  replicas: 5
```



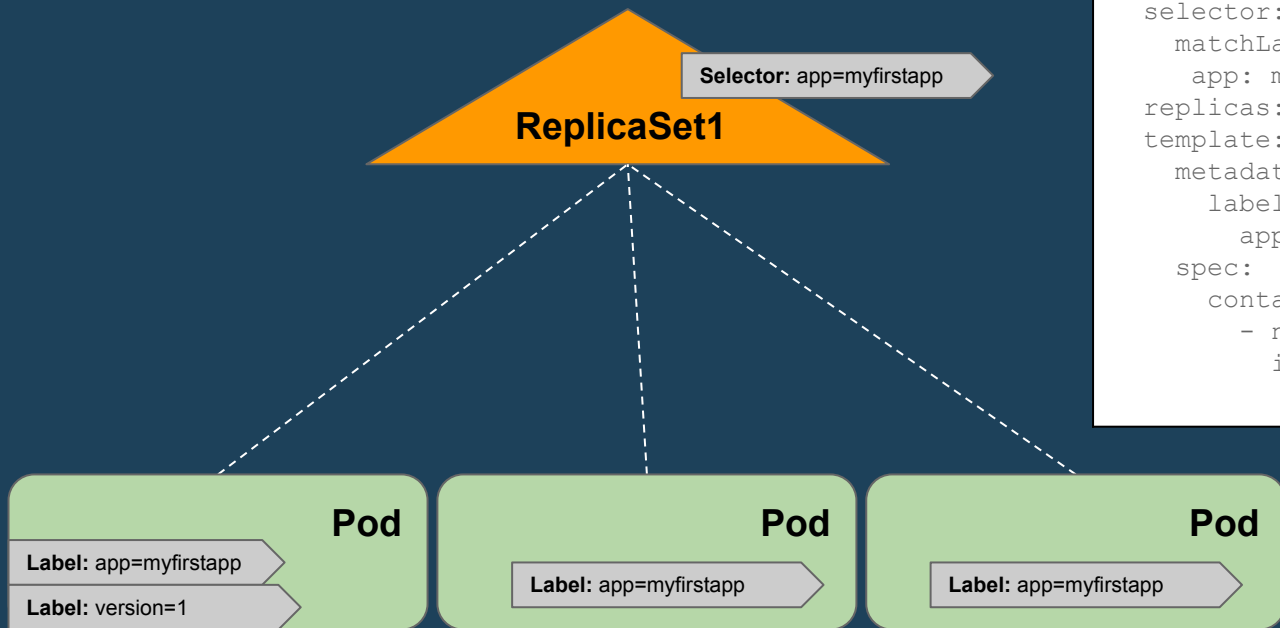
# ReplicaSets in Action!

```
kubectl run myfirstapp --image quay.io/coreosstrainme/hello-whoami:2.0.1 --restart=Never -l app=myfirstapp,version=1
```

```
kubectl create -f myfirstreplicaset.yaml
```

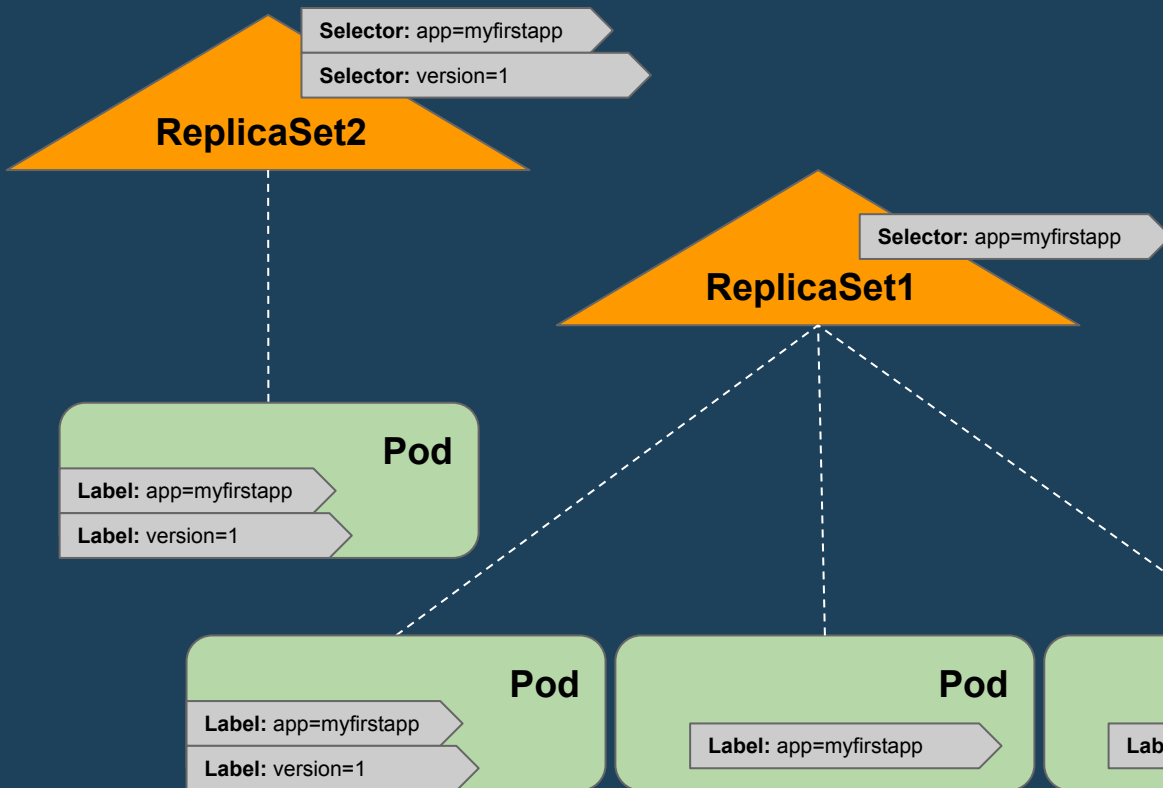
```
kubectl scale replicaset myfirstreplicaset --replicas=3
```

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: myfirstreplicaset
spec:
  selector:
    matchLabels:
      app: myfirstapp
  replicas: 3
  template:
    metadata:
      labels:
        app: myfirstapp
    spec:
      containers:
        - name: nodejs
          image: myimage
```



# Creating another ReplicaSet

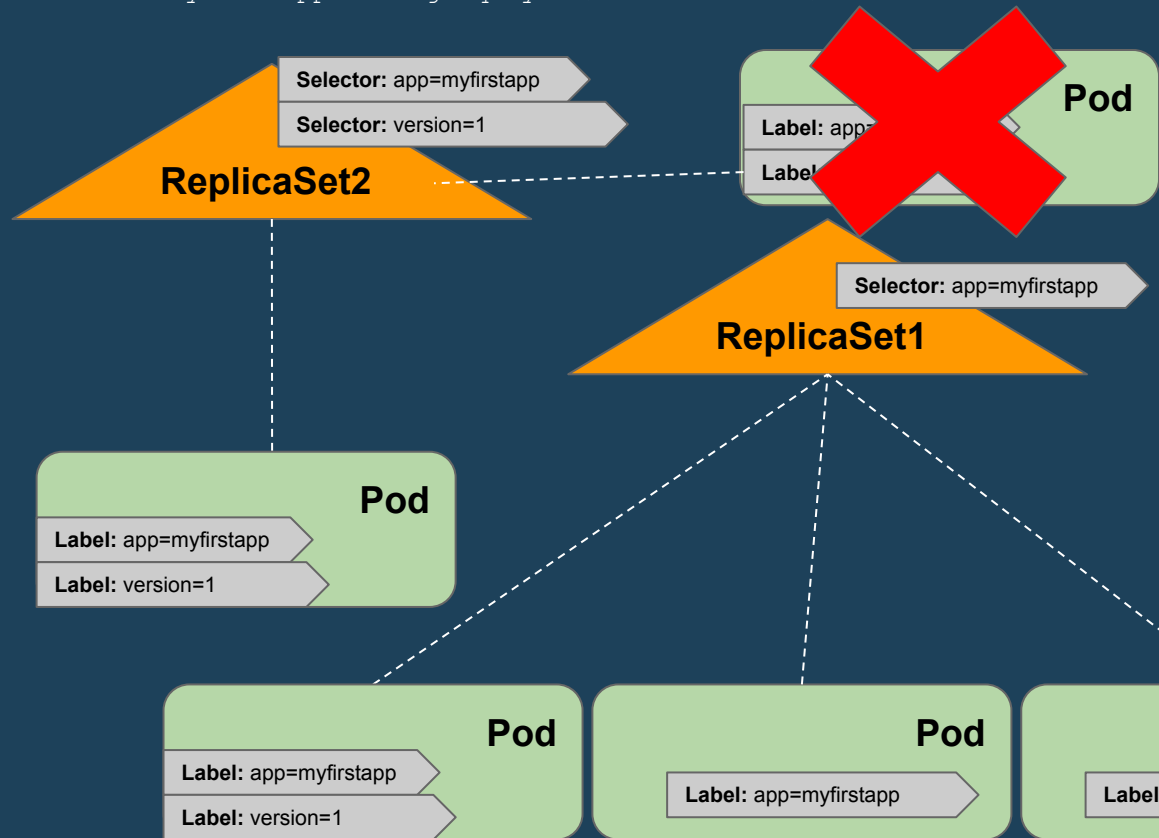
```
kubectl create -f mysecondreplicaset.yaml
```



```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: mysecondreplicaset
spec:
  selector:
    matchLabels:
      app: myfirstapp
      version: 1
  replicas: 1
  template:
    metadata:
      labels:
        app: myfirstapp
        version: 1
    spec:
      containers:
        - name: nodejs
          image: myimage
```

# Creating Another Orphan Pod

```
kubectl run myfirstapp --image quay.io/coreosstrainme/hello-whoami:2.0.1 --restart=Never -l app=myfirstapp,version=1
```



```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: myfirstreplicaset
spec:
  selector:
    matchLabels:
      app: myfirstapp
      version: 1
  replicas: 1
  template:
    metadata:
      labels:
        app: myfirstapp
        version: 1
    spec:
      containers:
        - name: nodejs
          image: myimage
```

# ReplicaSets

Behind the scenes, replicas (pod copies) in a deployment are represented and managed by ReplicaSets (RS)

- RS's job is to create/recreate/destroy pods when needed from a template in the manifest
- Results in self-healing and application high availability

```
$ kubectl get replicasets
```

NAME	DESIRED	CURRENT	AGE
nginx-684635458	2	2	31m

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-684635458-9p0kg	1/1	Running	0	31m
nginx-684635458-ibp6q	1/1	Running	0	31m

# Scaling your Application: Autoscale

Kubernetes automatically chooses the number of nginx replicas as needed, from one to three in this example

```
$ kubectl autoscale replicaset my-nginx --cpu-percent=80 --min=1 --max=3
```

```
replicaset "my-nginx" autoscaled
```

```
$ kubectl get pods -lapp=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
my-nginx-1jgkf	1/1	Running	0	3m
my-nginx-divi2	1/1	Running	0	3m

```
$ kubectl get horizontalpodautoscaler
```

NAME	REFERENCE	TARGET	CURRENT	MINPODS	MAXPODS	AGE
nginx	Replicaset/my-nginx	80%	39%	1	3	1m

kubectl scale

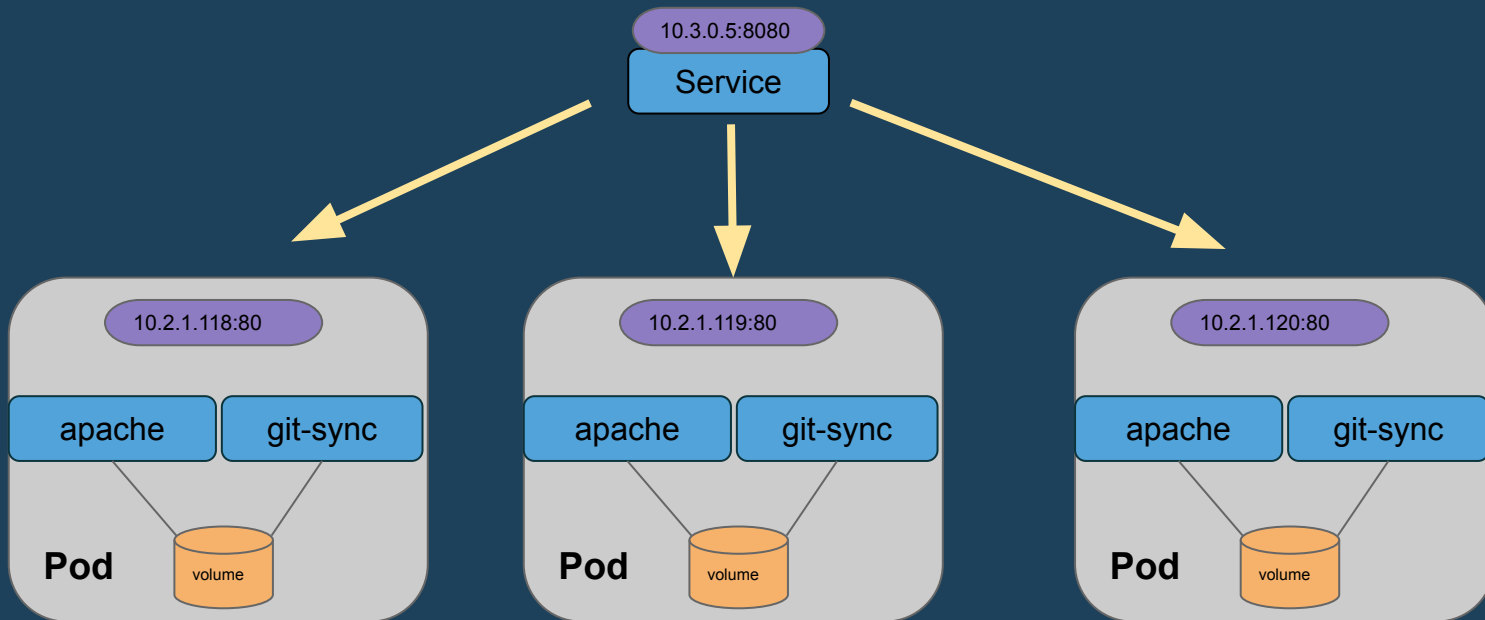
Update the size of the specified replication controller.

```
kubectl scale (-f FILENAME | TYPE NAME | TYPE/NAME) --replicas=COUNT  
[--resource-version=version] [--current-replicas=count] [flags]
```

# Services

# Services

- Persistent IPs for Pods
- Load Balances Between Replicas





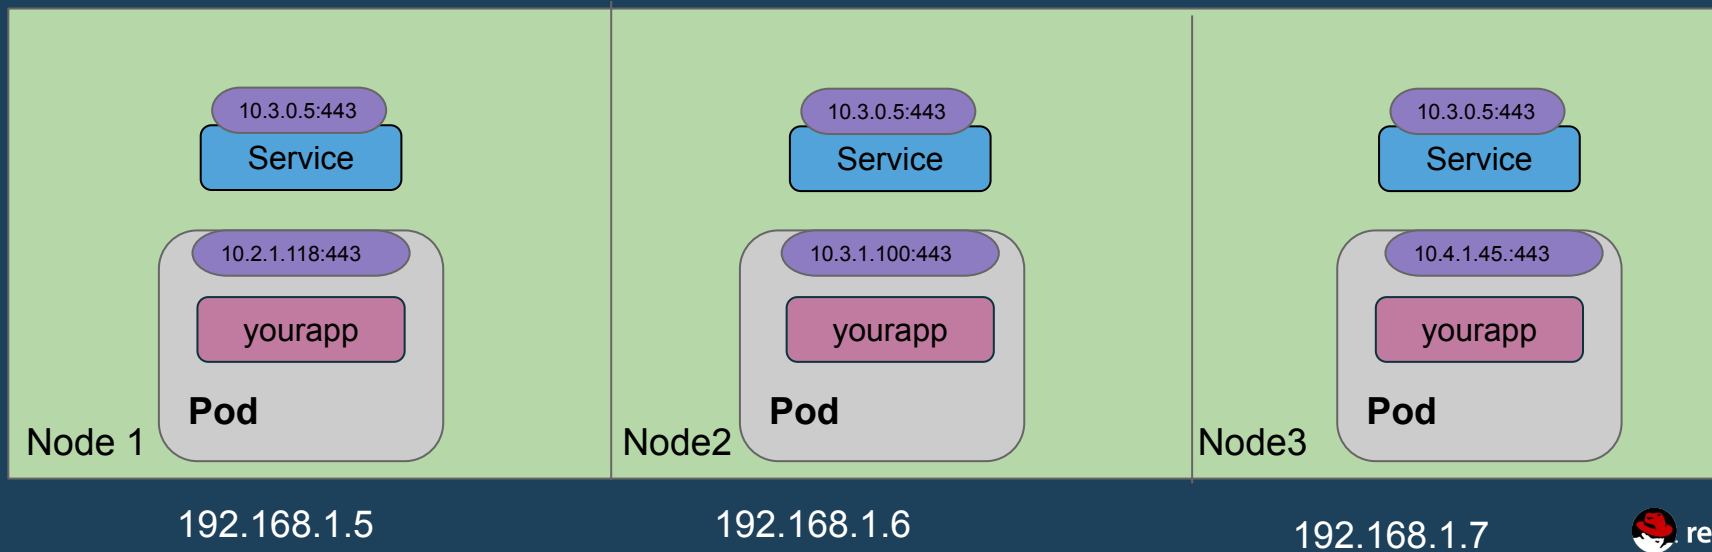
# Service Types

- ClusterIP (shown previously)
  - Exposes a service using an internal IP accessible only in the Kubernetes cluster
- NodePort
  - Exposes a service at every node <NODE\_IP>:<NODE\_PORT>
- LoadBalancer
  - Works with a cloud provider to create a load balancer and rules to expose the service as a layer on top of NodePort

# ClusterIP

```
Kubectl create -f deployment.yaml
```

```
Kubectl create -f myservice.yaml
```




# Service Discovery

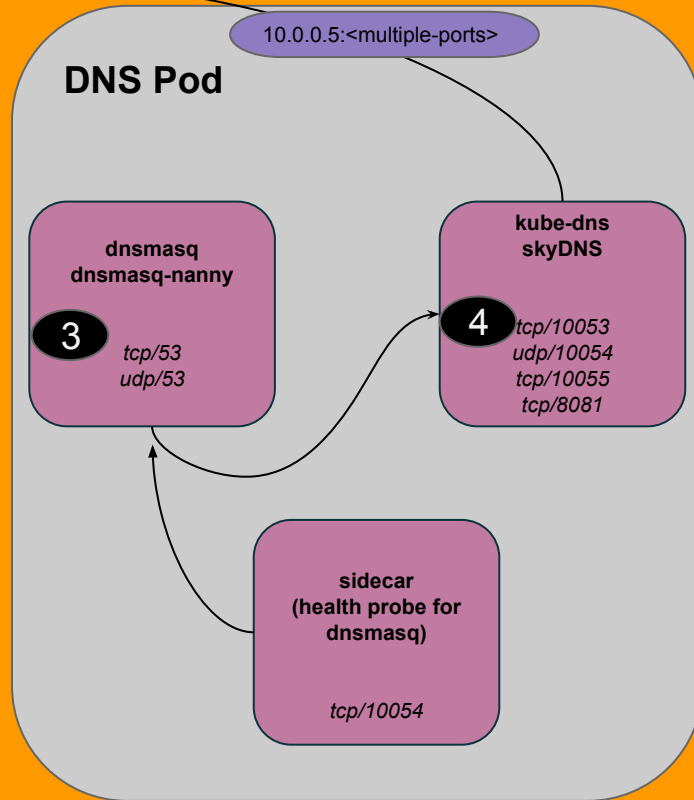
## Kubernetes DNS

- Schedules a DNS Pod and Service on the cluster.
- Configures the kubelets to tell individual containers to use the DNS Service's IP to resolve DNS names.

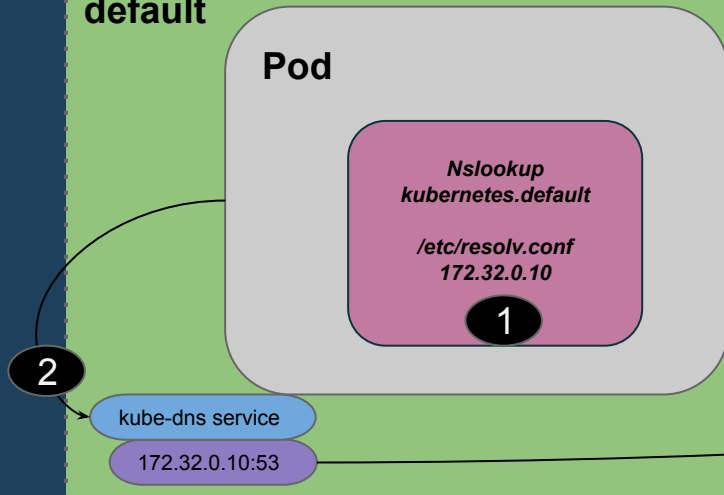
## Control Plane



## kube-system



## default

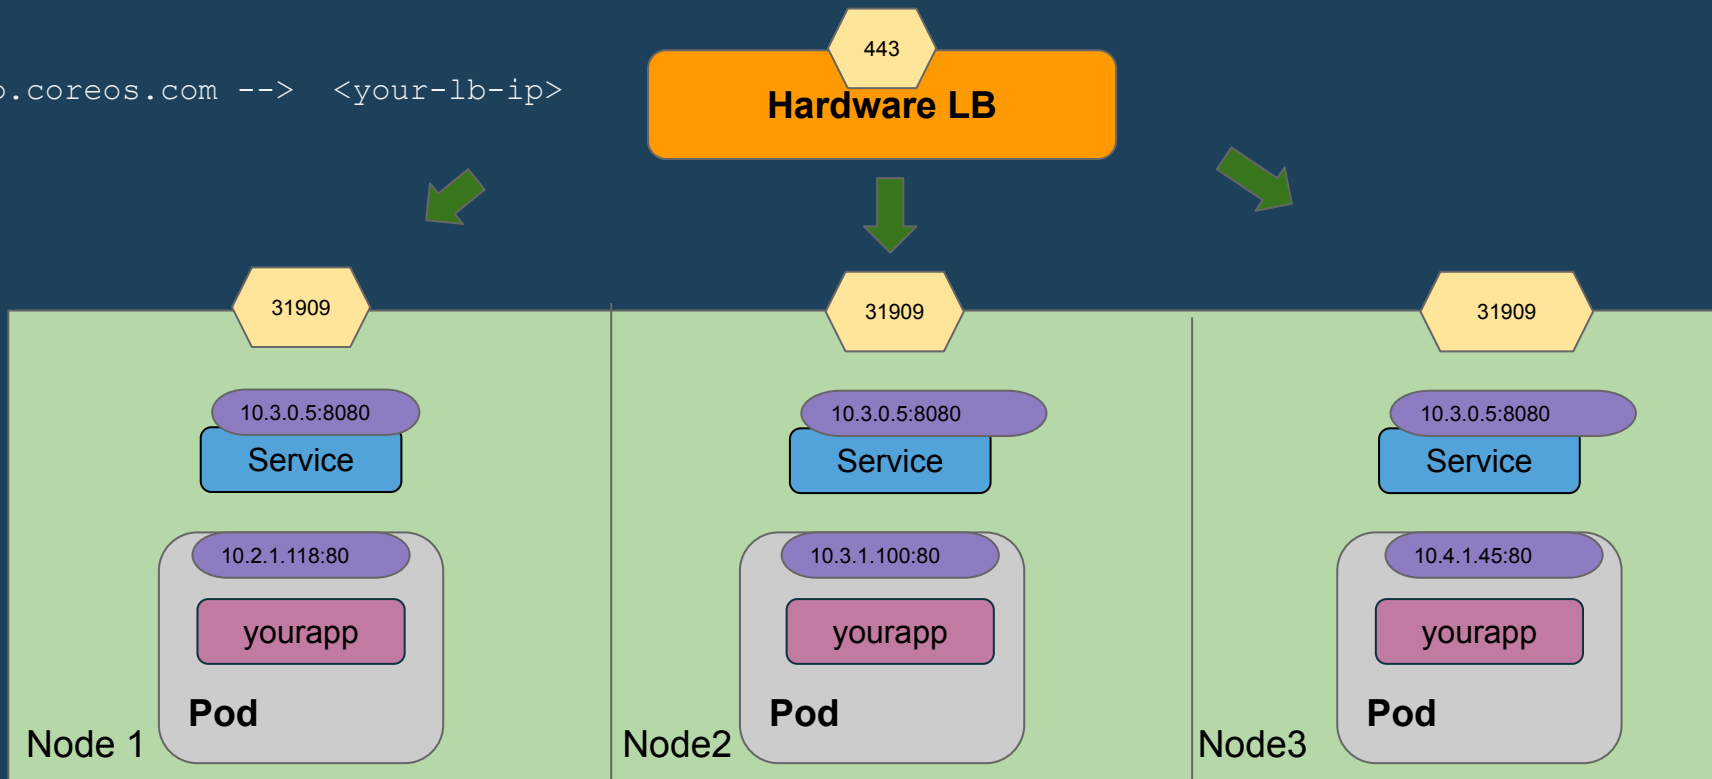


# NodePort

```
Kubectl create -f deployment.yaml
```

```
Kubectl create -f myservice.yaml
```

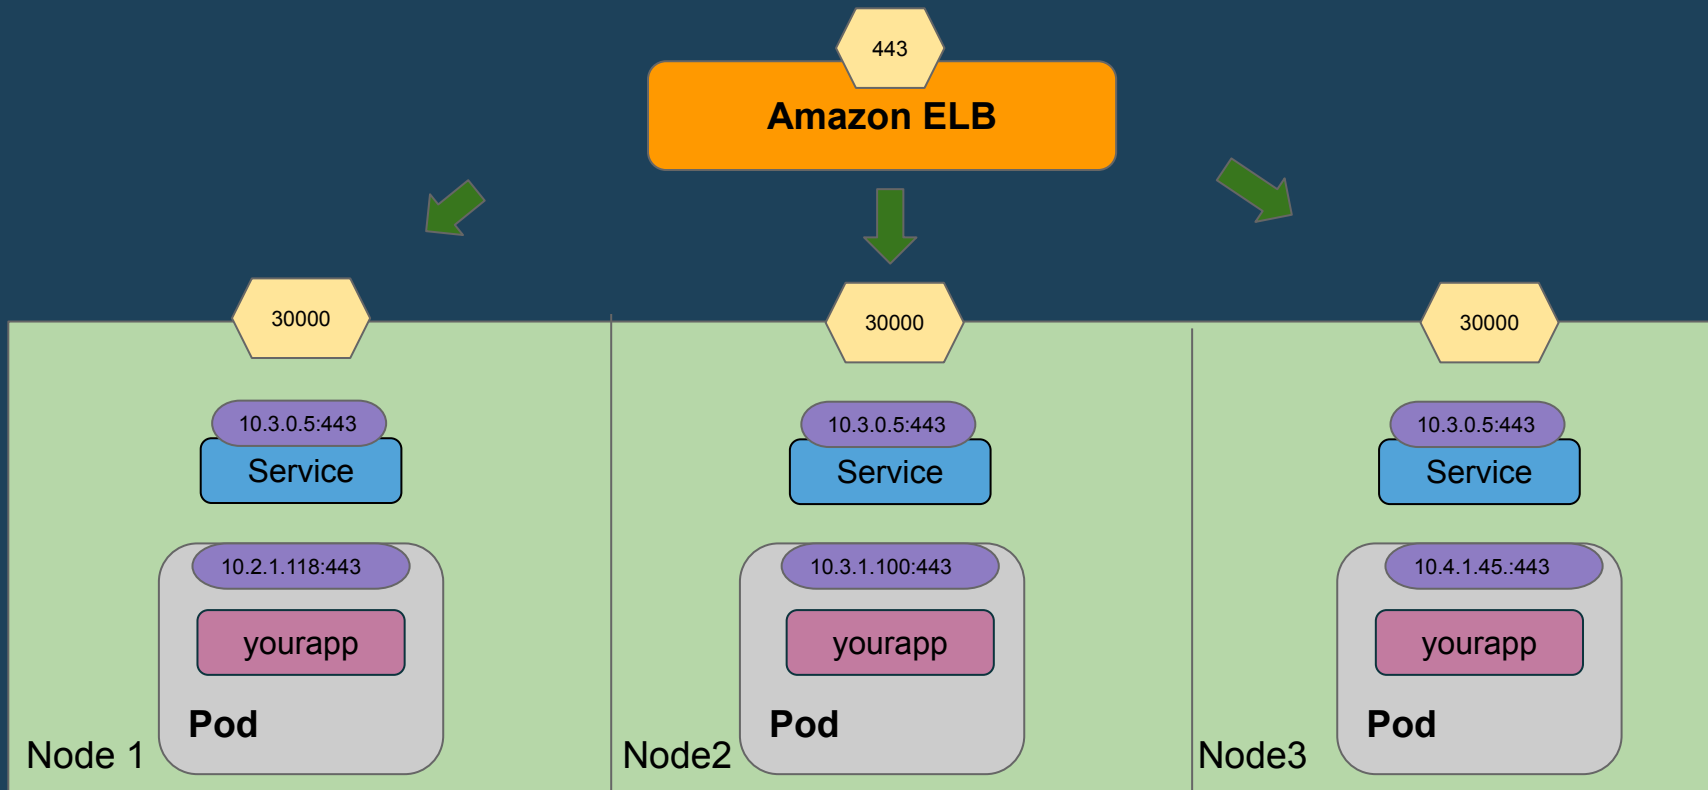
```
myapp.coreos.com --> <your-lb-ip>
```



# LoadBalancer

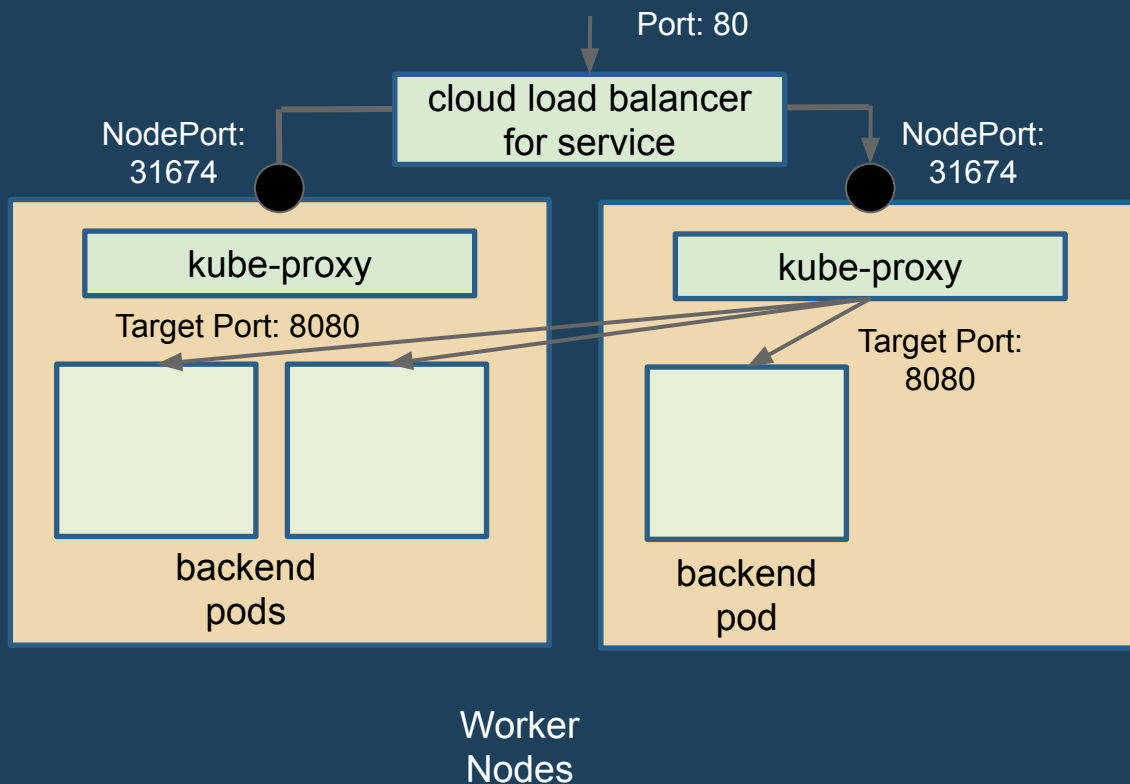
```
Kubectl create -f deployment.yaml
```

```
Kubectl expose... --port 443 --port 443
```



# Example Service with LoadBalancer

- A cloud load balancer is created for every service
  - myservice.mydomain.com
- The load balancer picks a worker node and sends the request at NodePort 31674
  - This is configured in the load balancer
- kube-proxy identifies the service by port and load balances among all service pods






# Deployments

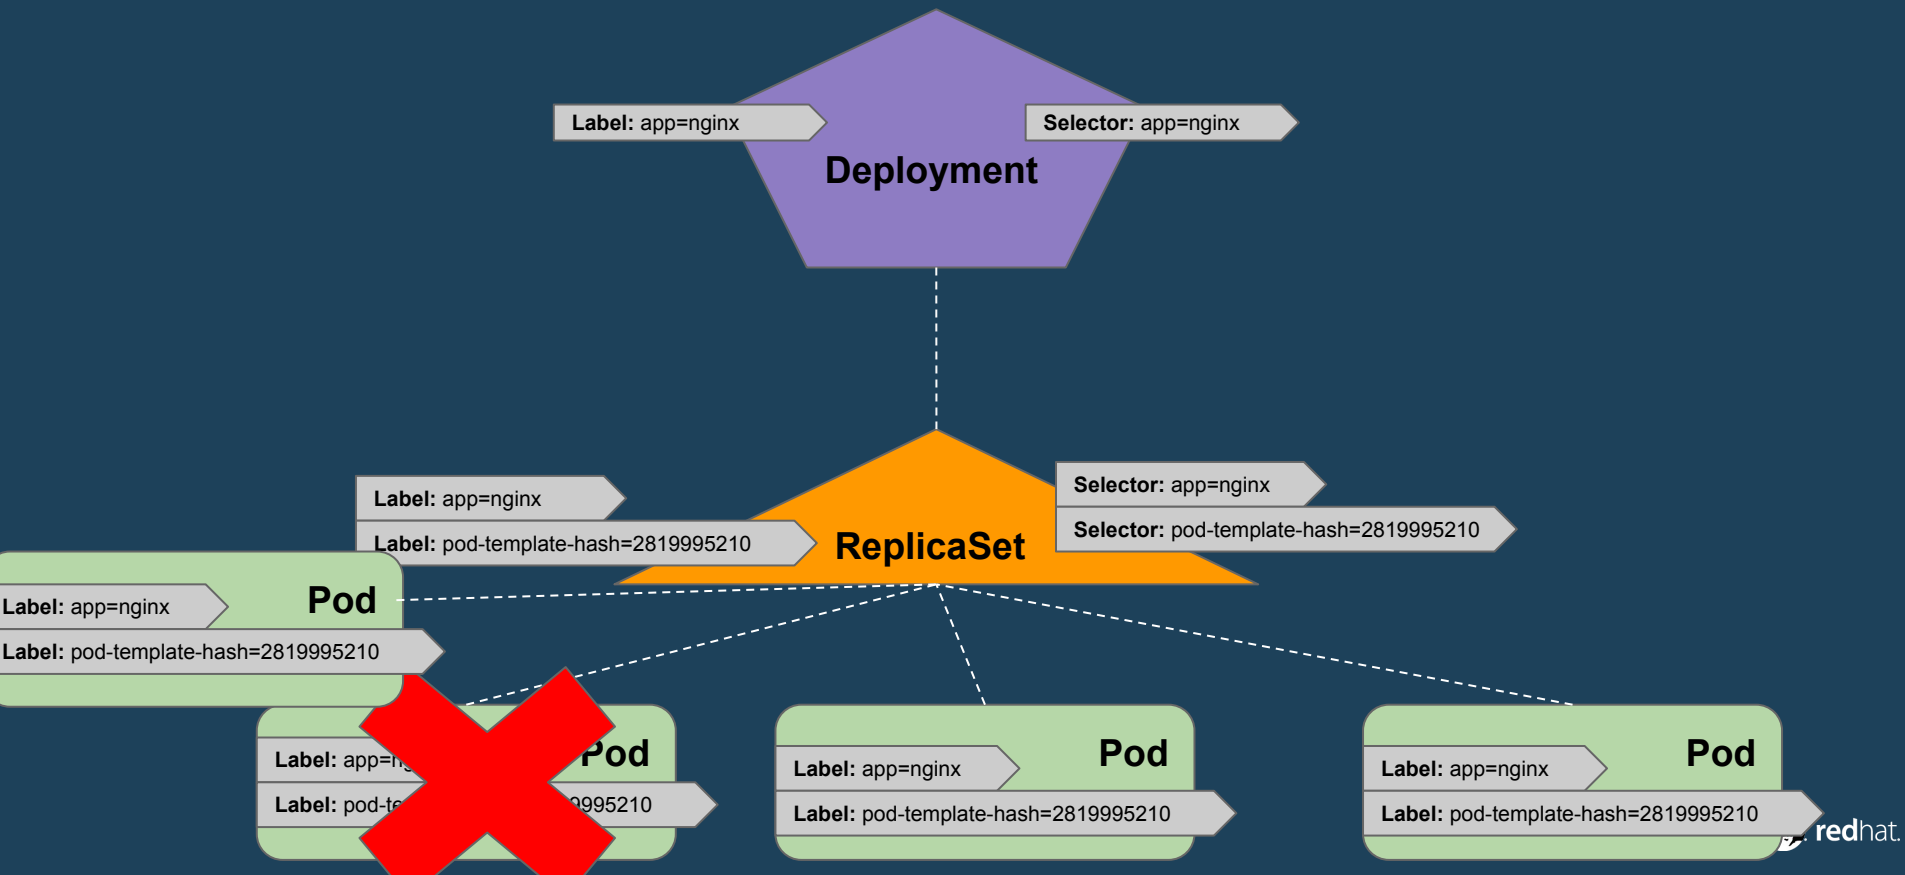
# Sample Deployment Manifest file.

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80
```

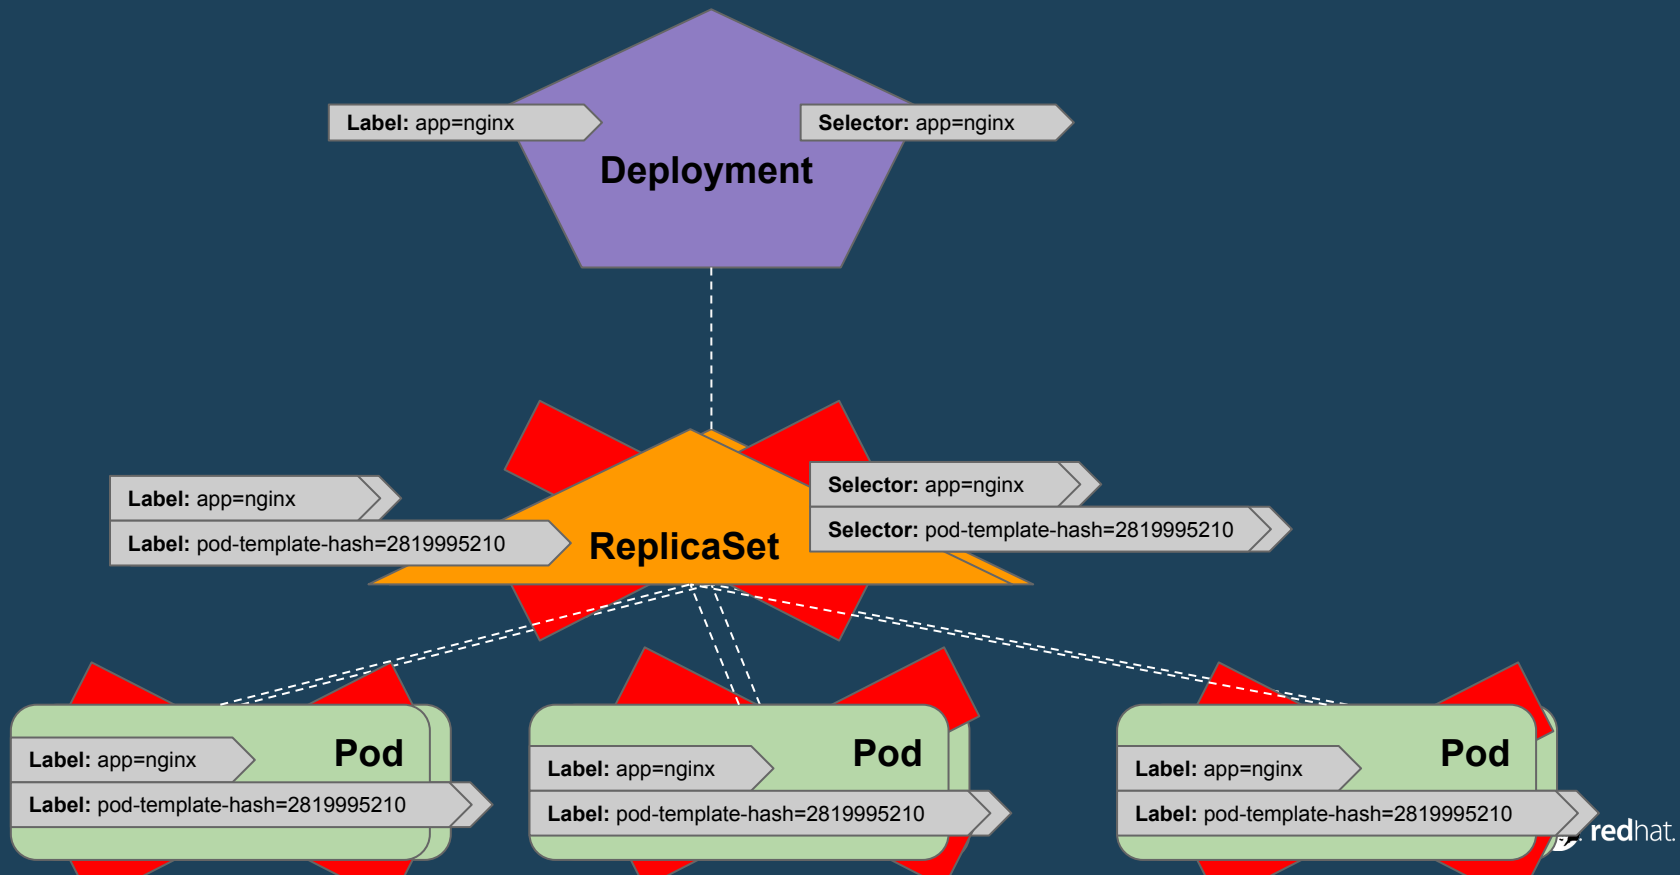
# Deployments: Labels & Selectors



# Deleting a Pod....



# Deleting a ReplicaSet...



This is nice.

But to really grasp the power of Deployments,  
consider how we would “roll out” a new application  
to an existing ReplicaSet.

# Rolling out a new app the “old way”.

- 1) Created ReplicaSet2.
- 2) Scaling ReplicaSet1 down to 2.
- 3) Scaling ReplicaSet2 up to 1.
- 4) Scaling ReplicaSet1 down to 1.
- 5) Scaling ReplicaSet2 up to 2.
- 6) Scaling ReplicaSet1 down to 0.
- 7) Scaling ReplicaSet2 up to 3.

kube-api

Service

ReplicaSet2

Pod - v2

Pod - v2

Pod - v2

What's wrong with this picture?



# Issues with “Old” Rolling Updates

kube-api

Update takes place client-side.

Limited control over update strategy.

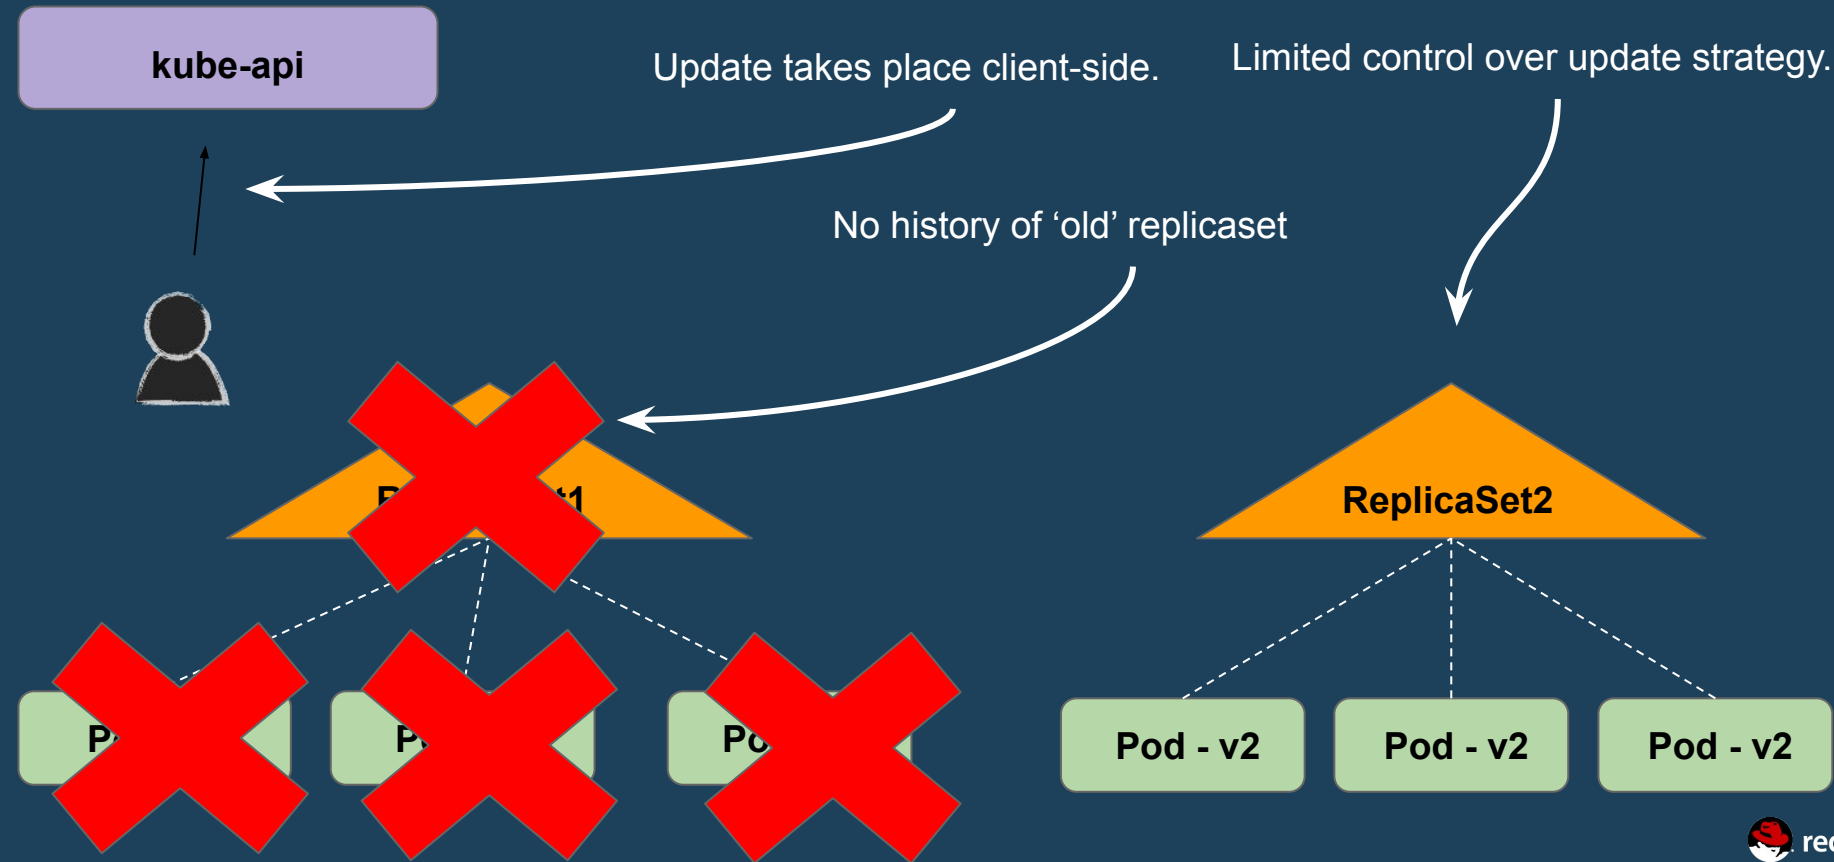
No history of ‘old’ replicaset

ReplicaSet2

Pod - v2

Pod - v2

Pod - v2



# Deployments to the Rescue.

```
kubectl run nginx-deployment --image=nginx:1.7.9 --replicas=3
```

# Deployment Rollout Strategies

# 2

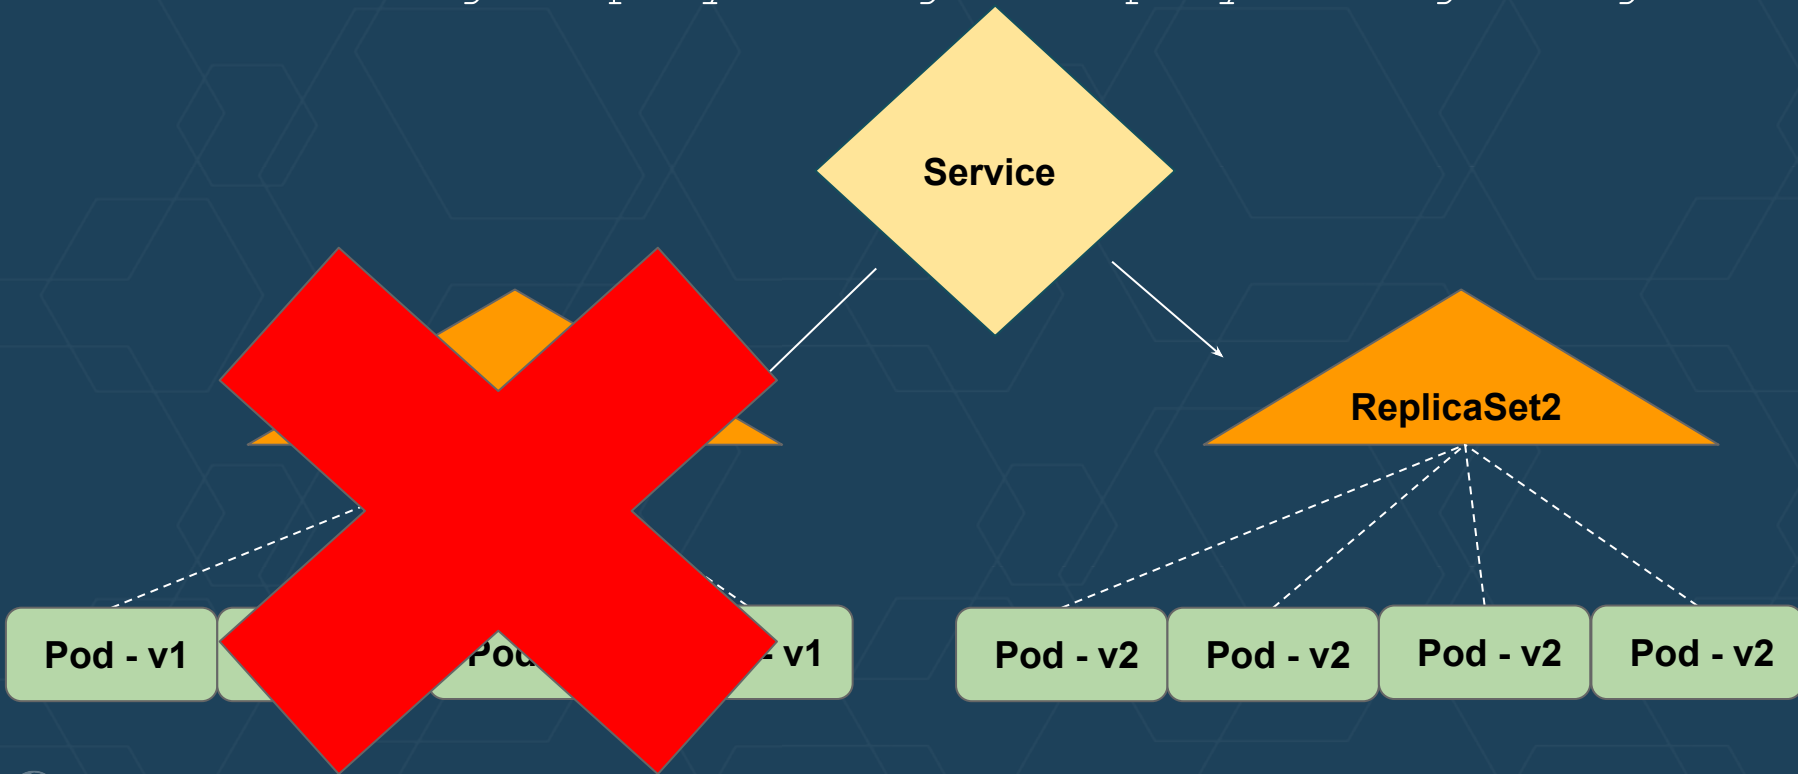
1) Recreate

2) RollingUpdate

If you don't care about downtime...

# Strategy: Recreate

```
kubectl set image deployment nginx-deployment nginx=nginx:1.12
```



If you care about downtime...

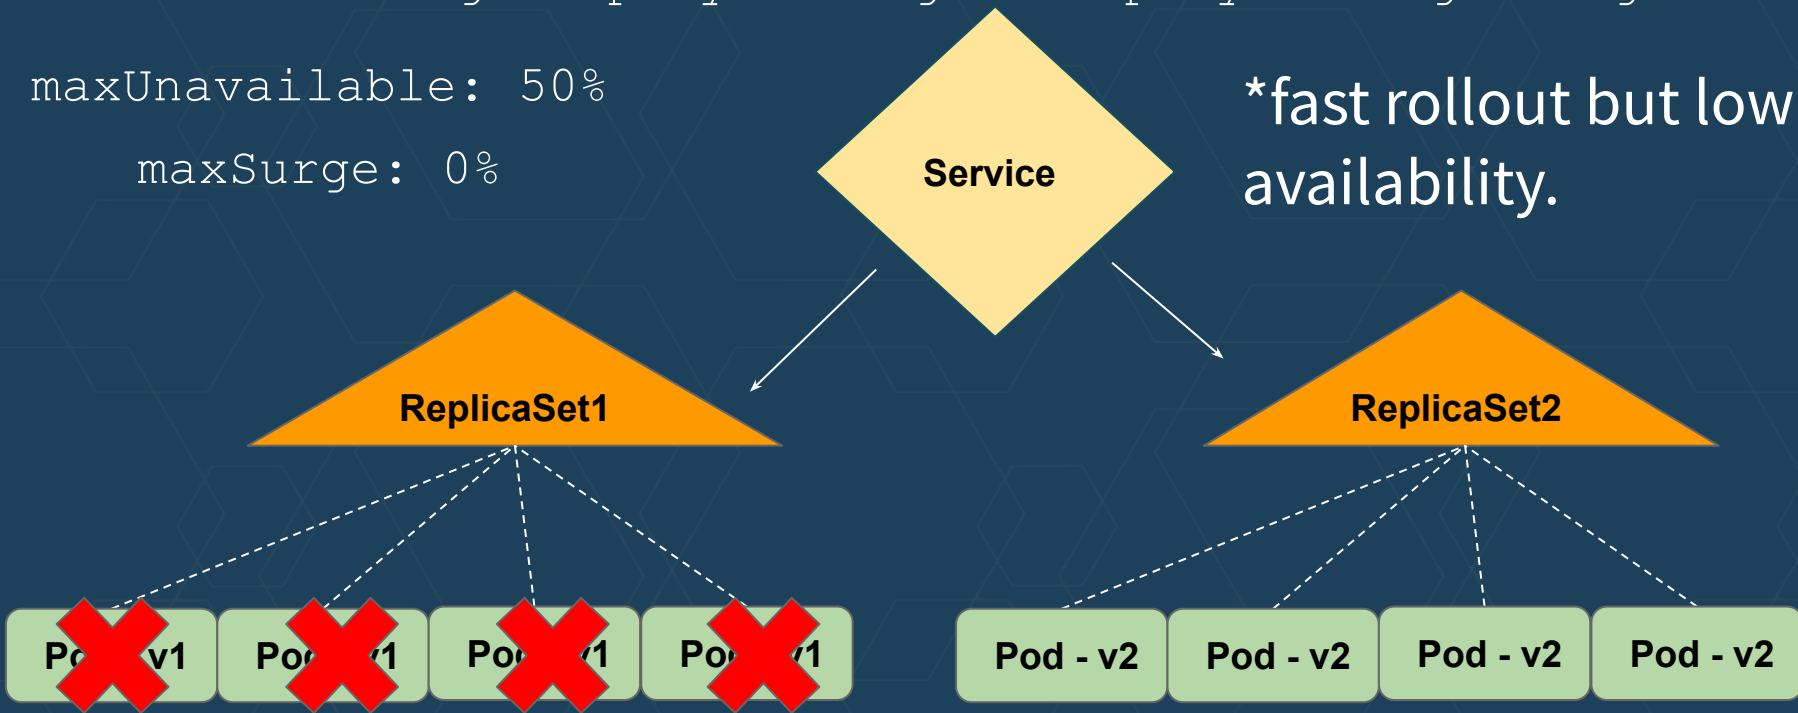
# Strategy: RollingUpdate

```
kubectl set image deployment nginx-deployment nginx=nginx:1.12
```

maxUnavailable: 50%

maxSurge: 0%

\*fast rollout but low availability.





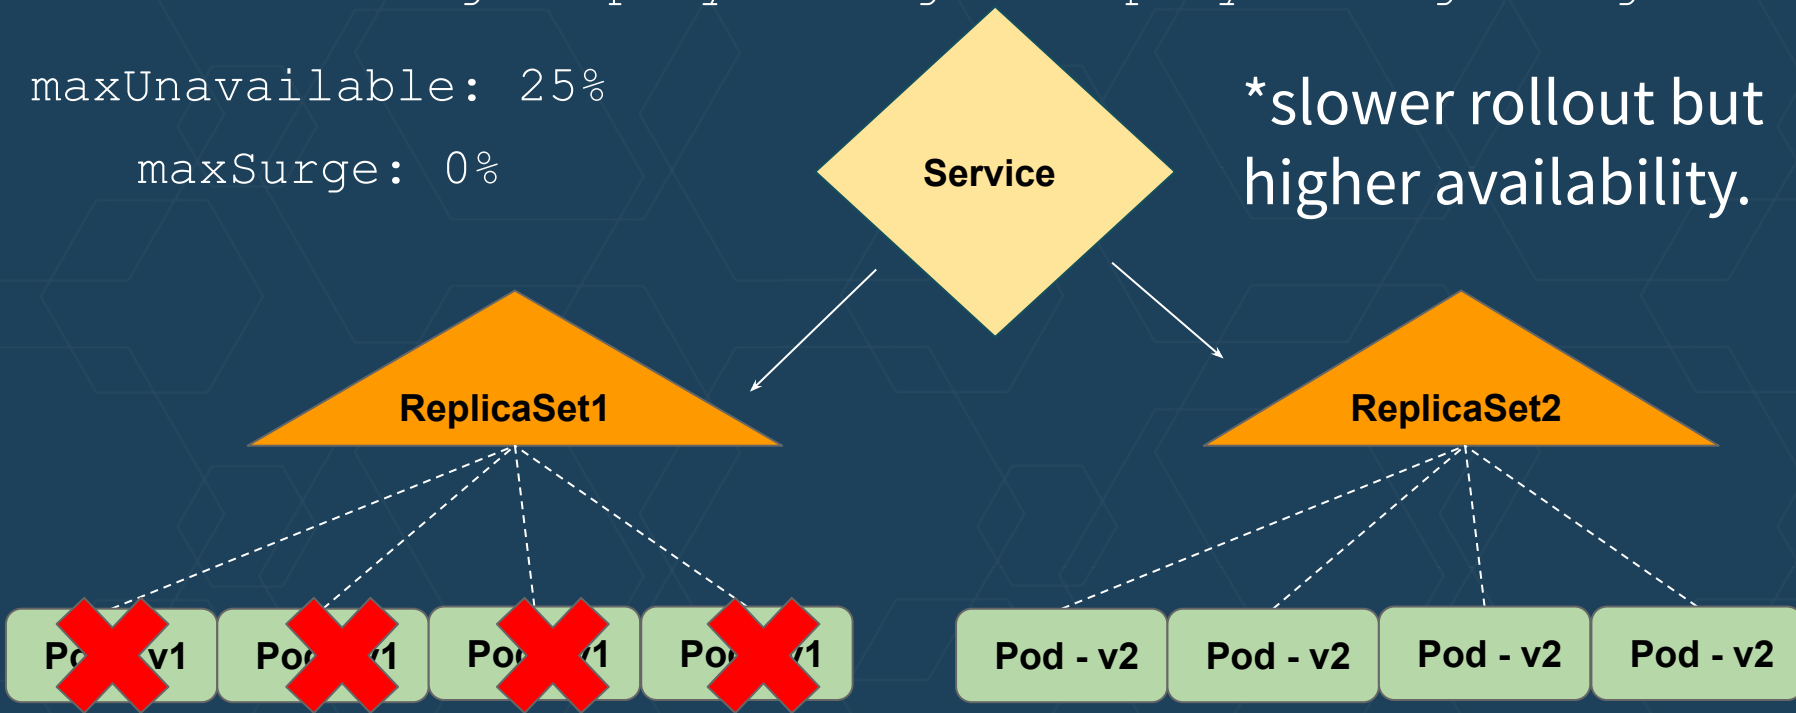
# Strategy: RollingUpdate

```
kubectl set image deployment nginx-deployment nginx=nginx:1.12
```

maxUnavailable: 25%

maxSurge: 0%

\*slower rollout but  
higher availability.



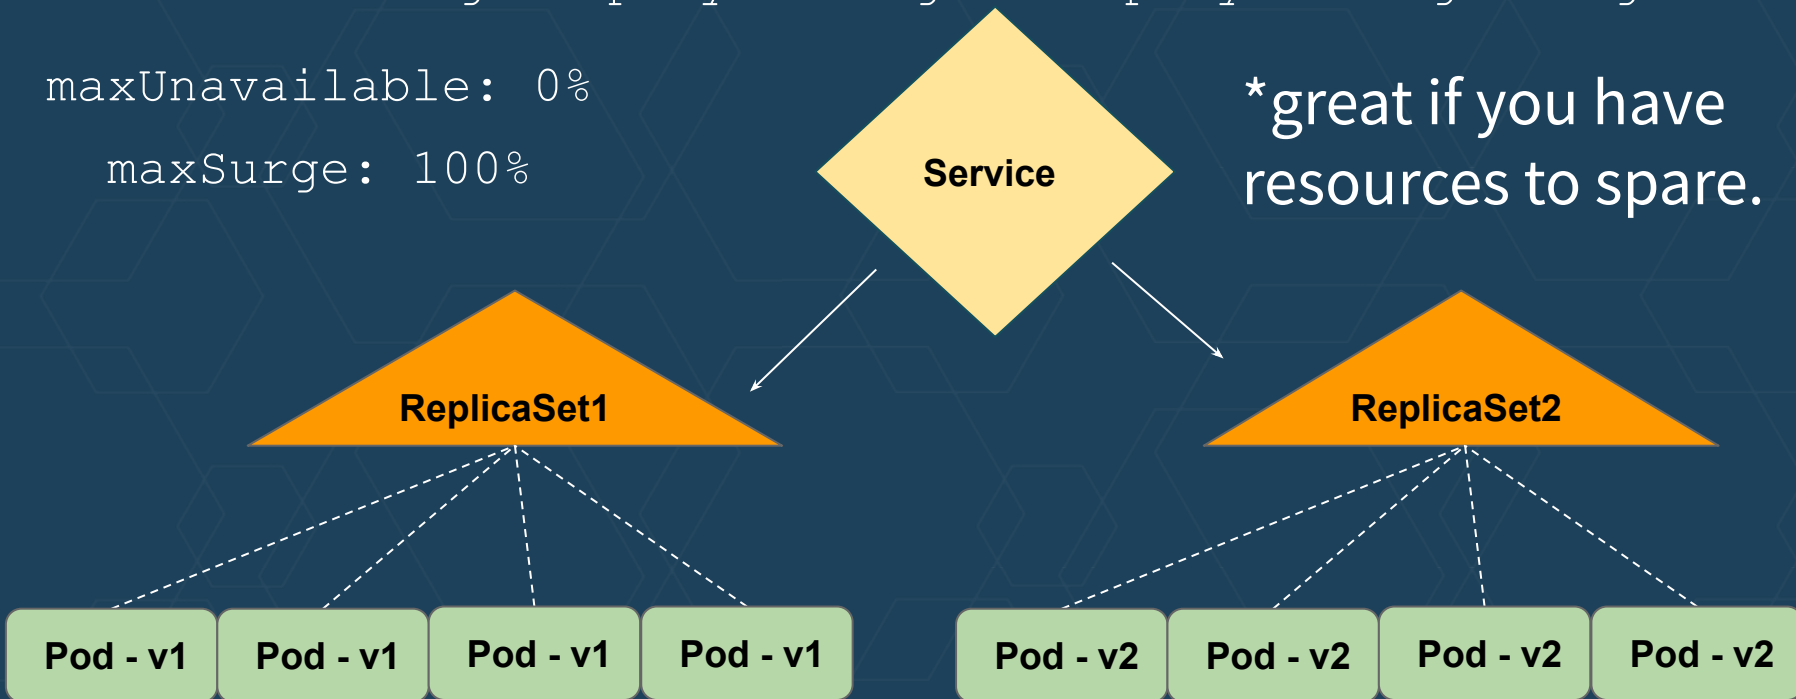
# Strategy: RollingUpdate

```
kubectl set image deployment nginx-deployment nginx=nginx:1.12
```

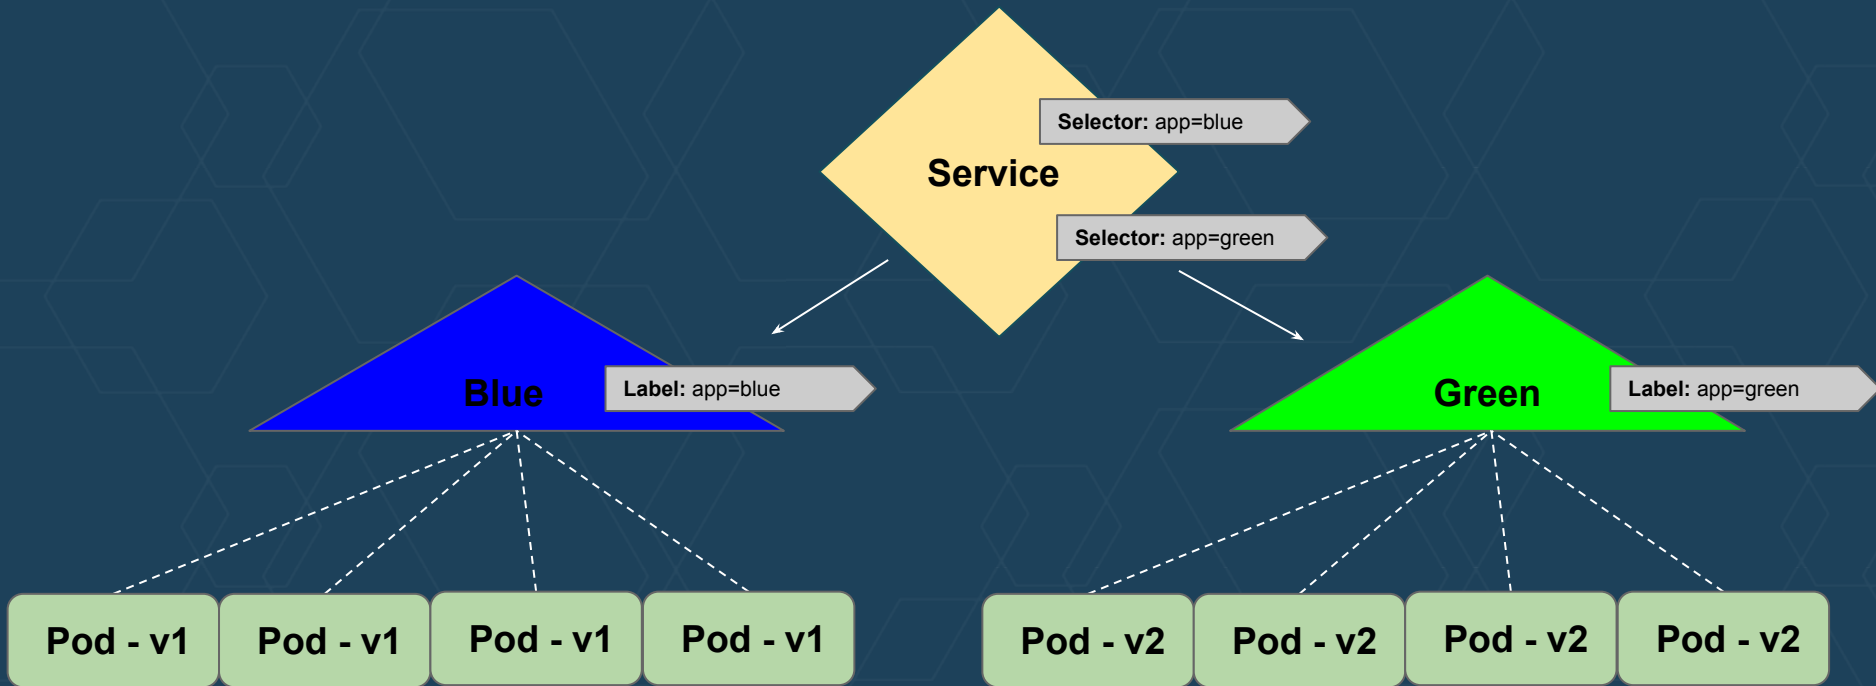
maxUnavailable: 0%

maxSurge: 100%

\*great if you have  
resources to spare.



# Blue-Green Deployments



# Canary Deployments

original deployment:

```
labels:  
  app: guestbook  
  tier: frontend  
  track: stable
```

canary deployment:

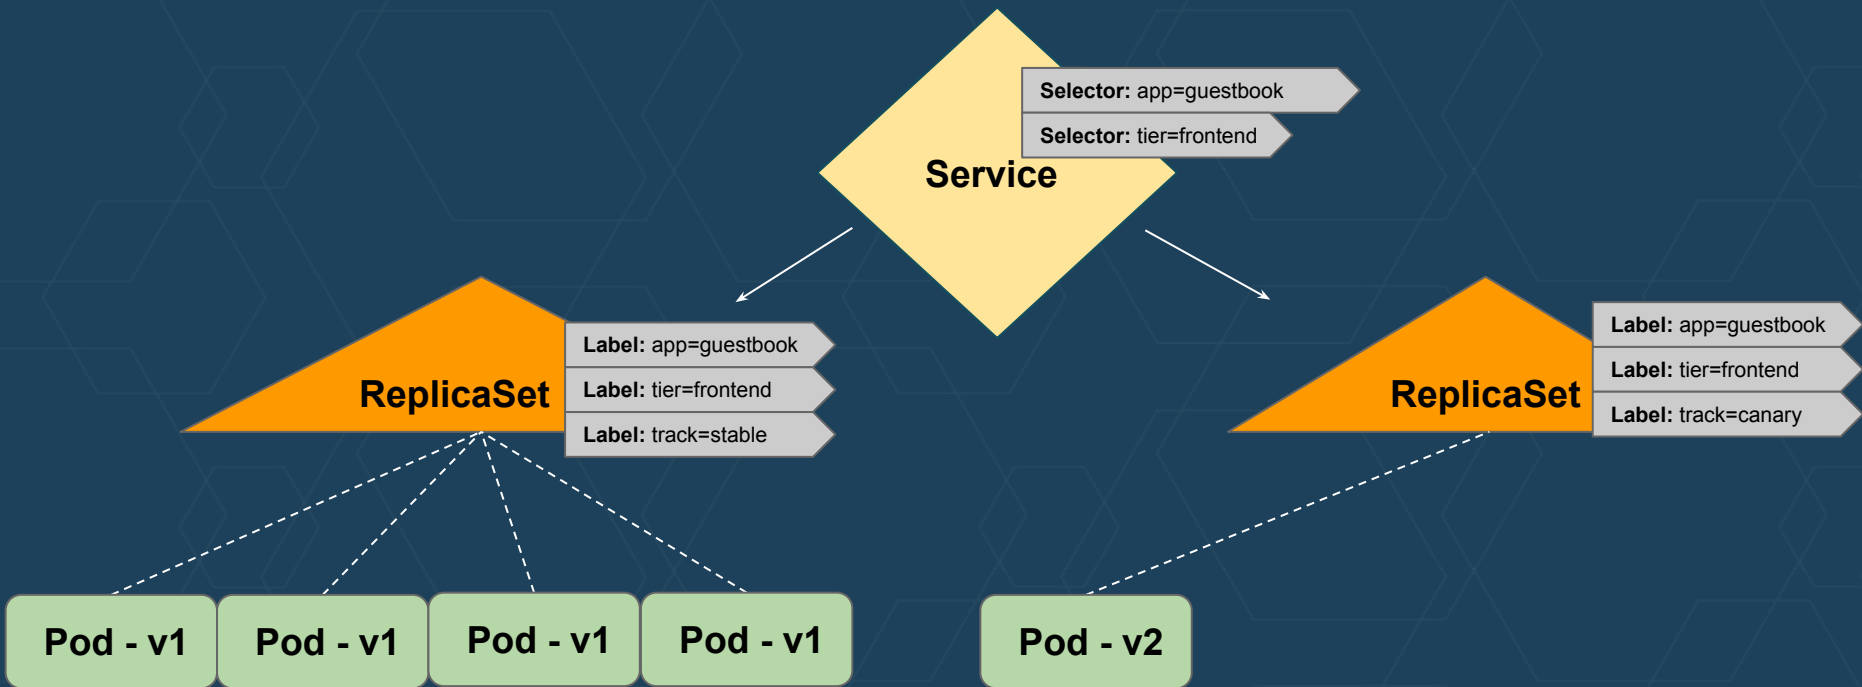
```
labels:  
  app: guestbook  
  tier: frontend  
  track: canary
```

service:

```
selector:  
  app: guestbook  
  tier: frontend
```

- When updating an application, it is common to start with a single "canary" to handle a small fraction of the existing traffic
- If the canary is successful, a full update can proceed
- Create a new deployment for the canary, with unique labels
- The frontend service can span both deployments by using a common label ("guestbook")

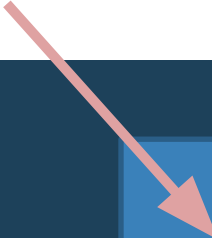
# Canary Deployments



# The Deployment Object

- The parameters of the `kubectl run` command provides the initial contents of a deployment object in the API Server

```
$ kubectl run hello-kube --image=gcr.io/google_containers/echoserver:1.4 --port=8080
```



API Server  
{deployment  
object}

# Desired State and Actual State

- The deployment object contains a desired state object (spec) from the user and the actual state from the deployment controller (status)



# Deployment Example (1/2)

Run three replicas (copies) of pods containing an nginx web server/proxy/load balancer

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
        tier: backend
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
```

Pod spec template defines the 'cookie cutter' used for creating new pods when necessary



# Deployment Example (2/2)

```
$ kubectl create -f nginx-deployment.yaml
deployment "nginx-deployment" created

$ kubectl describe deployments/nginx-deployment
Name:          nginx-deployment
Namespace:     default
CreationTimestamp:  Wed, 24 Aug 2016 13:17:36 -0400
Labels:        app=nginx
                tier=backend
Selector:      app=nginx,tier=backend
Replicas:      3 updated | 3 total | 3 available | 0 unavailable
StrategyType:  RollingUpdate
MinReadySeconds:  0
RollingUpdateStrategy:  1 max unavailable, 1 max surge
OldReplicaSets:  <none>
NewReplicaSet:   nginx-deployment-4253660899 (3/3 replicas created)
Events:
[...]
```

You can also use `'kubectl get deployment nginx-deployment -o json'` to view the JSON

# Deleting Deployments

kubectl delete deployment

```
$ kubectl delete deployment/my-nginx  
deployment "my-nginx" deleted
```

If you try to delete the pods or ReplicaSets before deleting the deployments, Kubernetes will just replace them

# Summary: Deployments, ReplicaSets and Pods



- A Deployment object results in the creation of a ReplicaSet (rs) object, which results in a number of Pod objects

