redhat

ReplicaSets & Deployments

(The Underrated, OG Operators)

redhat

Why do we care about ReplicaSets (formerly ReplicationControllers)?

Redundancy

Multiple running instances means failure can be tolerated.

Scale

Multiple running instances mean more requests can be handled.

Sharding

Multiple running instances can handle different parts of a computation in parallel.

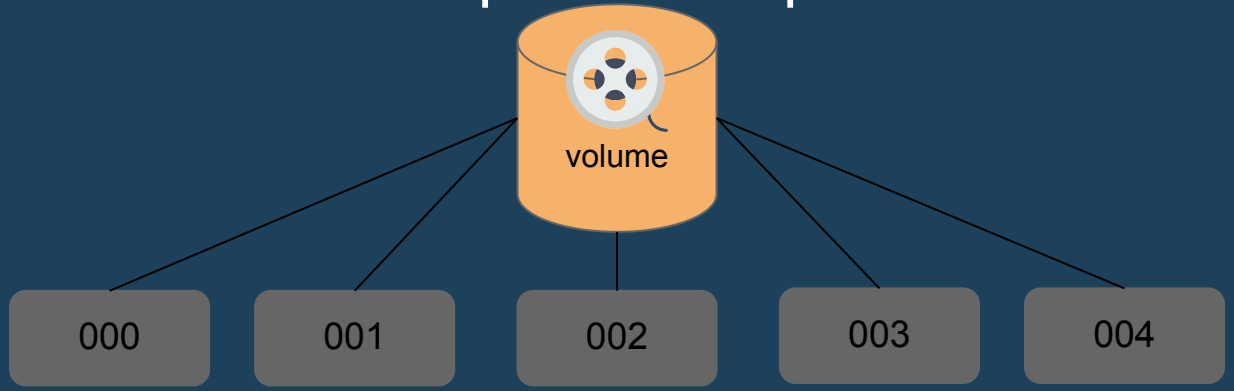

ReplicaSets in Action!

```
kubectl run myfirstapp --image quay.io/coreostrainme/hello-whoami:2.0.1  --restart=Never -l app=myfirstapp,version=1
```

```
kubectl create -f  myfirstreplicaset.yaml
kubectl scale replicaset myfirstreplicaset --replicas=3
```

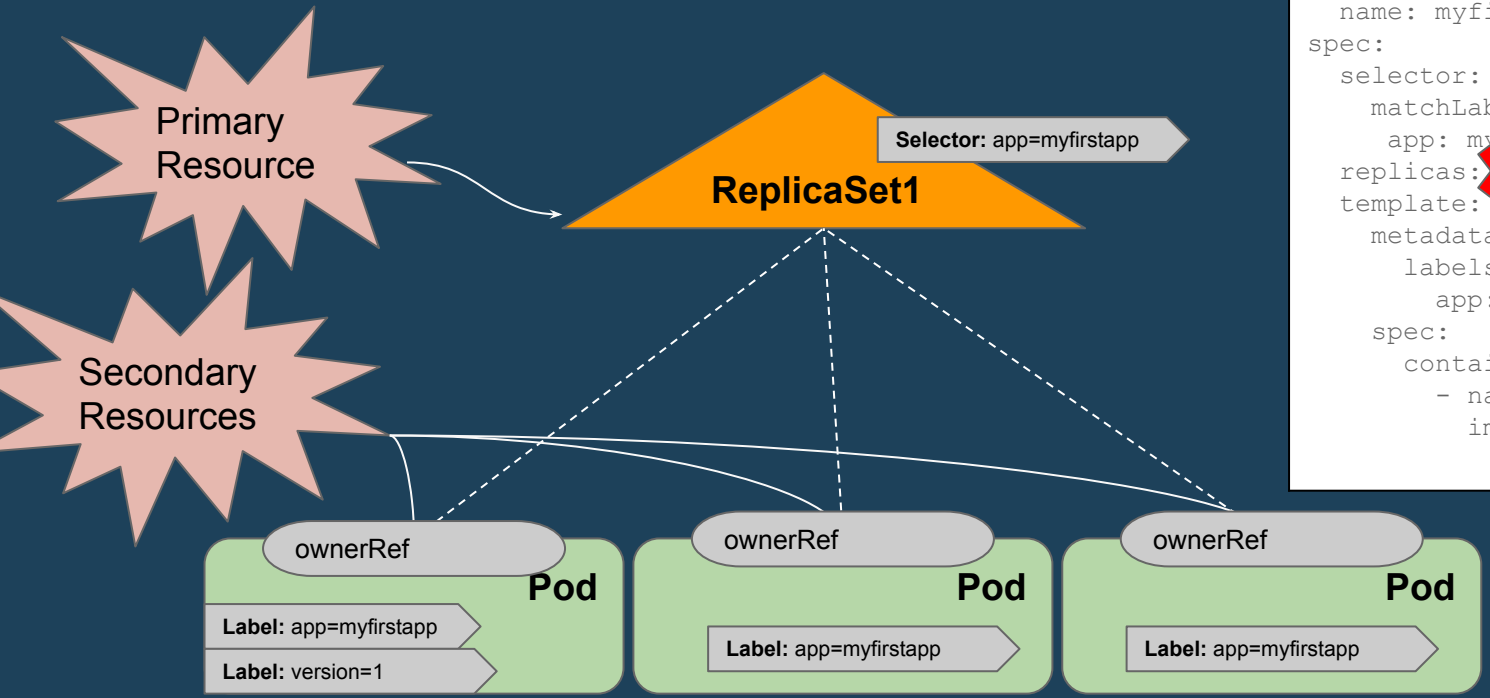


```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: myfirstreplicaset
spec:
  selector:
    matchLabels:
      app: myfirstapp
  replicas: 3
  template:
    metadata:
      labels:
        app: myfirstapp
    spec:
      containers:
        - name: nodejs
          image: myimage
```

redhat.

The ReplicaSet Control Loop

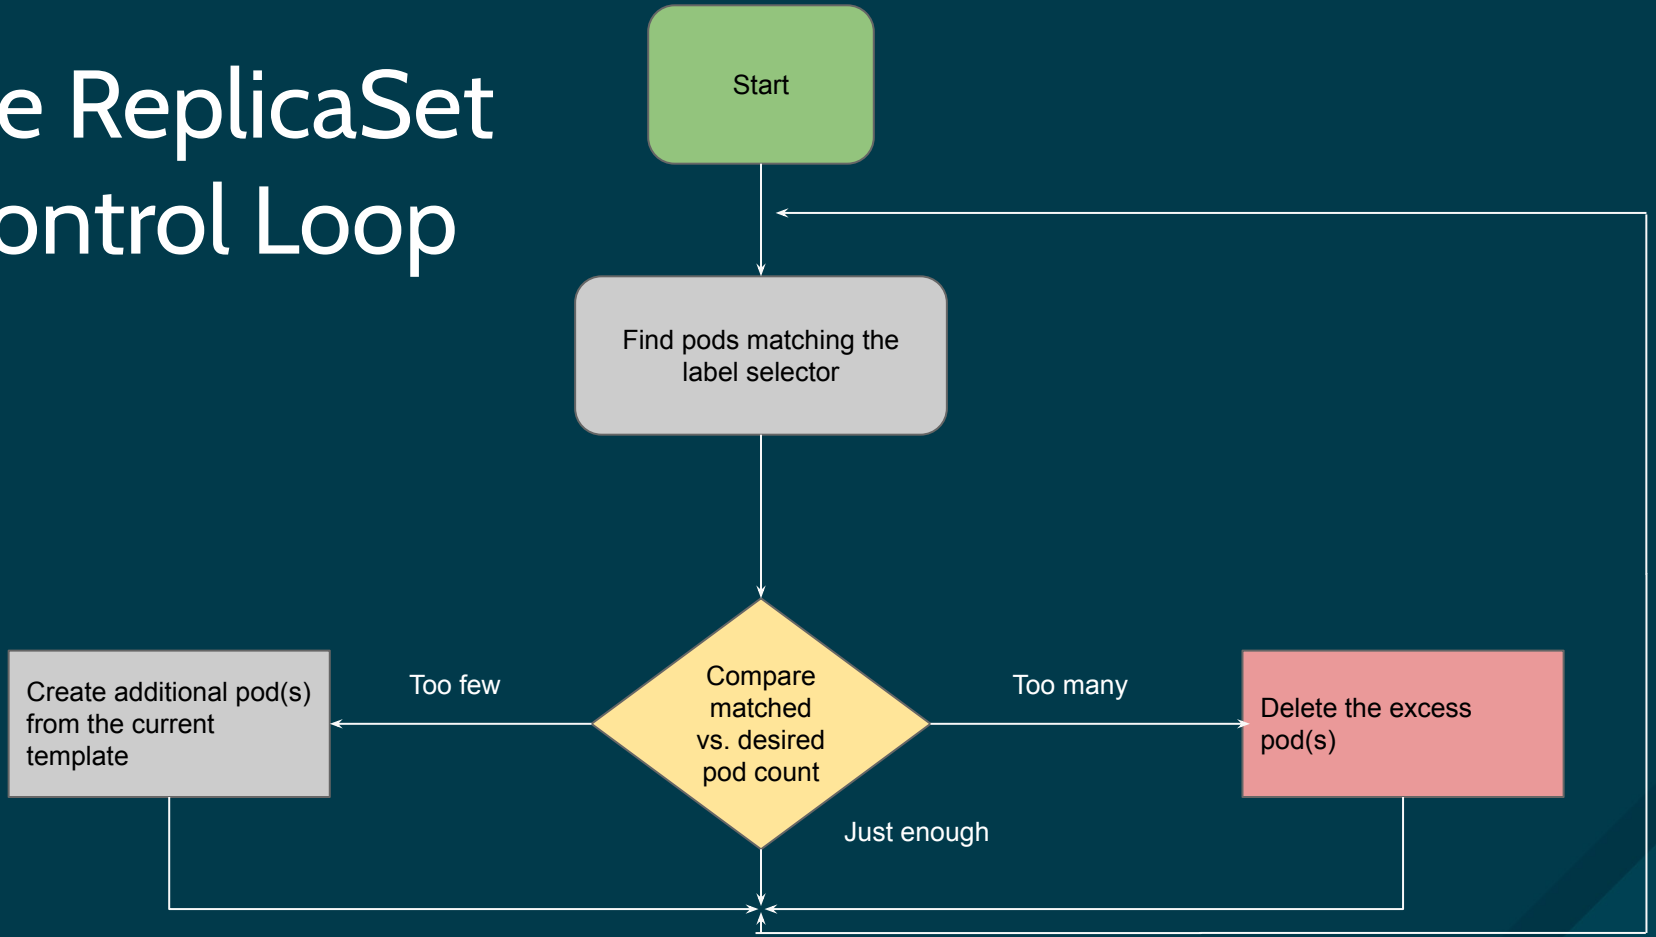

How do we accomplish this?

kube-controller-manager

kube-controller-manager

- Daemon that embeds the core control loops shipped with Kubernetes.
- The control loop is non-terminating loop that regulates the state of the system.
- In Kubernetes, there are multiple control loops running at all times that watch the shared state of the cluster through the apiserver and make changes attempting to move the current state towards the desired state.
- Examples include: ReplicaSets, Endpoints, Deployments, DaemonSets, etc.

Kubernetes Cluster

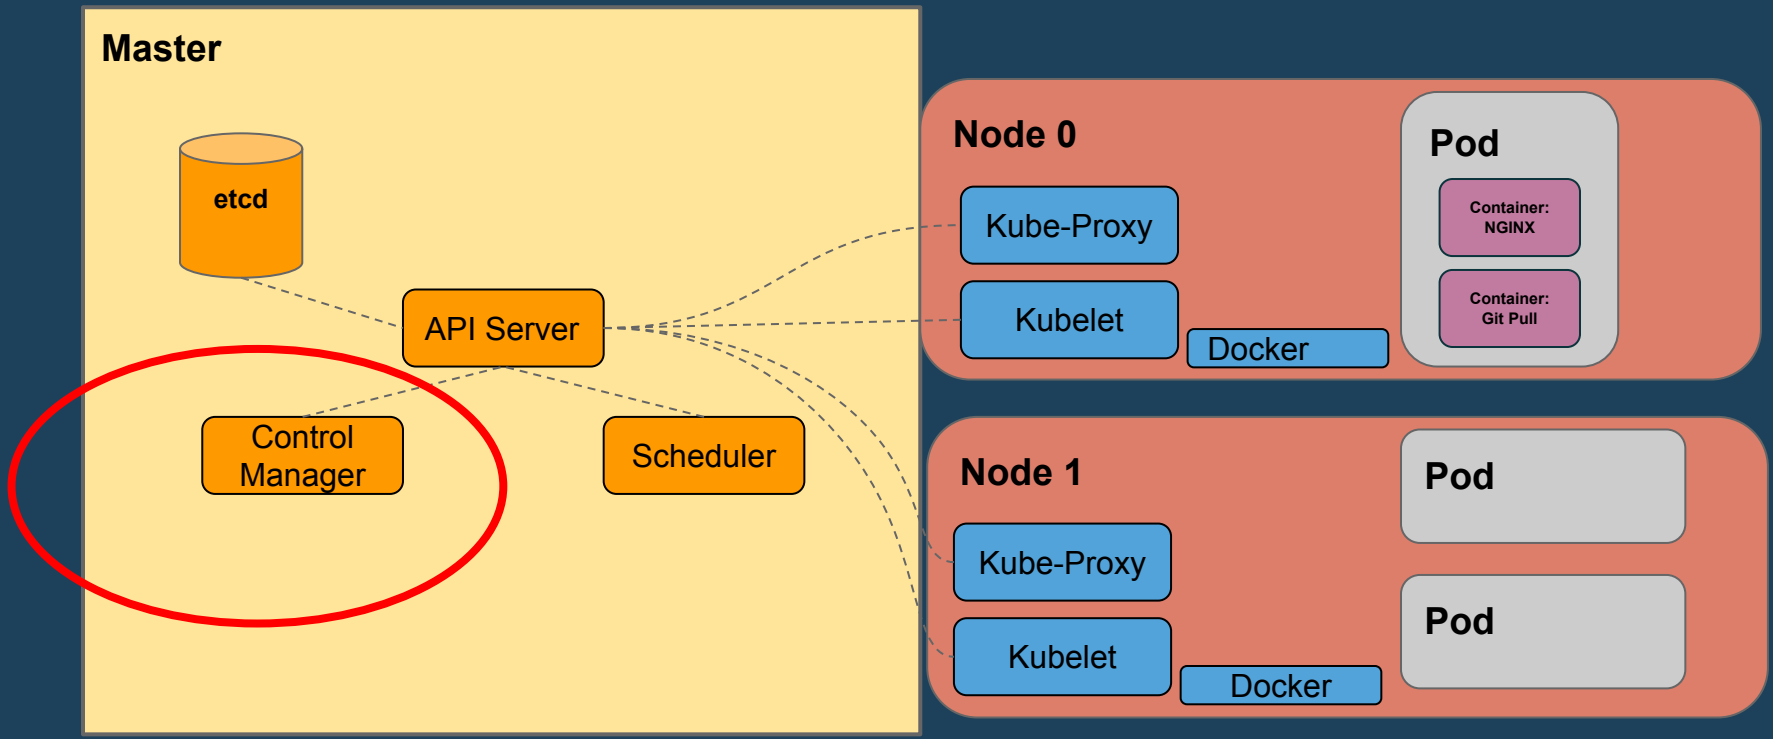

Self-Hosted Kubernetes

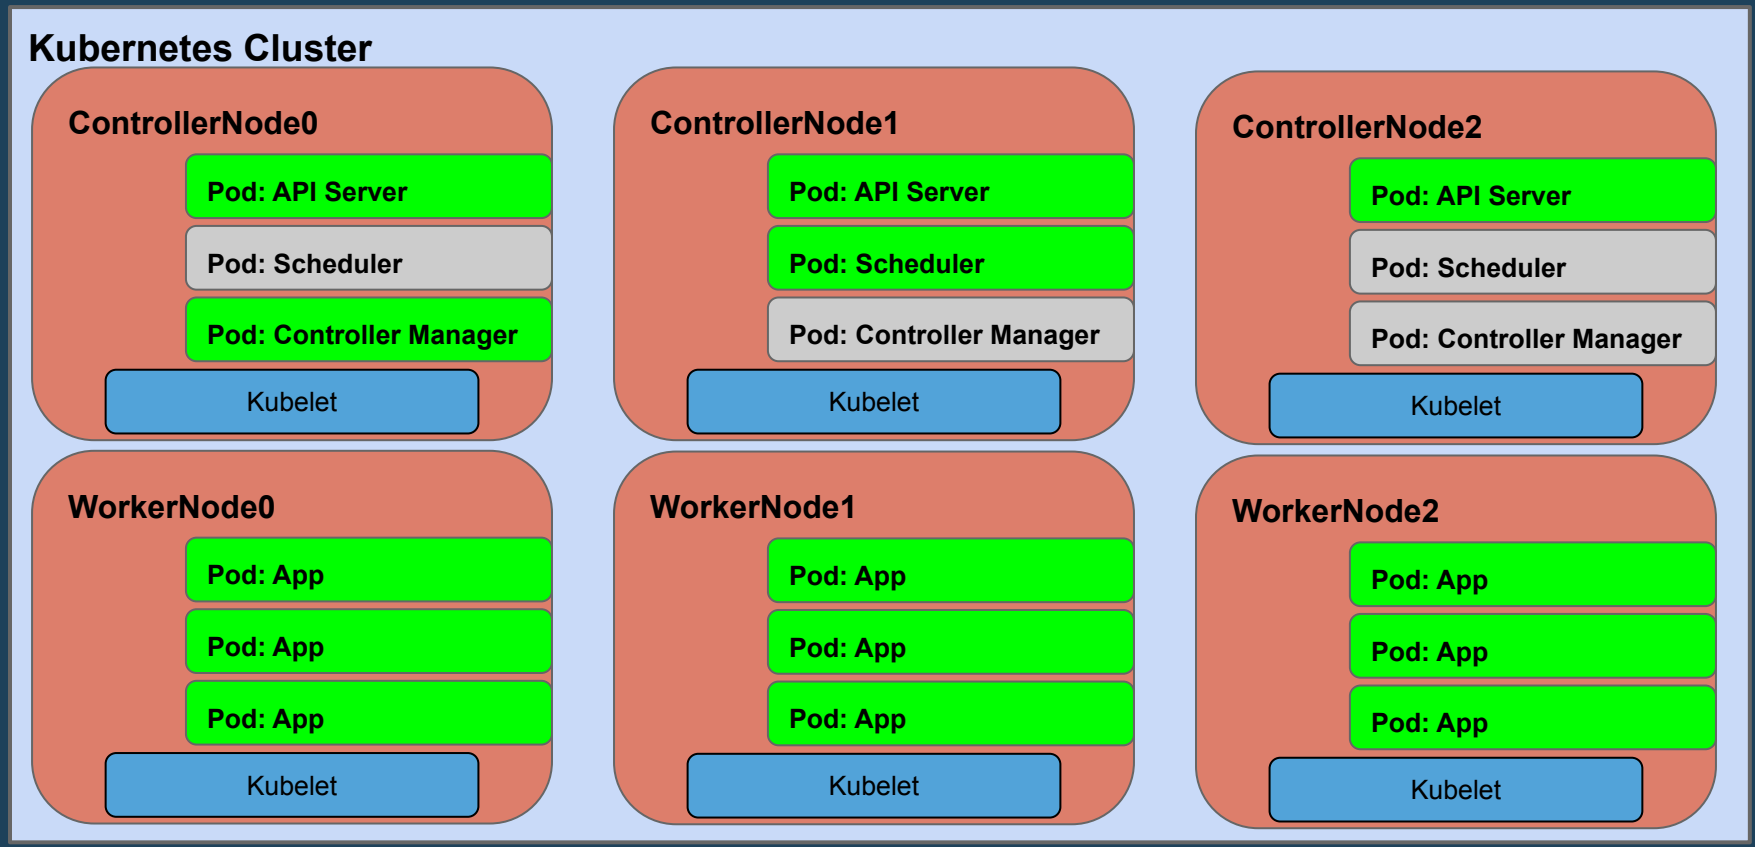

```
oc -n kube-system logs kube-controller-manager-localhost -f
```

Let's head to GitHub.

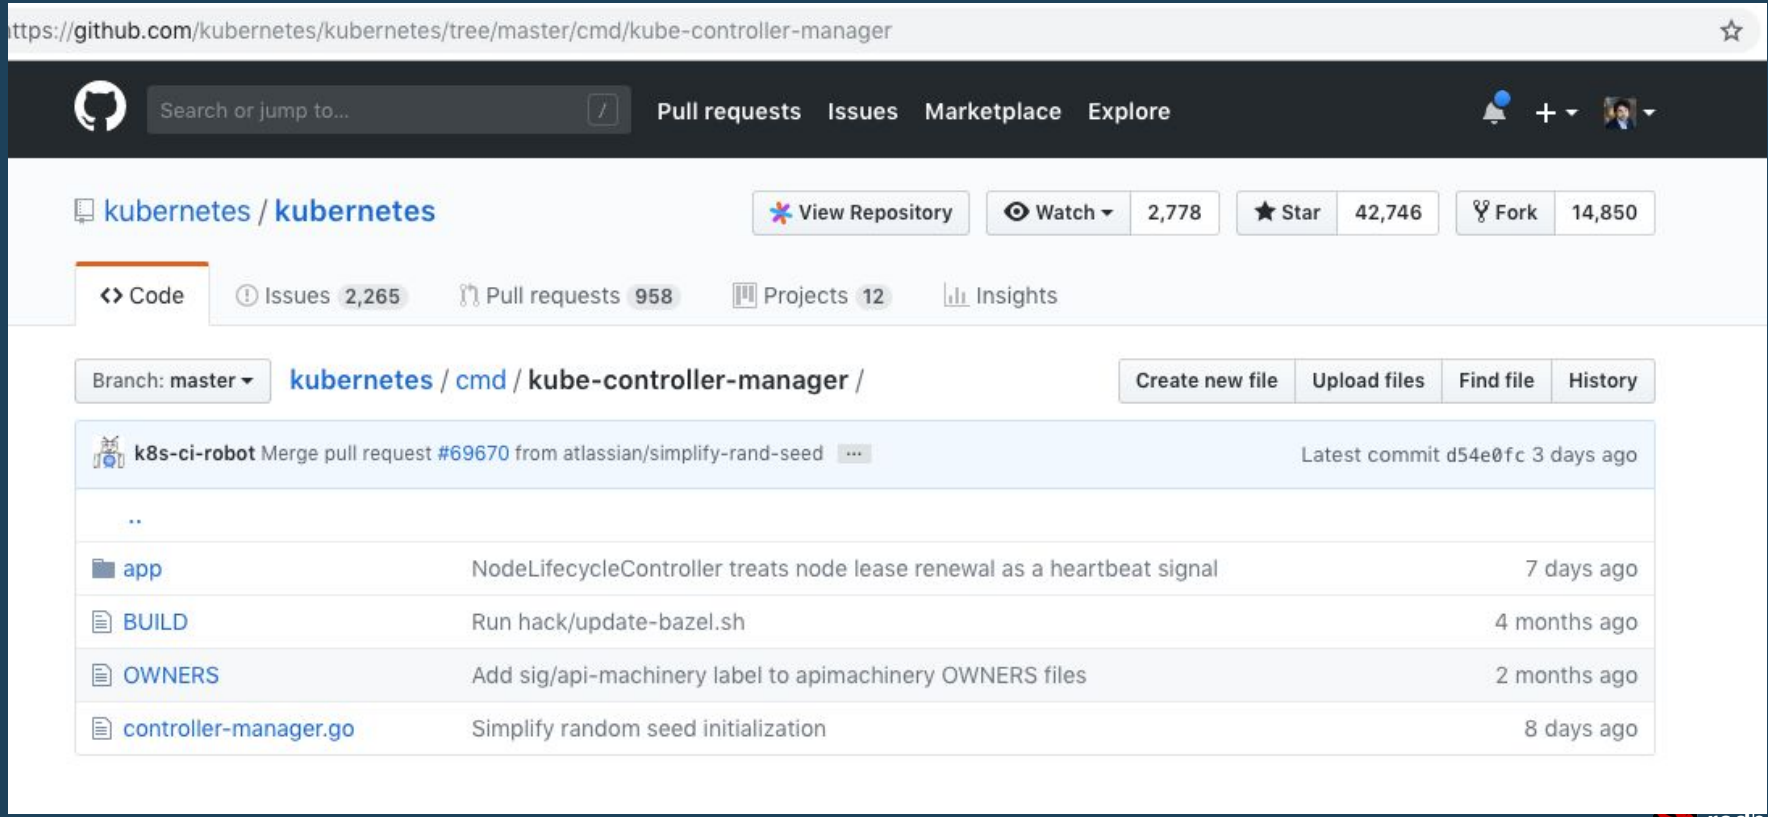

The ReplicaSet Controller Initialization.

https://github.com/kubernetes/kubernetes/blob/master/cmd/kube-controller-manager/app/controllermanager.go

```
342  // NewControllerInitializers is a public map of named controller groups (you can start more than one in an init func)
343  // paired to their InitFunc.  This allows for structured downstream composition and subdivision.
344  func NewControllerInitializers(loopMode ControllerLoopMode) map[string]InitFunc {
345          controllers := map[string]InitFunc{}
346          controllers["endpoint"] = startEndpointController
347          controllers["replicationcontroller"] = startReplicationController
348          controllers["podgc"] = startPodGCController
349          controllers["resourcequota"] = startResourceQuotaController
350          controllers["namespace"] = startNamespaceController
351          controllers["serviceaccount"] = startServiceAccountController
352          controllers["garbagecollector"] = startGarbageCollectorController
353          controllers["daemonset"] = startDaemonSetController
354          controllers["job"] = startJobController
355
356          controllers["replicaset"] = startReplicaSetController
357          controllers["horizontalpodautoscaling"] = startHPAController
358          controllers["disruption"] = startDisruptionController
359          controllers["statefulset"] = startStatefulSetController
```

startReplicaSetController

https://github.com/kubernetes/kubernetes/blob/master/cmd/kube-controller-manager/app/apps.go

```
59          go statefulset.NewStatefulSetController(
60                  ctx.InformerFactory.Core().V1().Pods(),
61                  ctx.InformerFactory.Apps().V1().StatefulSets(),
62                  ctx.InformerFactory.Core().V1().PersistentVolumeClaims(),
63                  ctx.InformerFactory.Apps().V1().ControllerRevisions(),
64                  ctx.ClientBuilder.ClientOrDie("statefulset-controller"),
65          ).Run(1, ctx.Stop)
66          return nil, true, nil
67  }
68
69  func startReplicaSetController(ctx ControllerContext) (http.Handler, bool, error) {
70          if !ctx.AvailableResources[schema.GroupVersionResource{Group: "apps", Version: "v1", Resource: "replicasets"}] {
71                  return nil, false, nil
72          }
73          go replicaset.NewReplicaSetController(
74                  ctx.InformerFactory.Apps().V1().ReplicaSets(),
75                  ctx.InformerFactory.Core().V1().Pods(),
76                  ctx.ClientBuilder.ClientOrDie("replicaset-controller"),
77                  replicaset.BurstReplicas,
78          ).Run(int(ctx.ComponentConfig.ReplicaSetController.ConcurrentRSSyncs), ctx.Stop)
79          return nil, true, nil
80  }
```

newReplicaSetController

https://github.com/kubernetes/kubernetes/blob/master/pkg/controller/replicaset/replica_set.go

```
108  // NewReplicaSetController configures a replica set controller with the specified event recorder
109  func NewReplicaSetController(rsInformer appsinformers.ReplicaSetInformer, podInformer coreinformers.PodInformer, kubeClient cl
110          eventBroadcaster := record.NewBroadcaster()
111          eventBroadcaster.StartLogging(glog.Infof)
112          eventBroadcaster.StartRecordingToSink(&v1core.EventSinkImpl{Interface: kubeClient.CoreV1().Events("")})
113          return NewBaseController(rsInformer, podInformer, kubeClient, burstReplicas,
114              apps.SchemeGroupVersion.WithKind("ReplicaSet"),
115              "replicaset_controller",
116              "replicaset",
117              controller.RealPodControl{
118                      KubeClient: kubeClient,
119                      Recorder:   eventBroadcaster.NewRecorder(scheme.Scheme, v1.EventSource{Component: "replicaset-controlle
120              },
121          )
122  }
```

1

2

newBaseController

https://github.com/kubernetes/kubernetes/blob/master/pkg/controller/replicaset/replica_set.go#L255:34

```
126  func NewBaseController(rsInformer appsinformers.ReplicaSetInformer, podInformer coreinformers.PodInformer, kubeClient clientse
127          gvk schema.GroupVersionKind, metricOwnerName, queueName string, podControl controller.PodControlInterface) *ReplicaSetC
128          if kubeClient != nil && kubeClient.CoreV1().RESTClient().GetRateLimiter() != nil {
129                  metrics.RegisterMetricAndTrackRateLimiterUsage(metricOwnerName, kubeClient.CoreV1().RESTClient().GetRateLimiter
130          }
131
132          rsc := &ReplicaSetController{
133                  GroupVersionKind: gvk,
134                  kubeClient:       kubeClient,
135                  podControl:       podControl,
136                  burstReplicas:    burstReplicas,
137                  expectations:     controller.NewUIDTrackingControllerExpectations(controller.NewControllerExpectations()),
138                  queue:            workqueue.NewNamedRateLimitingQueue(workqueue.DefaultControllerRateLimiter(), queueName),
139          }
```

1

newBaseController

```
141         rsInformer.Informer().AddEventHandler(cache.ResourceEventHandlerFuncs{
142                 AddFunc:    rsc.enqueueReplicaSet,
143                 UpdateFunc: rsc.updateRS,
144                 // This will enter the sync loop and no-op, because the replica set has been deleted from the store.
145                 // Note that deleting a replica set immediately after scaling it to 0 will not work. The recommended
146                 // way of achieving this is by performing a `stop` operation on the replica set.
147                 DeleteFunc: rsc.enqueueReplicaSet,
148         })
149         rsc.rsLister = rsInformer.Lister()
150         rsc.rsListerSynced = rsInformer.Informer().HasSynced
```

2

Watching Primary Resource: kind:ReplicaSet

newBaseController

```
152         podInformer.Informer().AddEventHandler(cache.ResourceEventHandlerFuncs{
153                 AddFunc: rsc.addPod,
154                 // This invokes the ReplicaSet for every pod change, eg: host assignment. Though this might seem like
155                 // overkill the most frequent pod update is status, and the associated ReplicaSet will only list from
156                 // local storage, so it should be ok.
157                 UpdateFunc: rsc.updatePod,
158                 DeleteFunc: rsc.deletePod,
159         })
```

3

Watching Secondary Resource: kind:Pod

Creating another ReplicaSet

```
kubectl create -f  mysecondreplicaset.yaml
```

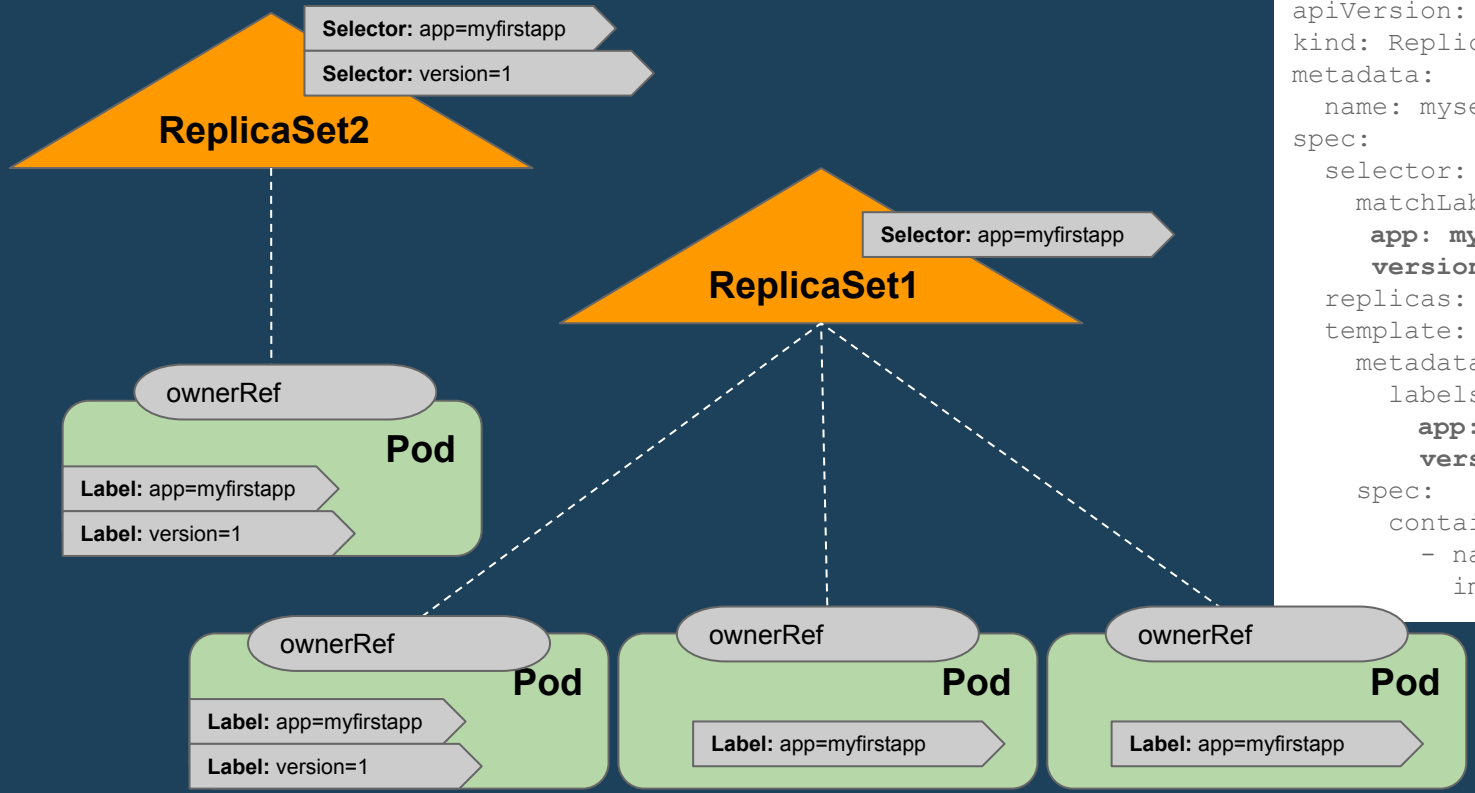


```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: mysecondreplicaset
spec:
  selector:
    matchLabels:
     app: myfirstapp
     version: 1
  replicas: 1
  template:
    metadata:
      labels:
        app: myfirstapp
        version: 1
    spec:
      containers:
        - name: nodejs
          image: myimage
```

redhat.

Creating Another Orphan Pod

```
kubectl run myfirstapp --image quay.io/coreostrainme/hello-whoami:2.0.1  --restart=Never -l app=myfirstapp,version=1
```

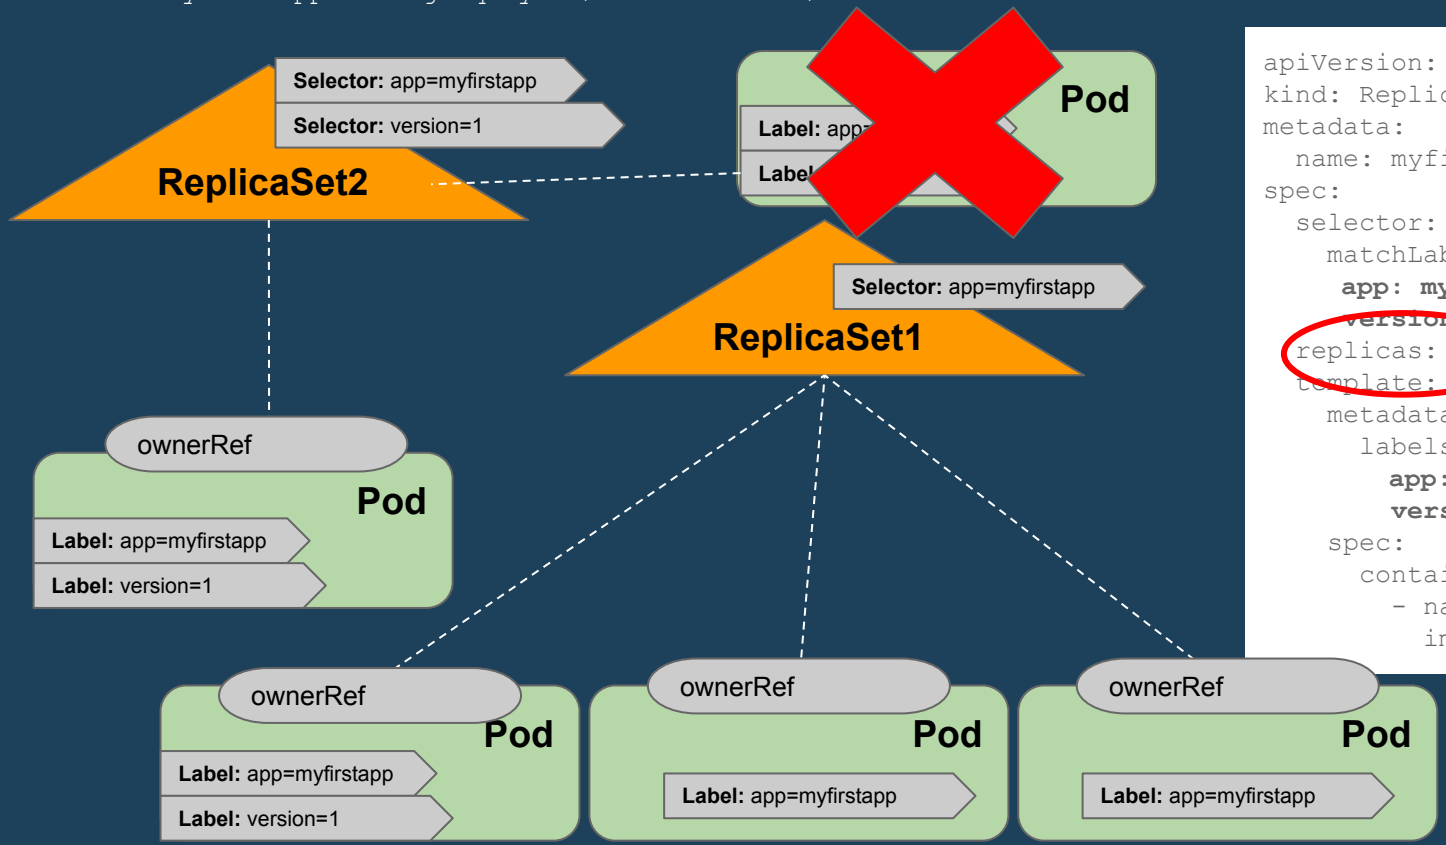


```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: myfirstreplicaset
spec:
  selector:
    matchLabels:
      app: myfirstapp
      version: 1
  replicas: 1
  template:
    metadata:
      labels:
        app: myfirstapp
        version: 1
    spec:
      containers:
        - name: nodejs
          image: myimage
```

redhat.

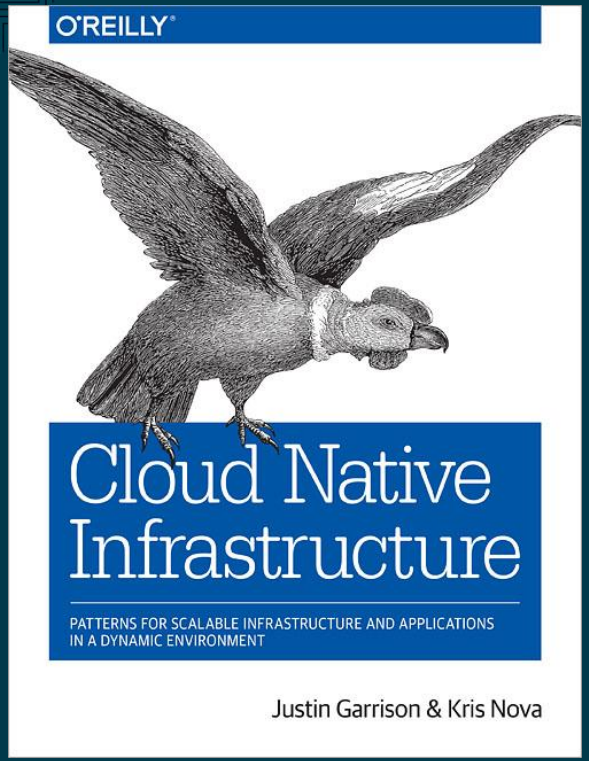

Chapter 4
Designing Infrastructure Applications

*The **reconciler pattern** is a software pattern that can be used or expanded upon for managing cloud native infrastructure. The pattern enforces the idea of having two representations of the infrastructure—the first being the actual state of the infrastructure, and the second being the expected state of the infrastructure.*

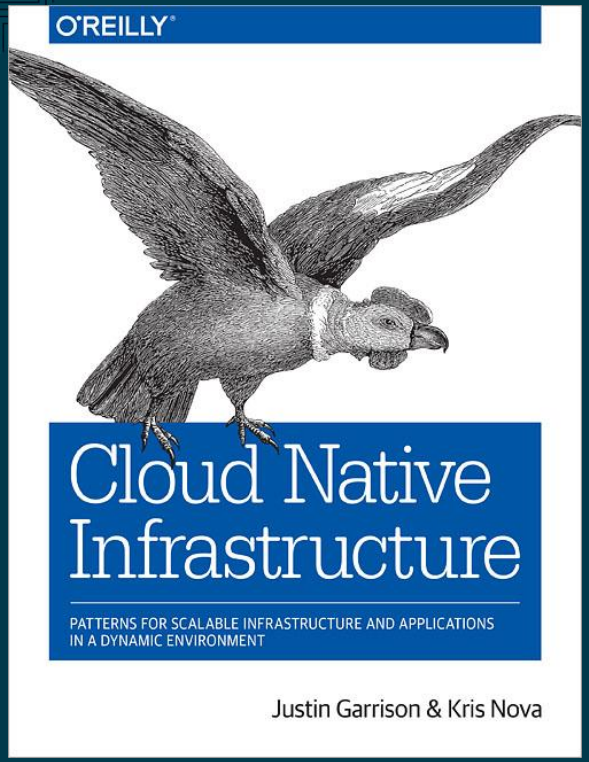

*The **reconciler pattern** will force the engineer to have two independent avenues for getting either of these representations, as well as to implement a solution to reconcile the actual state into the expected state.*

ReplicaSets in Action!

```
kubectl create -f  myfirstreplicaset.yaml
```

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: myfirstreplicaset
spec:
  selector:
    matchLabels:
     app: myfirstapp
  replicas: 3
  template:
    metadata:
      labels:
        app: myfirstapp
    spec:
      containers:
        - name: nodejs
          image: myimage
status:
  availableReplicas: 0
  fullyLabeledReplicas: 3
  observedGeneration: 1
  readyReplicas: 0
  replicas: 3
status:
  availableReplicas: 3
  fullyLabeledReplicas: 3
  observedGeneration: 1
  readyReplicas: 3
  replicas: 3
```

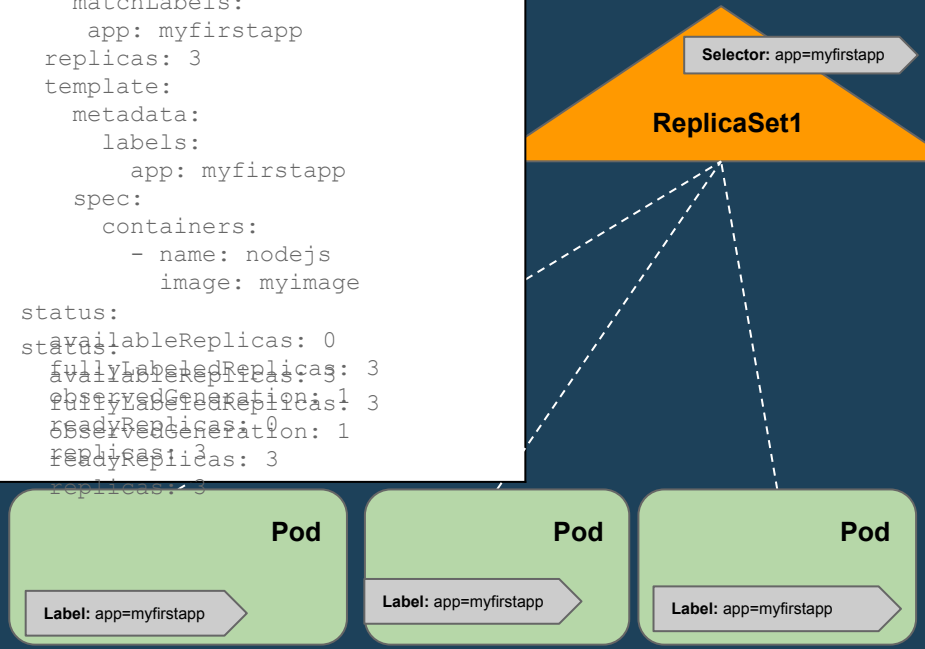

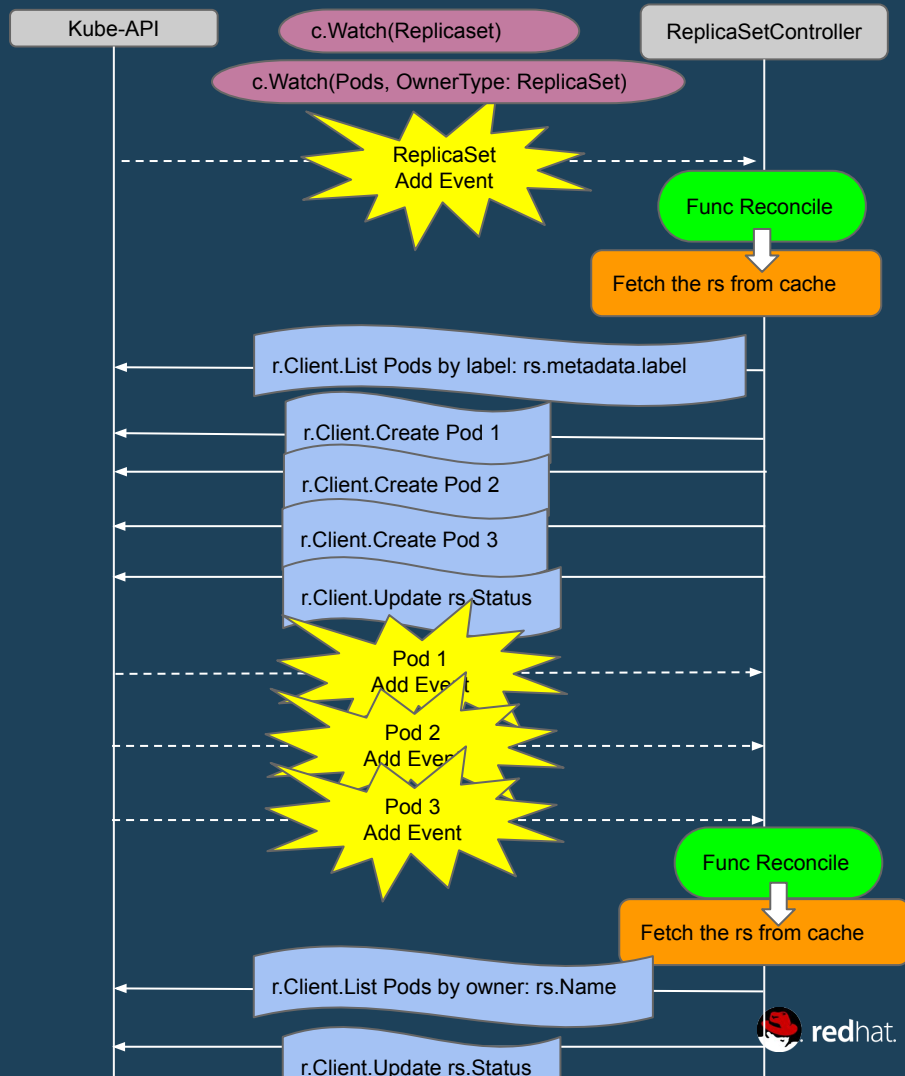

redhat.

ReplicaSets in Action!

```
kubectl create -f  myfirstreplicaset.yaml
```

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: myfirstreplicaset
spec:
  selector:
    matchLabels:
     app: myfirstapp
  replicas: 3
  template:
    metadata:
      labels:
        app: myfirstapp
    spec:
      containers:
        - name: nodejs
          image: myimage

status:
  availableReplicas: 3
  fullyLabeledReplicas: 3
  observedGeneration: 1
  readyReplicas: 3
  replicas: 3
```

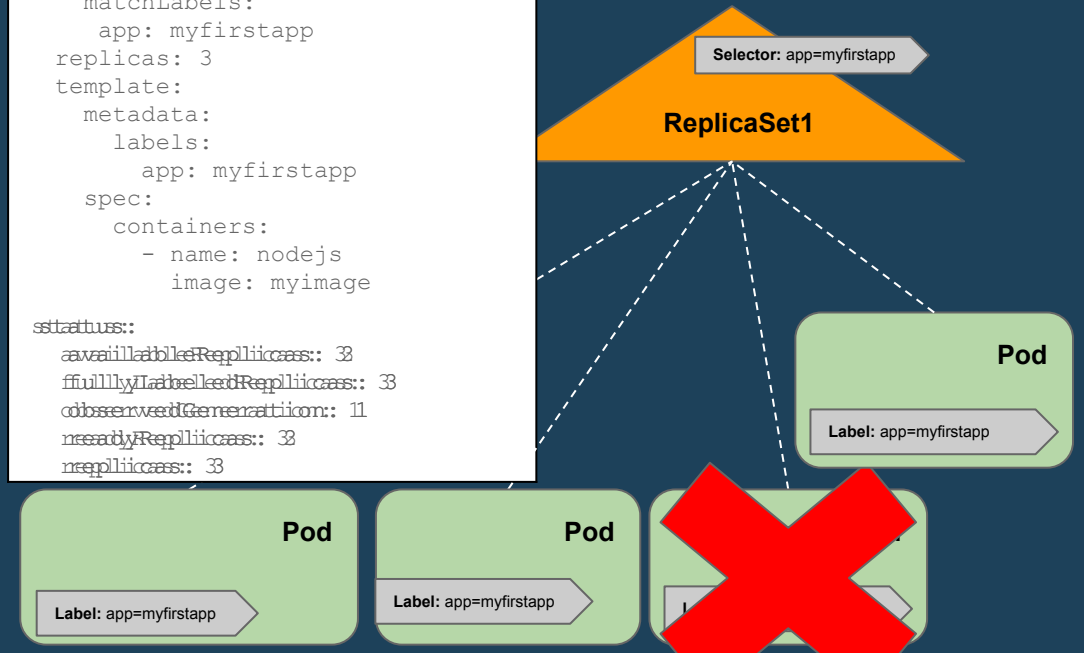

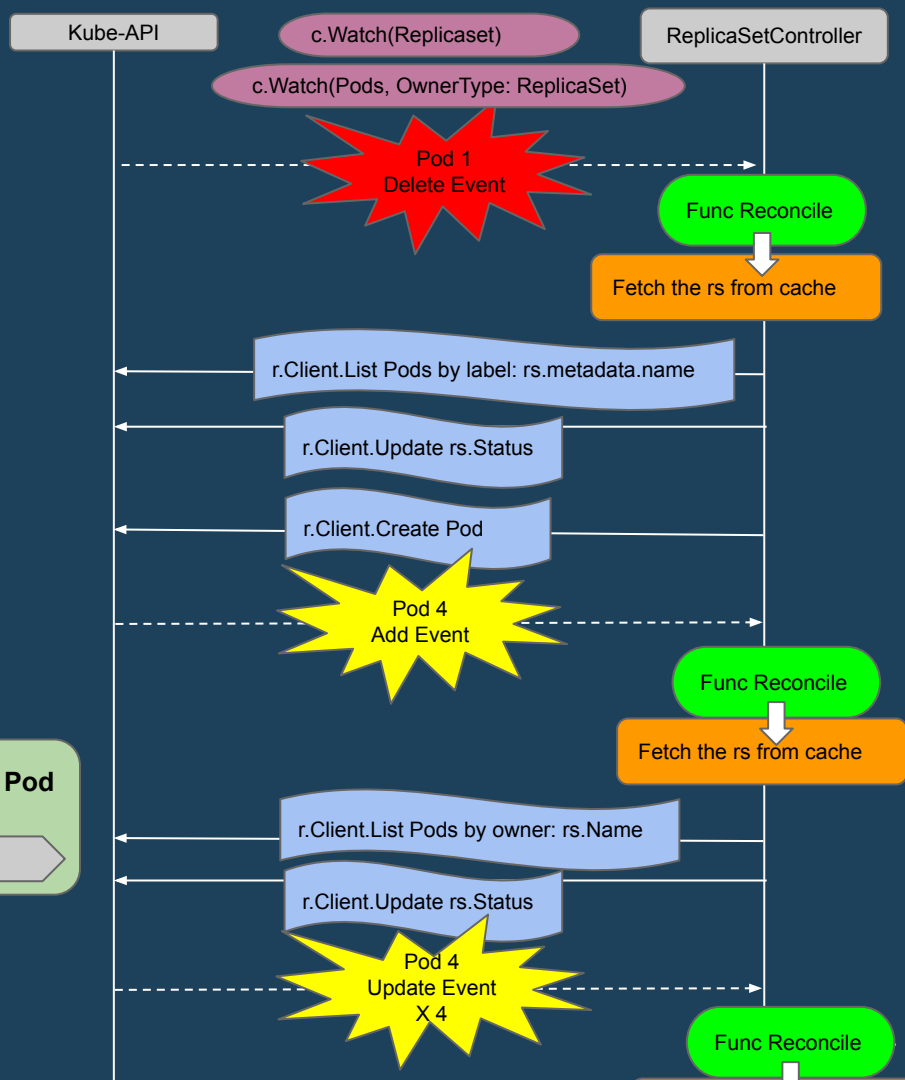

Let's Identify Primary/Secondary Resources for Existing Kubernetes Controllers!

(without looking at the code!)

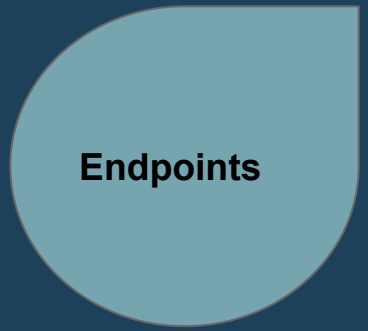

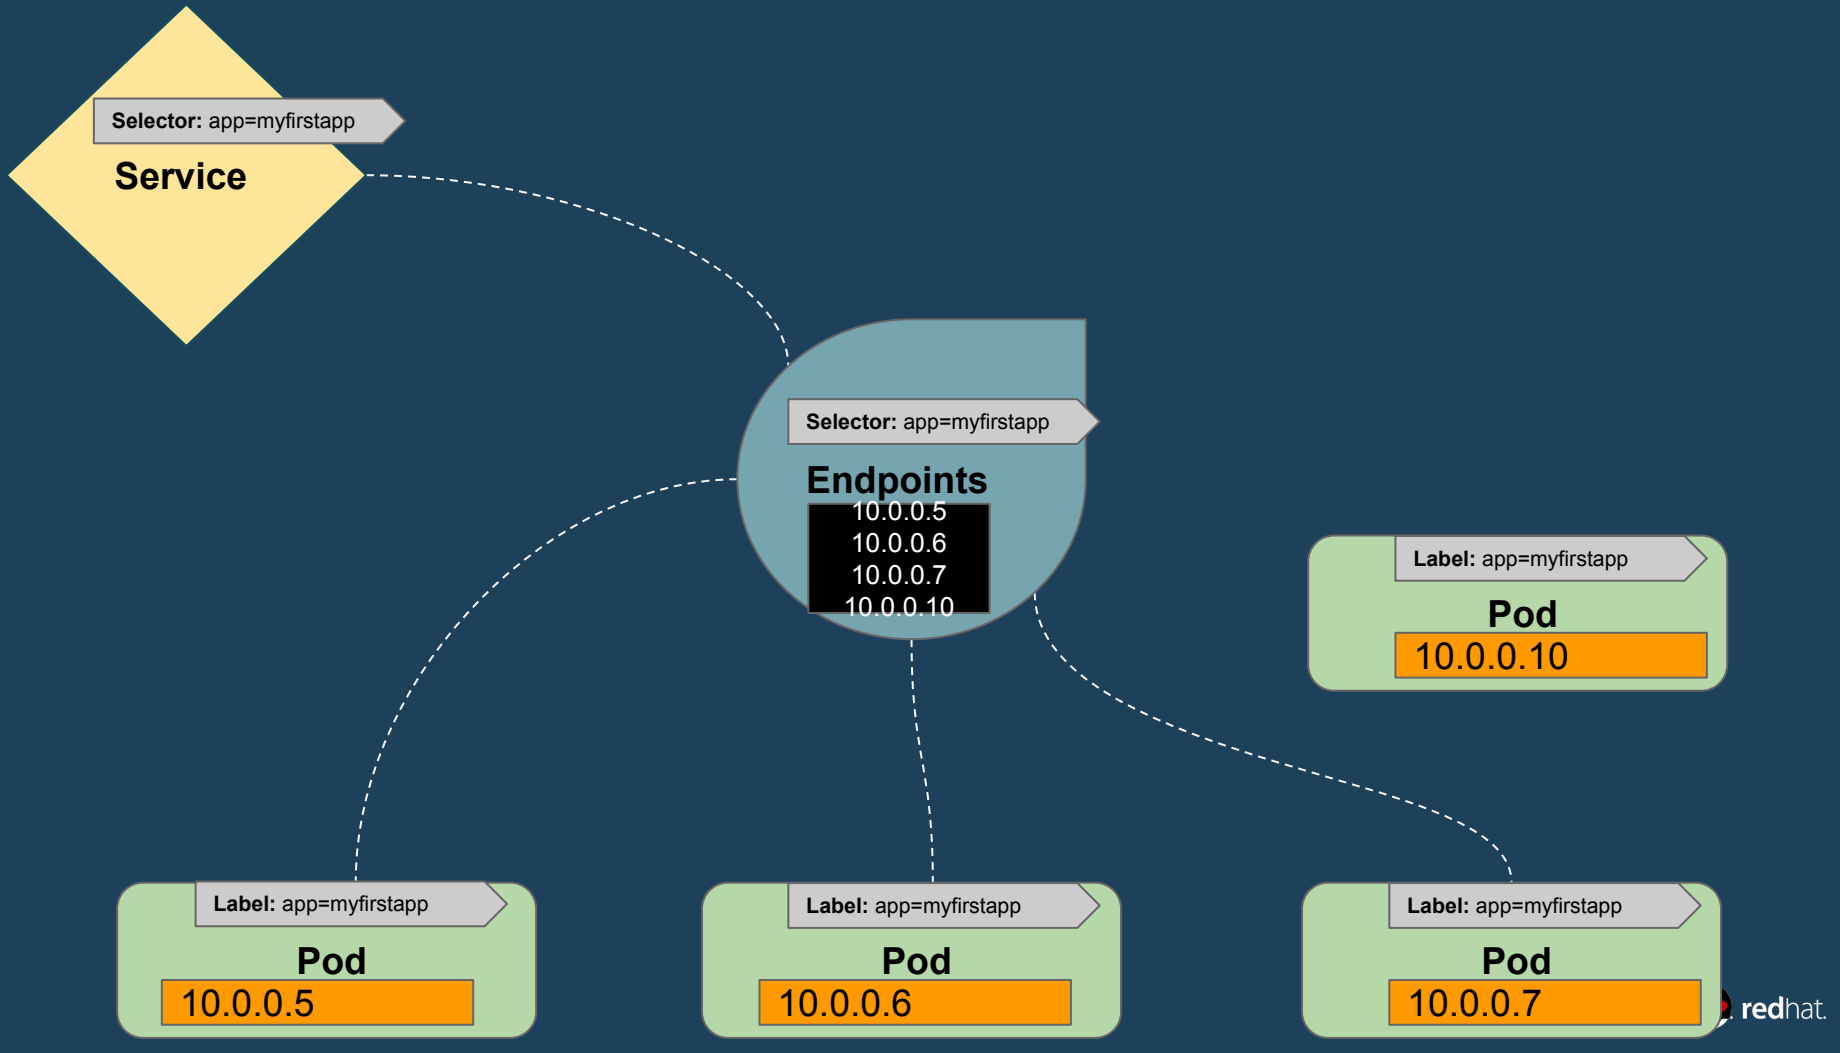
Selector: app=myfirstapp
Service
Selector: app=myfirstapp
Endpoints
10.0.0.5
10.0.0.6
10.0.0.7
10.0.0.10
Label: app=myfirstapp
Pod
10.0.0.10
Label: app=myfirstapp
Pod
10.0.0.5
Label: app=myfirstapp
Pod
10.0.0.6
Label: app=myfirstapp
Pod
10.0.0.7
redhat.

EndPoint Controller

```
71  // NewEndpointController returns a new *EndpointController.
72  func NewEndpointController(podInformer coreinformers.PodInformer, serviceInformer coreinformers.ServiceInformer,
73          endpointsInformer coreinformers.EndpointsInformer, client clientset.Interface) *EndpointController {
74          if client != nil && client.CoreV1().RESTClient().GetRateLimiter() != nil {
75                  metrics.RegisterMetricAndTrackRateLimiterUsage("endpoint_controller", client.CoreV1().RESTClient().GetRateLimit
76          }
77          e := &EndpointController{
78                  client:           client,
79                  queue:            workqueue.NewNamedRateLimitingQueue(workqueue.DefaultControllerRateLimiter(), "endpoint"),
80                  workerLoopPeriod: time.Second,
81          }
82
83          serviceInformer.Informer().AddEventHandler(cache.ResourceEventHandlerFuncs{
84                  AddFunc: e.enqueueService,
85                  UpdateFunc: func(old, cur interface{}) {
86                          e.enqueueService(cur)
87                  },
88                  DeleteFunc: e.enqueueService,
89          })
90          e.serviceLister = serviceInformer.Lister()
91          e.servicesSynced = serviceInformer.Informer().HasSynced
92
93          podInformer.Informer().AddEventHandler(cache.ResourceEventHandlerFuncs{
94                  AddFunc:    e.addPod,
95                  UpdateFunc: e.updatePod,
96                  DeleteFunc: e.deletePod,
97          })
98          e.podLister = podInformer.Lister()
99          e.podsSynced = podInformer.Informer().HasSynced
100
101         e.endpointsLister = endpointsInformer.Lister()
102         e.endpointsSynced = endpointsInformer.Informer().HasSynced
103
104         return e
105 }
```

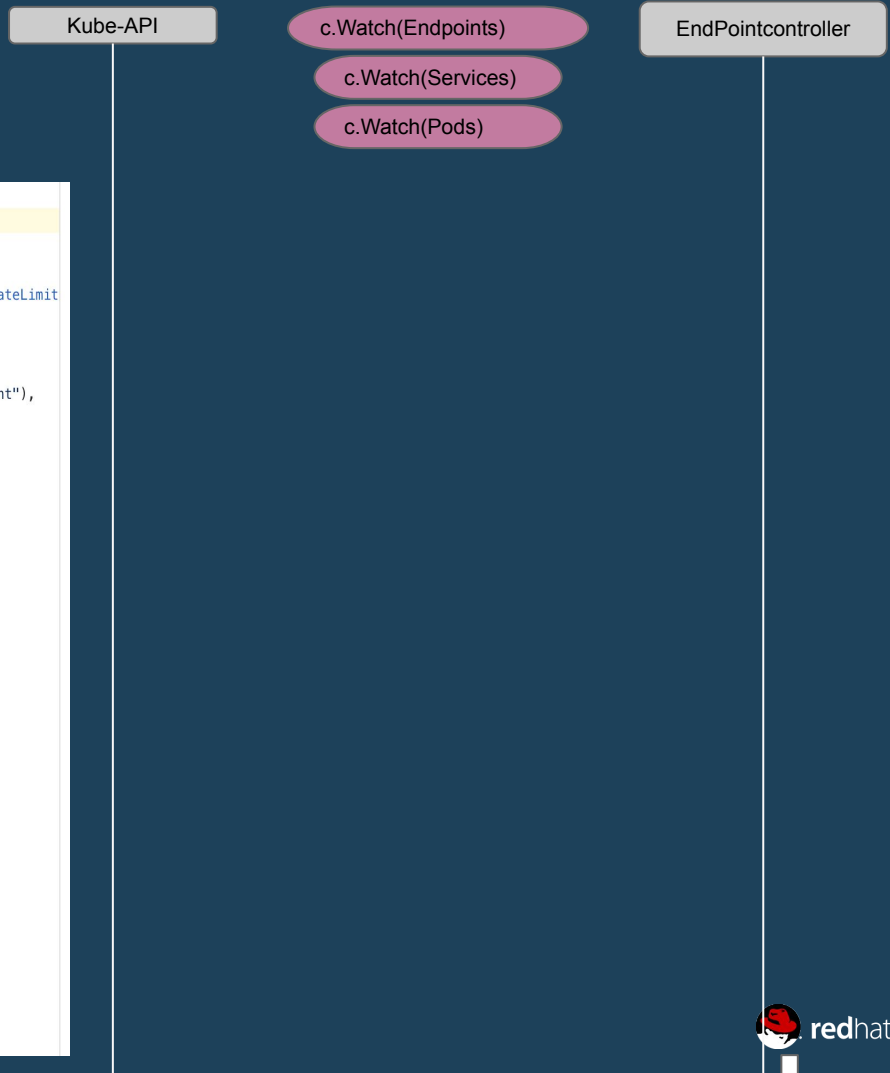

redhat.

DaemonSets

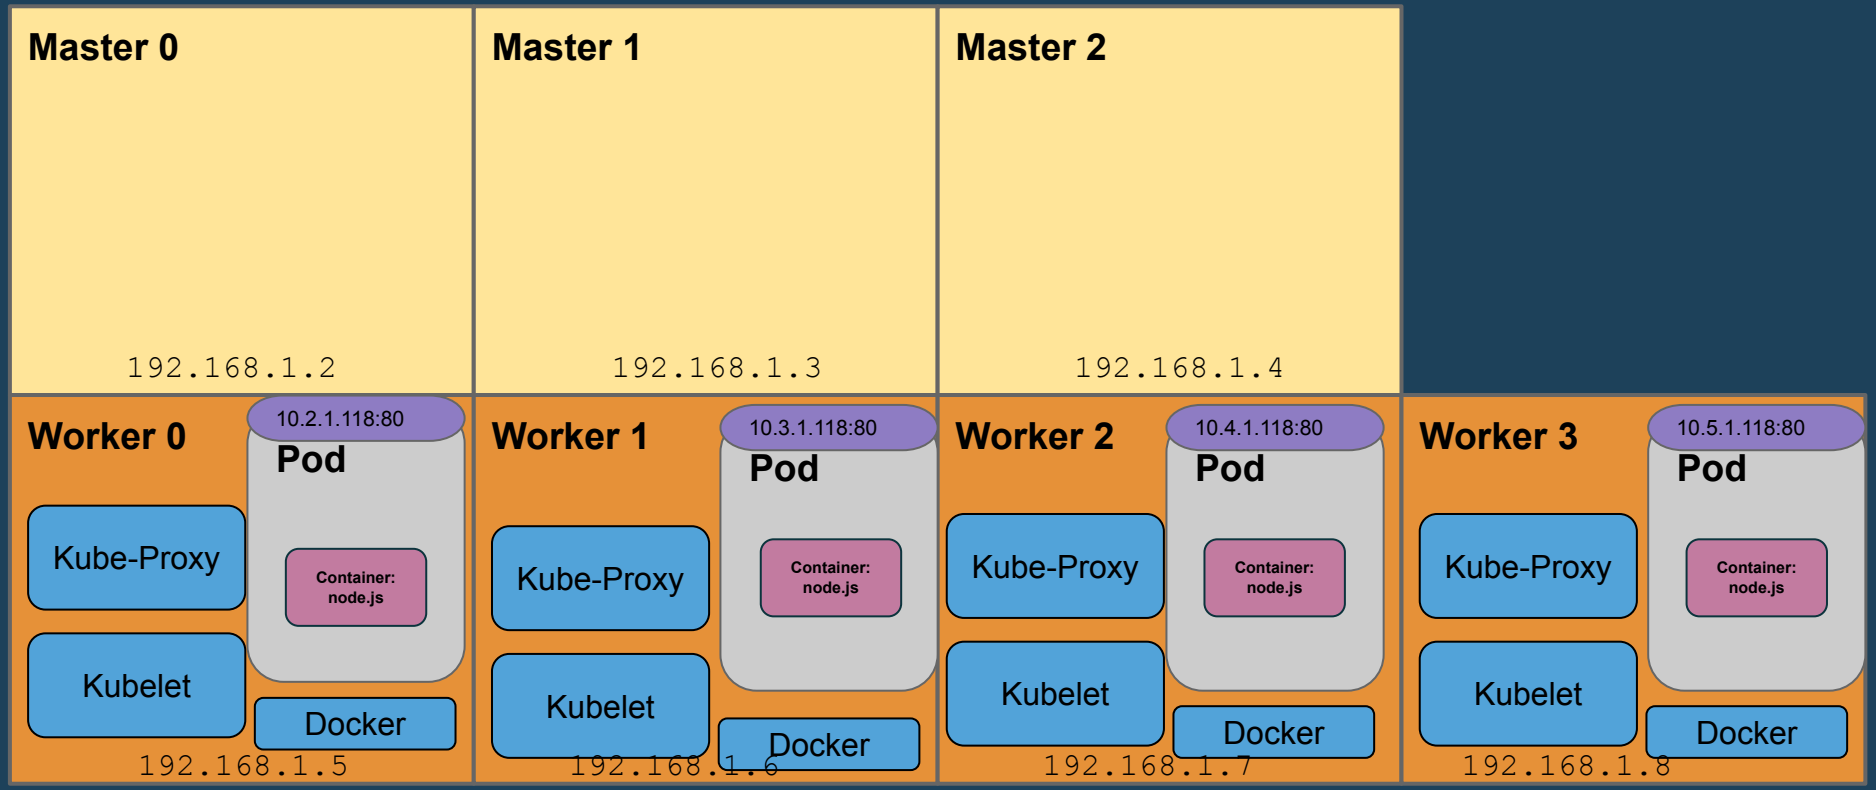

DaemonSetController

```
143 func NewDaemonSetsController(
144         daemonSetInformer appsinformers.DaemonSetInformer,
145         historyInformer appsinformers.ControllerRevisionInformer,
146         podInformer coreinformers.PodInformer,
147         nodeInformer coreinformers.NodeInformer,
148         kubeClient clientset.Interface,
149         failedPodsBackoff *flowcontrol.Backoff,
```

Worker 0	Worker 1	Worker 2	Worker 3
Pod	Pod	Pod	Pod

Administrator adds a new node!

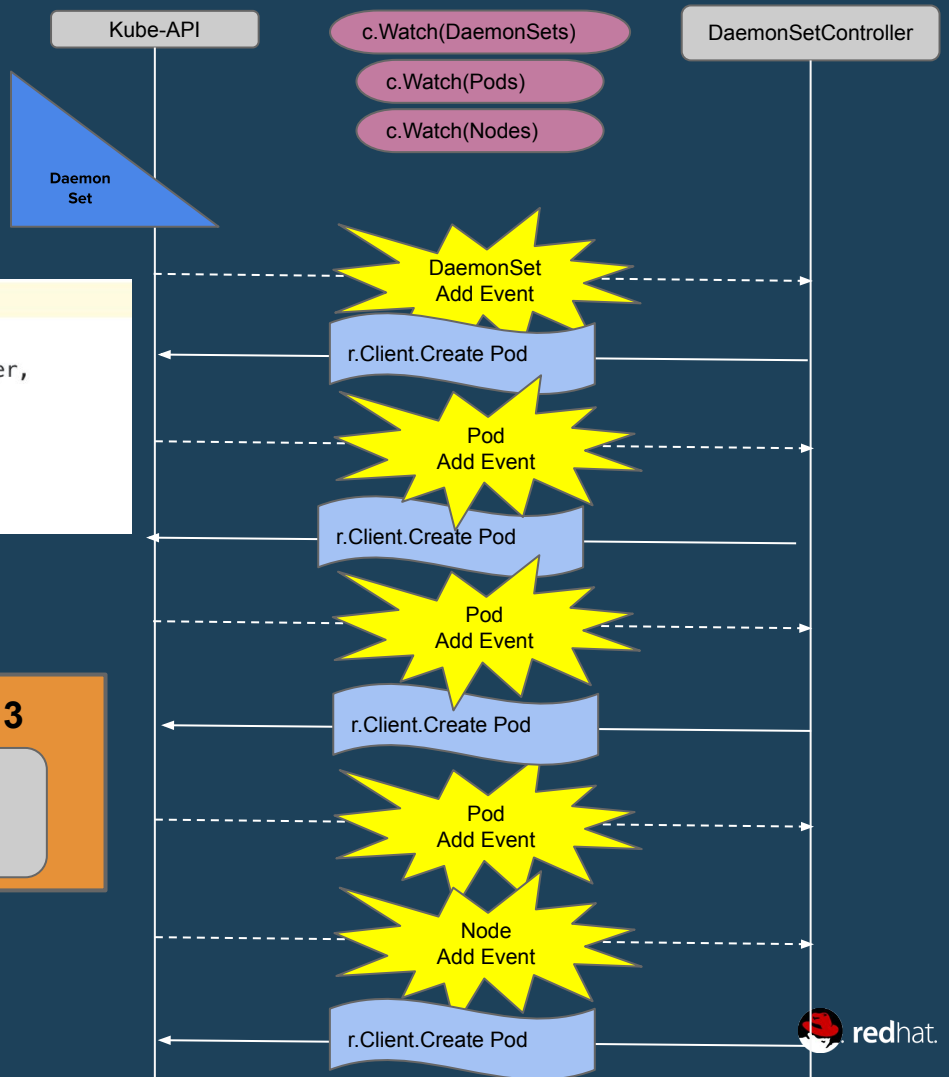

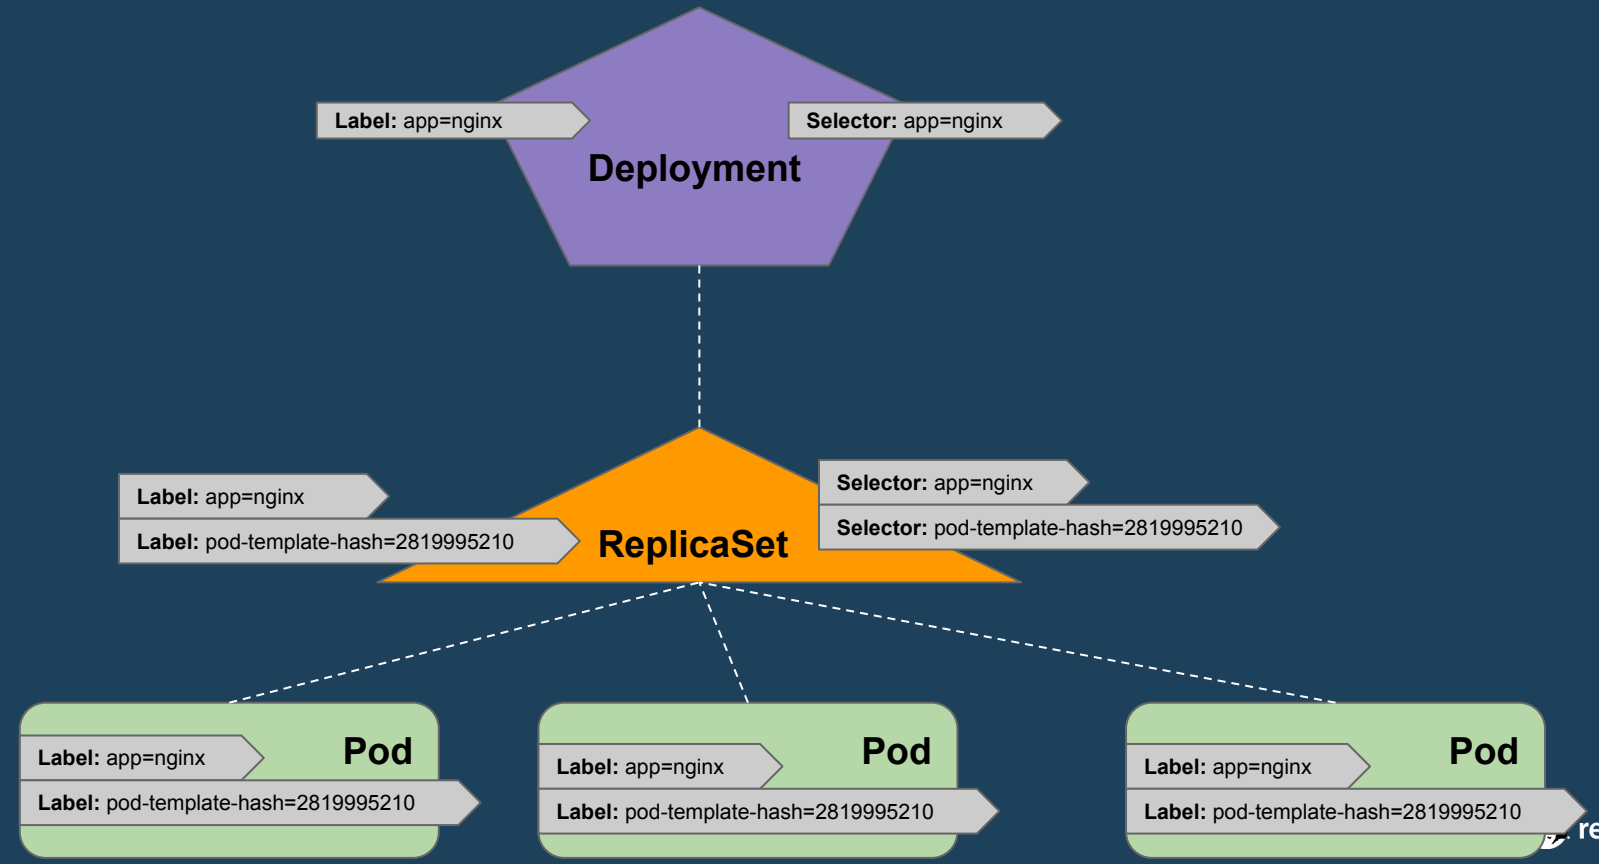
Label: app=nginx
Selector: app=nginx
Deployment
Label: app=nginx
Label: pod-template-hash=2819995210
ReplicaSet
Selector: app=nginx
Selector: pod-template-hash=2819995210
Label: app=nginx
Label: pod-template-hash=2819995210
Pod
Label: app=nginx
Label: pod-template-hash=2819995210
Pod
Label: app=nginx
Label: pod-template-hash=2819995210
Pod
redhat.

DeploymentControll

```
func NewDeploymentController(dInformer appsinformers.DeploymentInformer, rsInformer appsinformers.ReplicaSetInformer, podInfor
        eventBroadcaster := record.NewBroadcaster()
        eventBroadcaster.StartLogging(klog.Infof)
        eventBroadcaster.StartRecordingToSink(&v1core.EventSinkImpl{Interface: client.CoreV1().Events("")})

        if client != nil && client.CoreV1().RESTClient().GetRateLimiter() != nil {
                if err := metrics.RegisterMetricAndTrackRateLimiterUsage("deployment_controller", client.CoreV1().RESTClient().
                        return nil, err
                }
        }
        dc := &DeploymentController{
                client:        client,
                eventRecorder: eventBroadcaster.NewRecorder(scheme.Scheme, v1.EventSource{Component: "deployment-controller"}),
                queue:         workqueue.NewNamedRateLimitingQueue(workqueue.DefaultControllerRateLimiter(), "deployment"),
        }
        dc.rsControl = controller.RealRSControl{
                KubeClient: client,
                Recorder:   dc.eventRecorder,
        }

        dInformer.Informer().AddEventHandler(cache.ResourceEventHandlerFuncs{
                AddFunc:    dc.addDeployment,
                UpdateFunc: dc.updateDeployment,
                // This will enter the sync loop and no-op, because the deployment has been deleted from the store.
                DeleteFunc: dc.deleteDeployment,
        })
        rsInformer.Informer().AddEventHandler(cache.ResourceEventHandlerFuncs{
                AddFunc:    dc.addReplicaSet,
                UpdateFunc: dc.updateReplicaSet,
                DeleteFunc: dc.deleteReplicaSet,
        })
        podInformer.Informer().AddEventHandler(cache.ResourceEventHandlerFuncs{
                DeleteFunc: dc.deletePod,
        })
```

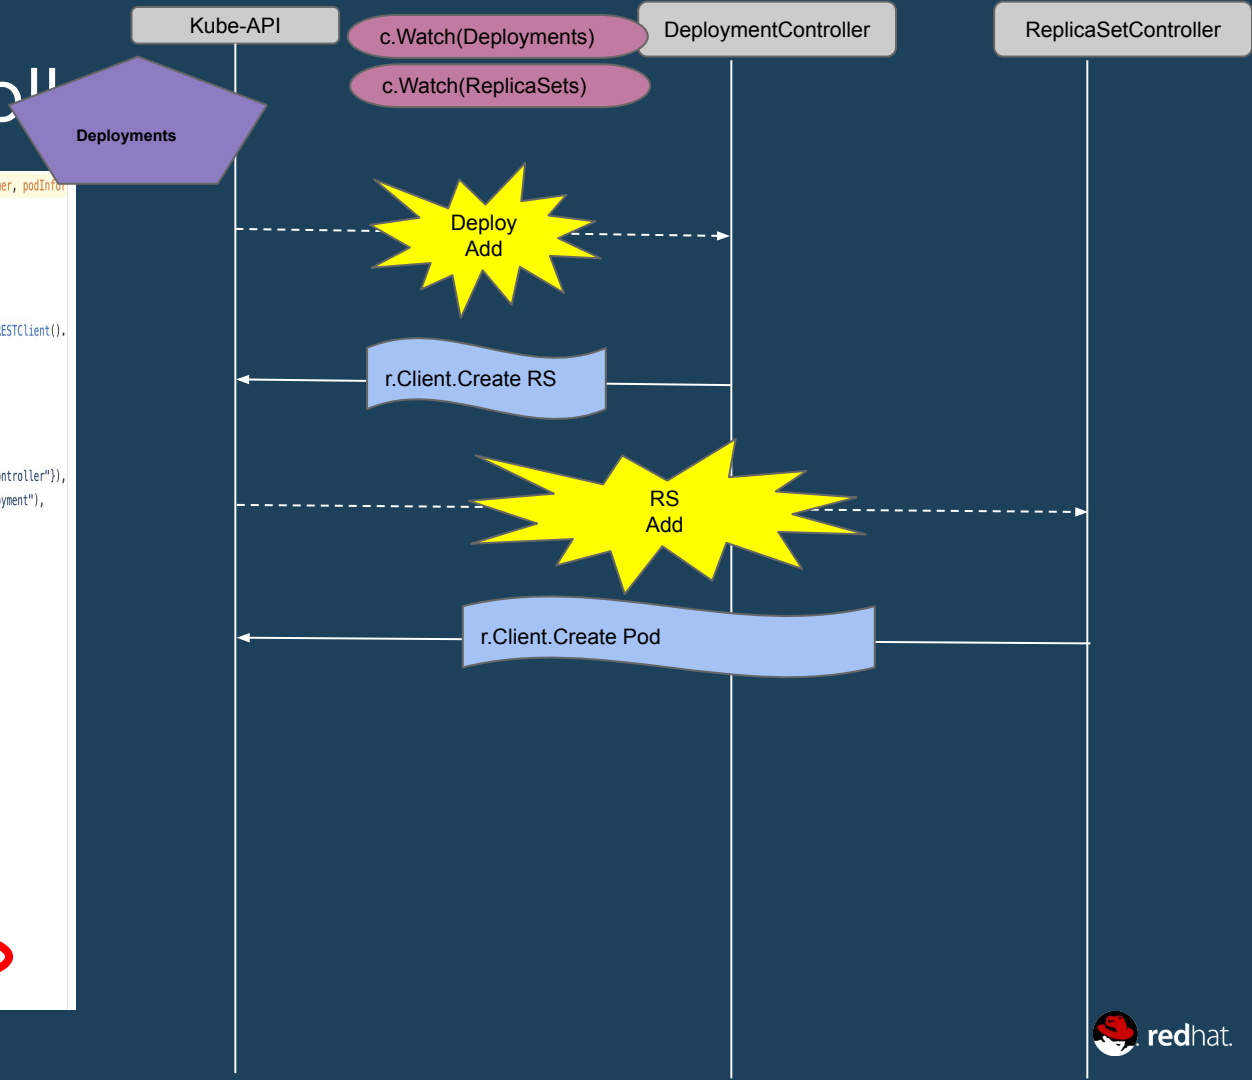

redhat.

Let’s Identify Primary/Secondary Resources for Existing Kubernetes Controllers!

redhat

Garbage Collection (GC)

redhat

Garbage Collection assists in deleting objects that have an owner that **no longer** exists.

controllermanager.go

https://github.com/kubernetes/kubernetes/blob/master/cmd/kube-controller-manager/app/controllermanager.go

```
336 )
337
338 const (
339         saTokenControllerName = "serviceaccount-token"
340 )
341
342 // NewControllerInitializers is a public map of named controller groups (you can start more than one in an init func)
343 // paired to their InitFunc.  This allows for structured downstream composition and subdivision.
344 func NewControllerInitializers(loopMode ControllerLoopMode) map[string]InitFunc {
345         controllers := map[string]InitFunc{}
346         controllers["endpoint"] = startEndpointController
347         controllers["replicationcontroller"] = startReplicationController
348         controllers["podgc"] = startPodGCController
349         controllers["resourcequota"] = startResourceQuotaController
350         controllers["namespace"] = startNamespaceController
351
352         controllers["garbagecollector"] = startGarbageCollectorController
353
354         controllers["job"] = startJobController
355         controllers["deployment"] = startDeploymentController
356         controllers["replicaset"] = startReplicaSetController
357         controllers["horizontalpodautoscaling"] = startHPAController
```

OwnerReferences

```
kubectl create -f  myfirstreplicaset.yaml
```

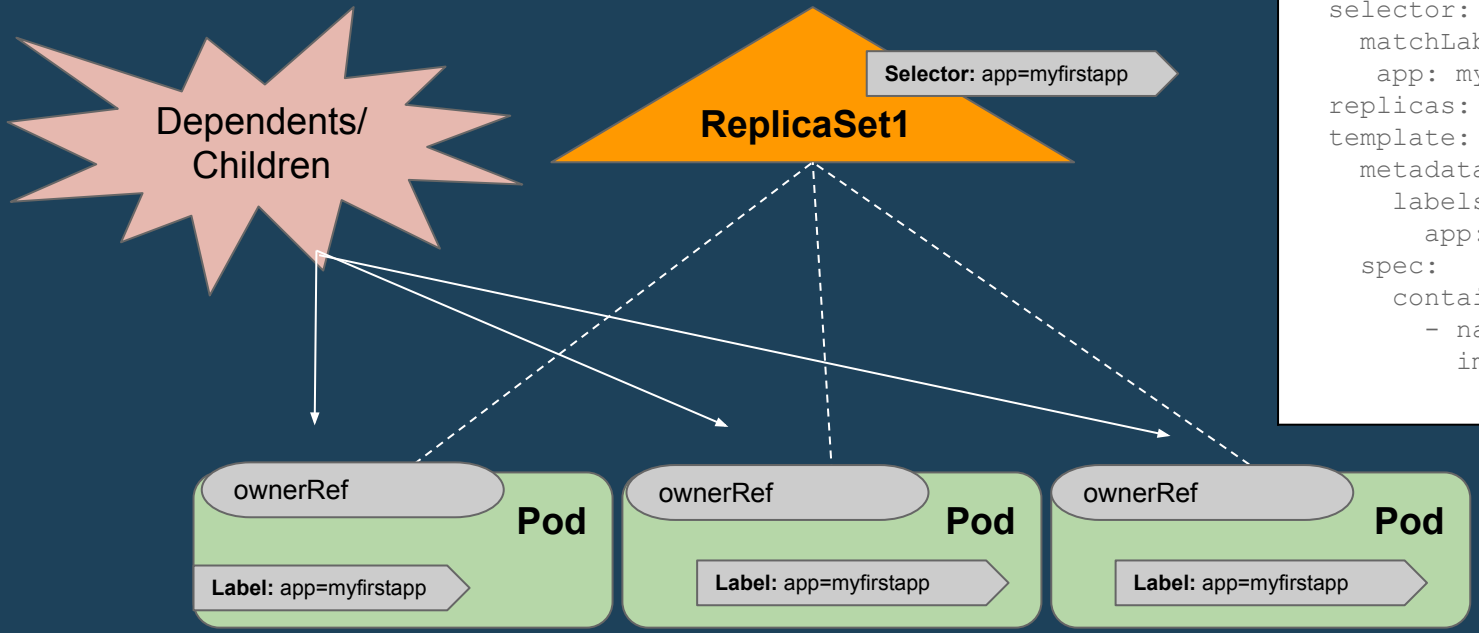


```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: myfirstreplicaset
spec:
  selector:
    matchLabels:
     app: myfirstapp
  replicas: 3
  template:
    metadata:
      labels:
        app: myfirstapp
    spec:
      containers:
        - name: nodejs
          image: myimage
```

redhat.

OwnerReferences

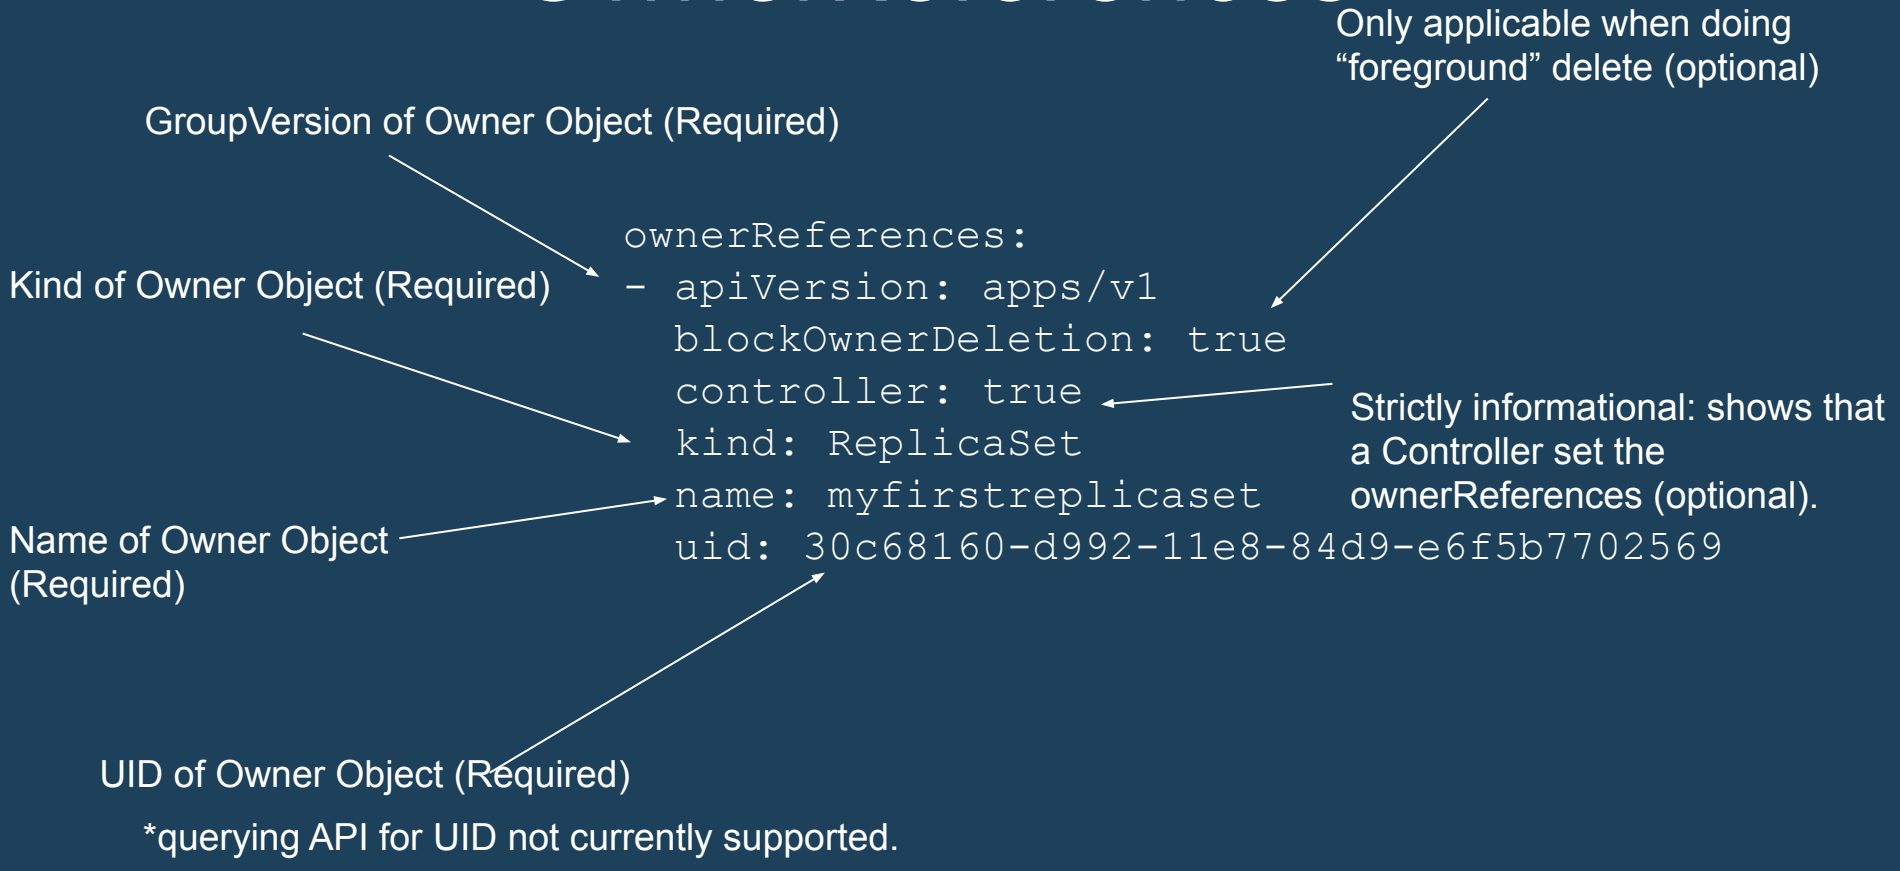

Background Delete

```
oc delete -f myfirstreplicaset
```

```
curl -X DELETE localhost:8080/apis/apps/v1/namespaces/default/replicasets/my-repset \
-d '{"kind":"DeleteOptions","apiVersion":"v1","propagationPolicy":"Background"}' \
-H "Content-Type: application/json"
```

30s..

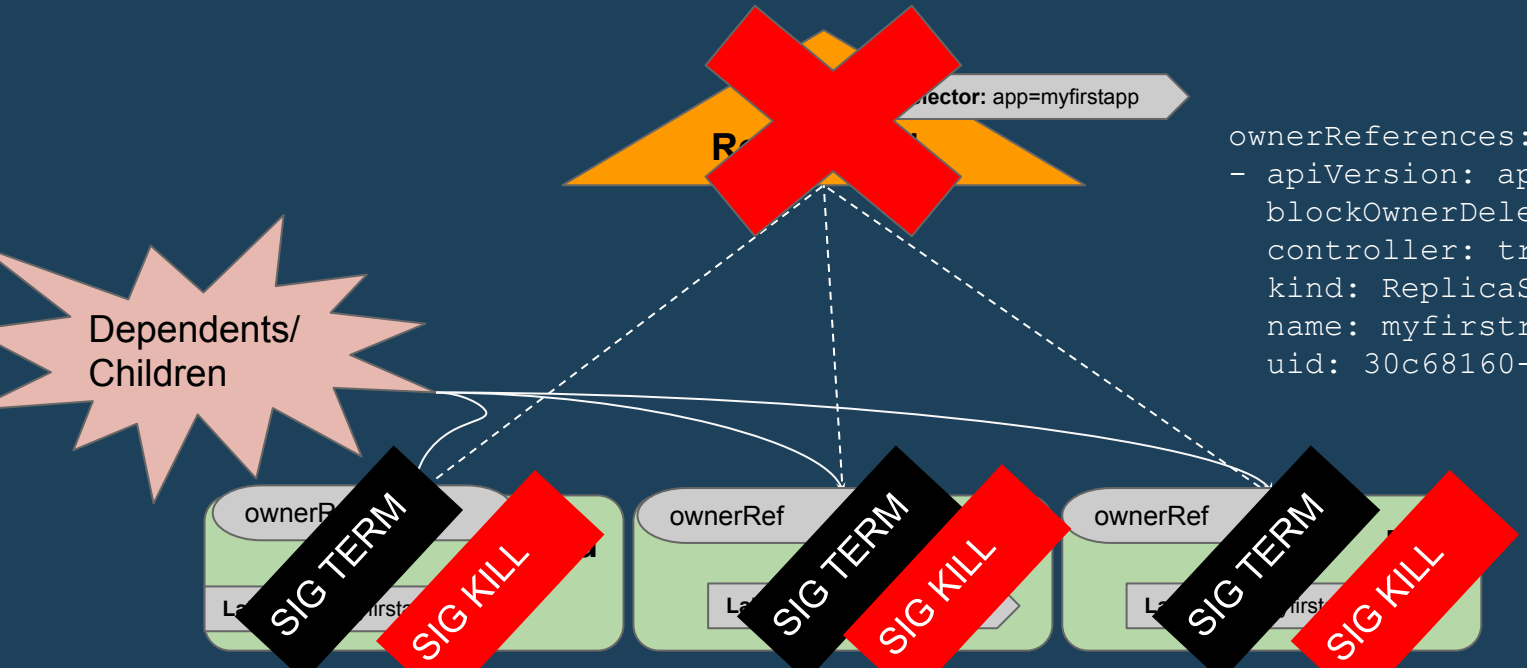


```
ownerReferences:
- apiVersion: apps/v1
  blockOwnerDeletion: true
  controller: true
  kind: ReplicaSet
  name: myfirstreplicaset
  uid: 30c68160-d992-11e8-84d9-e6f5b
```

redhat.

Foreground Delete

```
oc delete -f myfirstreplicaset
```

```
curl -X DELETE localhost:8080/apis/apps/v1/namespaces/default/replicasets/my-repset \
-d '{"kind":"DeleteOptions","apiVersion":"v1","propagationPolicy":"Foreground"}' \
-H "Content-Type: application/json"
```

30s..

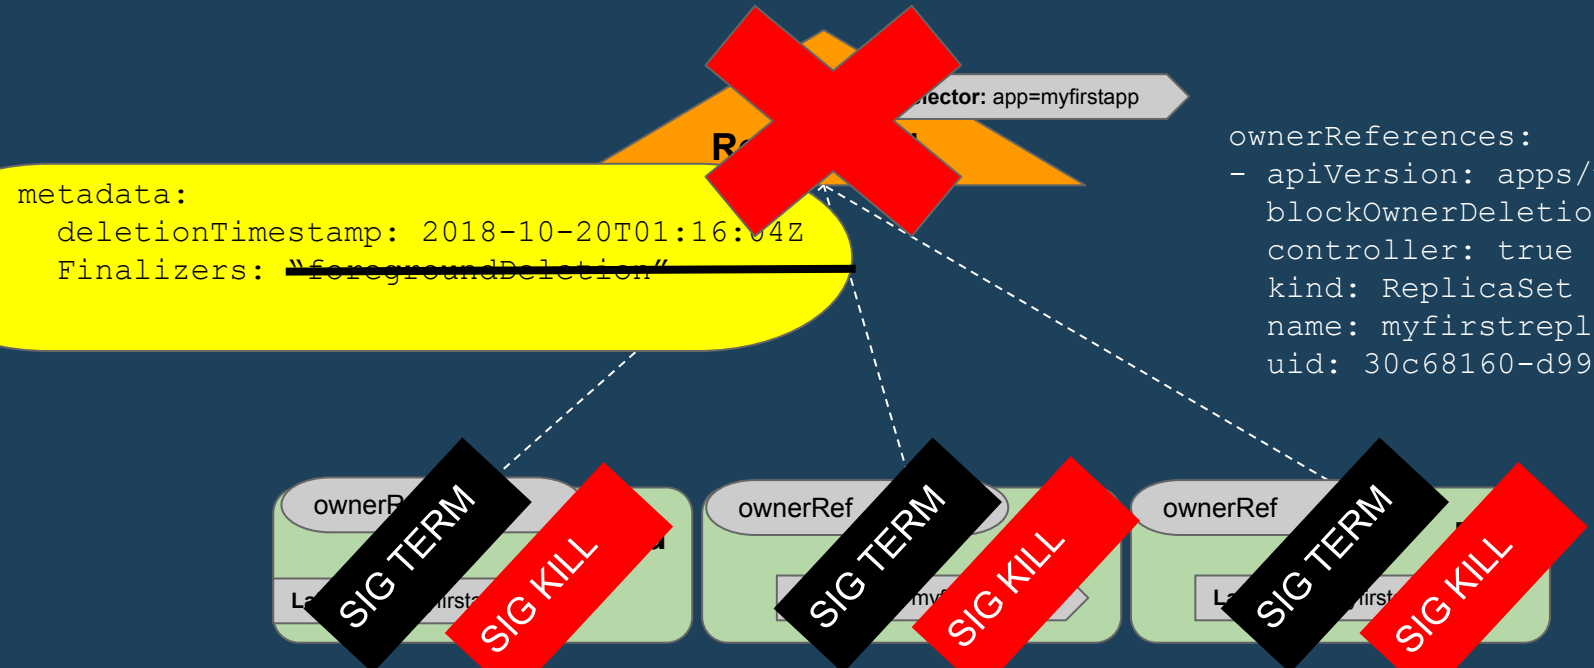

redhat.

Finalizers

Allows controllers to implement conditions that must be completed before the object can be deleted.

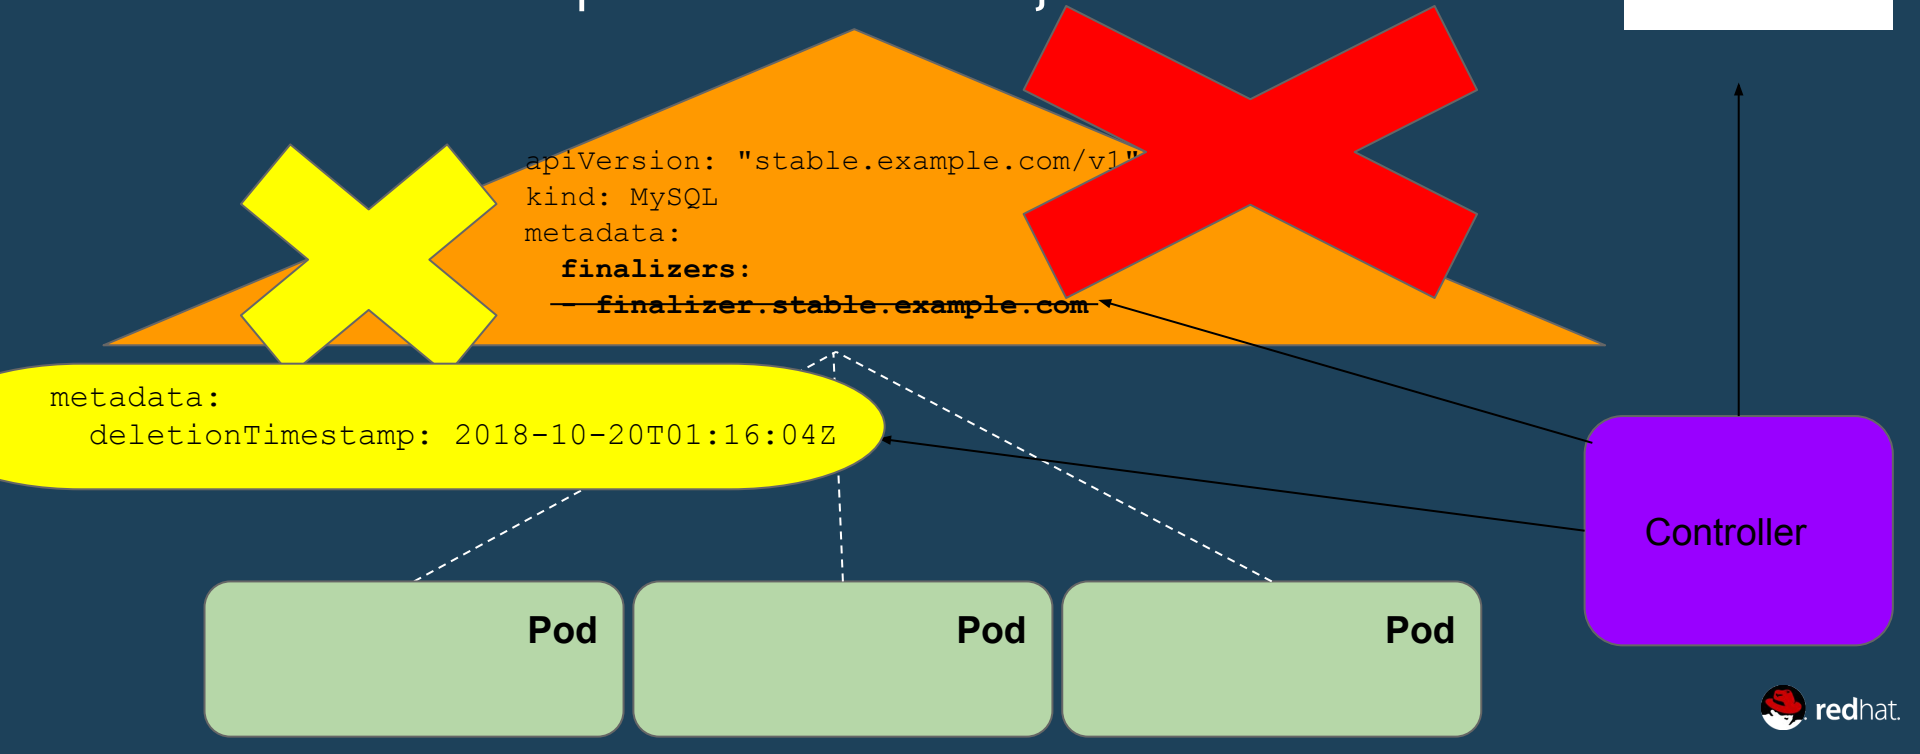

redhat.