

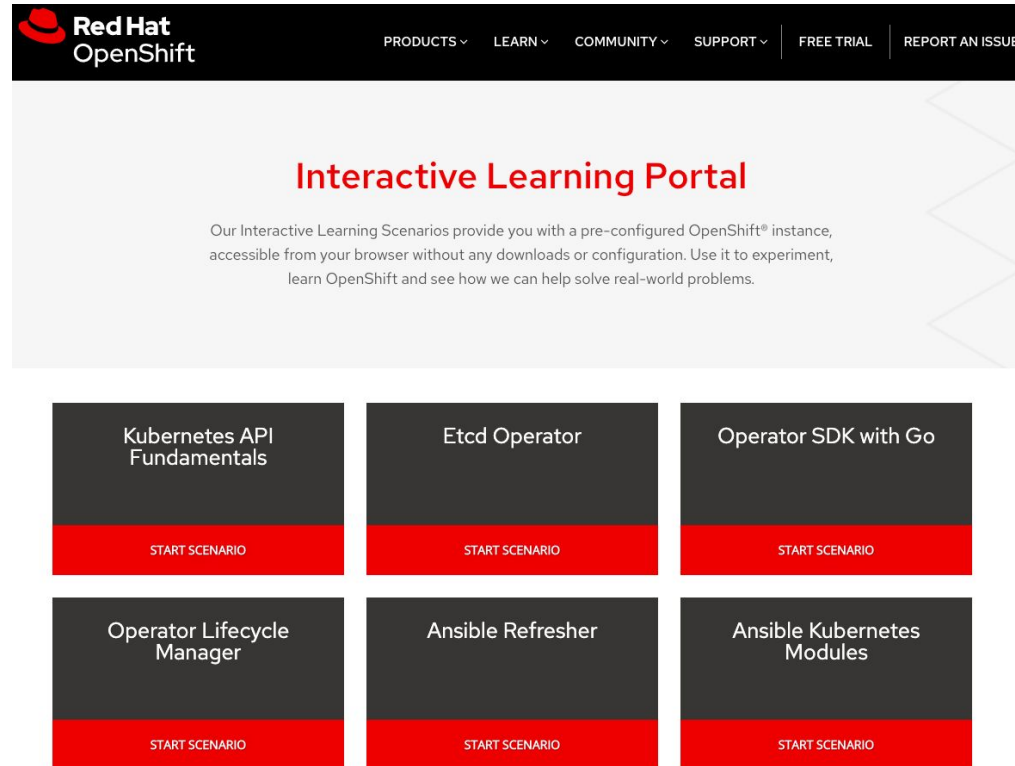


Introducing Operator-SDK 1.0.0

Part of the Operator Framework – a toolkit to manage kubernetes native applications in an effective, automated, scalable way.

Brought to you by the Red Hat Operator Enablement Team

learn.openshift.com/operatorframework



The screenshot shows the Red Hat OpenShift Interactive Learning Portal. At the top is a dark navigation bar with the Red Hat OpenShift logo on the left and links for PRODUCTS, LEARN, COMMUNITY, SUPPORT, FREE TRIAL, and REPORT AN ISSUE on the right. Below the navigation bar is a large light gray section with the title "Interactive Learning Portal" in red. Underneath the title is a paragraph explaining that the interactive learning scenarios provide a pre-configured OpenShift instance accessible from a browser without any downloads or configuration, intended for experimentation and learning. Below this section is a grid of six dark gray cards, each with a red "START SCENARIO" button at the bottom. The cards are titled: "Kubernetes API Fundamentals", "Etcd Operator", "Operator SDK with Go", "Operator Lifecycle Manager", "Ansible Refresher", and "Ansible Kubernetes Modules".

Red Hat OpenShift

PRODUCTS ▾ LEARN ▾ COMMUNITY ▾ SUPPORT ▾ FREE TRIAL REPORT AN ISSUE

Interactive Learning Portal

Our Interactive Learning Scenarios provide you with a pre-configured OpenShift® instance, accessible from your browser without any downloads or configuration. Use it to experiment, learn OpenShift and see how we can help solve real-world problems.

Kubernetes API Fundamentals	Etcd Operator	Operator SDK with Go
START SCENARIO	START SCENARIO	START SCENARIO
Operator Lifecycle Manager	Ansible Refresher	Ansible Kubernetes Modules
START SCENARIO	START SCENARIO	START SCENARIO

HERE TO HELP YOU **SUCCEED**
WITH BATTLE-TESTED TOOLS.




OPERATOR FRAMEWORK



**OPERATOR
SDK**



Building/Dev



**OPERATOR
LIFECYCLE MANAGER**



Install/Manage

WHAT IS AN OPERATOR?

Operators

An operator represents human operational knowledge in software, to reliably manage an application.

[← Back to All Blogs](#)

Introducing Operators: Putting Operational Knowledge into Software

November 03, 2016 • By Brandon Philips


Tags: [announcements](#) [Operators](#)


A Site Reliability Engineer (SRE) is a person that operates an application by writing software. They are an engineer, a developer, who knows how to develop software specifically for a particular application domain. The resulting piece of software has an application's operational domain knowledge programmed into it.

Our team has been busy in the Kubernetes community designing and implementing this concept to reliably create, configure, and manage complex application instances atop Kubernetes.

We call this new class of software Operators. An Operator is an application-specific controller that extends the Kubernetes API to create, configure, and manage instances ¹ of complex ² applications on behalf of a Kubernetes ³ user. It builds upon the basic Kubernetes resource and controller concepts but includes domain or application-specific knowledge to automate common tasks.

It builds upon the basic Kubernetes resource and controller concepts but includes domain or application-specific knowledge to automate common tasks.





Resource

an endpoint in the
Kubernetes API that
stores a collection of API
objects of a certain kind



Pod

the basic execution unit of a Kubernetes application—the smallest and simplest unit in the Kubernetes object model that you create or deploy. A Pod represents processes running on your Cluster.




ConfigMap

provides a way to inject configuration data into Pods. The data stored in a ConfigMap object can be referenced in a volume of type configMap and then consumed by containerized applications running in a Pod.



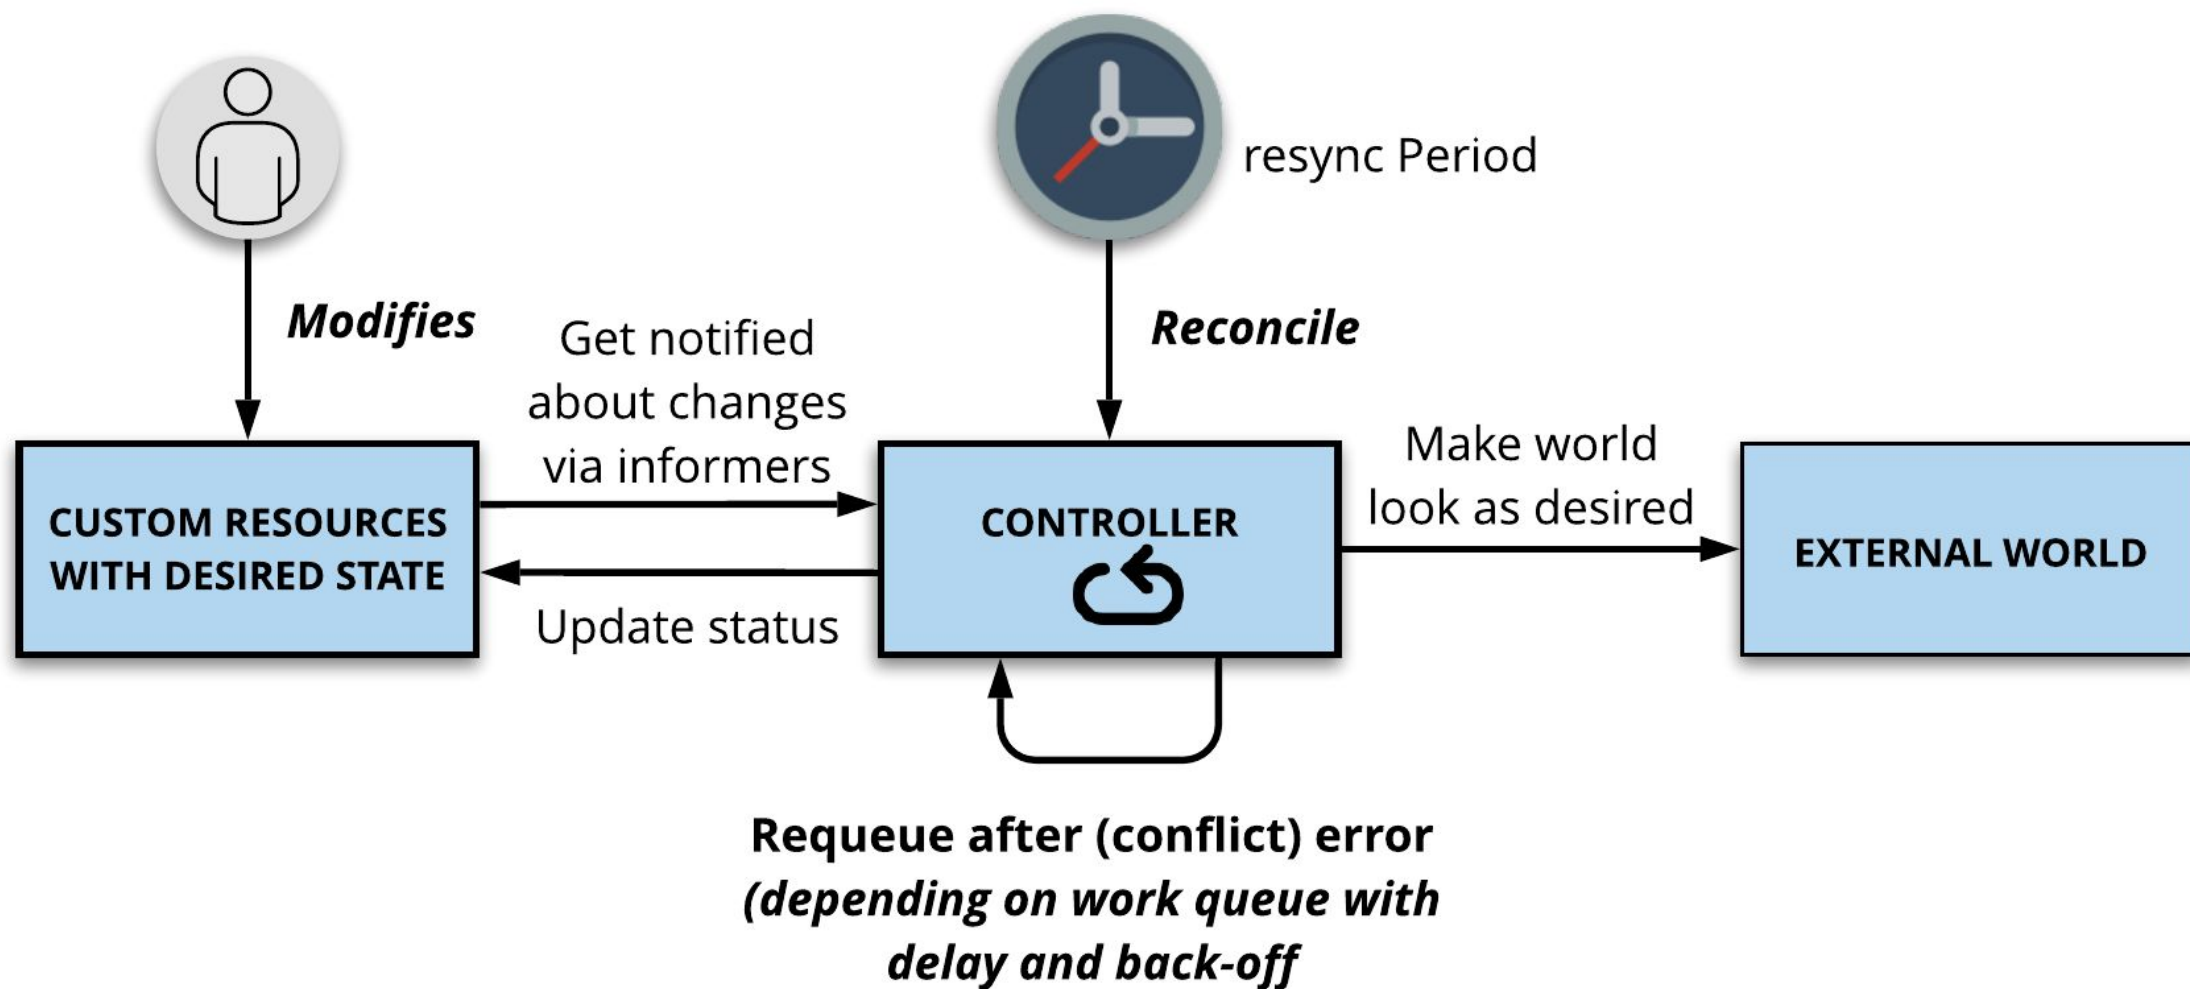
Route (Ingress)

a way to expose a service by giving it an externally-reachable hostname like `www.example.com`.



Controller

control loop that watches the state of your cluster and moves the current cluster state closer to the desired state





ReplicaSet Controller

defined with fields, including a selector that specifies how to identify Pods it can acquire, a number of replicas indicating how many Pods it should be maintaining, and a pod template specifying the data of new Pods it should create to meet the number of replicas criteria.




Deployment Controller

provides declarative updates for Pods and ReplicaSets. You describe a desired state in a Deployment, and the Deployment Controller changes the actual state to the desired state at a controlled rate.



DaemonSet Controller

ensures that all (or some) Nodes run a copy of a Pod. As nodes are added to the cluster, Pods are added to them. As nodes are removed from the cluster, those Pods are garbage collected.



Knowledge

domain or application specific; usually must be learned from users and/or administrators rather than developers

Domain or Application Specific Knowledge

real-world experience with managing your application(s)



Install

Self Heal

Scale

Update


Backup

Clean Up

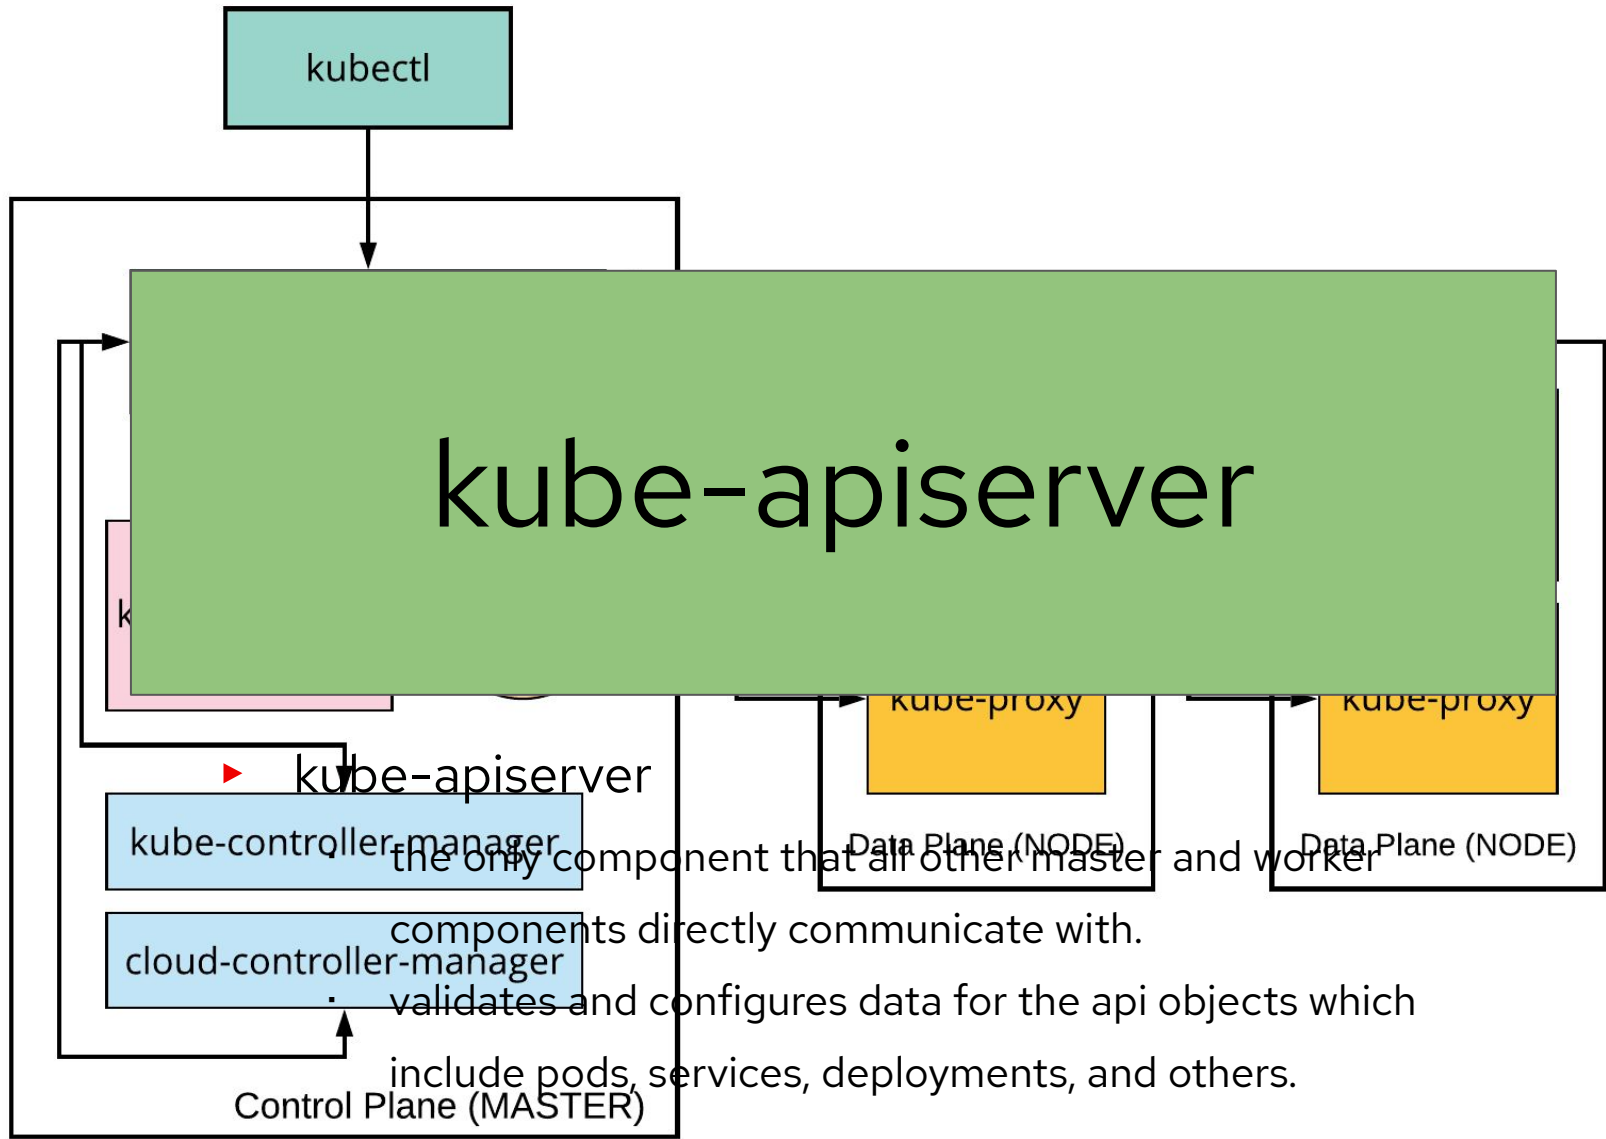
Observability

Resiliency

It builds upon the basic Kubernetes resource and controller concepts but includes domain or application-specific knowledge to automate common tasks.



An Operator takes
advantage of what
Kubernetes does best



the only component that all other master and worker components directly communicate with.
validates and configures data for the api objects which include pods, services, deployments, and others.

```
curl -s localhost:8001/api/v1 | jq -r .resources[].name
```

```
bindings  
componentstatuses  
configmaps  
endpoints  
events  
limitranges  
namespaces  
namespaces/finalize  
namespaces/status  
nodes  
...
```



```
redhat:mhillsma deploy $ oc get -n openshift-dns pods
NAME                READY STATUS  RESTARTS  AGE
...
dns-default-vxvth  3/3    Running  0         5d8h
```

```
(curl -s -XGET localhost:8001/api/v1/namespaces/openshift-dns/pods | jq -r
.items[].metadata.name)
```

```
"dns-default-478pn"
"dns-default-4fv5s"
"dns-default-vxvth"
"dns-default-7k289"
"dns-default-fw7gv"
"dns-default-j7mzv"
```

```
redhat:mhillsma deploy $ oc get -n openshift-dns pod/dns-default-vxvth -o yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
...
```

```
  name: dns-default-vxvth
```

```
(curl -XGET localhost:8001/api/v1/namespaces/openshift-dns/pods/dns-default-vxvth)
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: dns-default-vxvth
```

```
  namespace: openshift-dns
```

```
  ownerReferences:
```

```
...
```

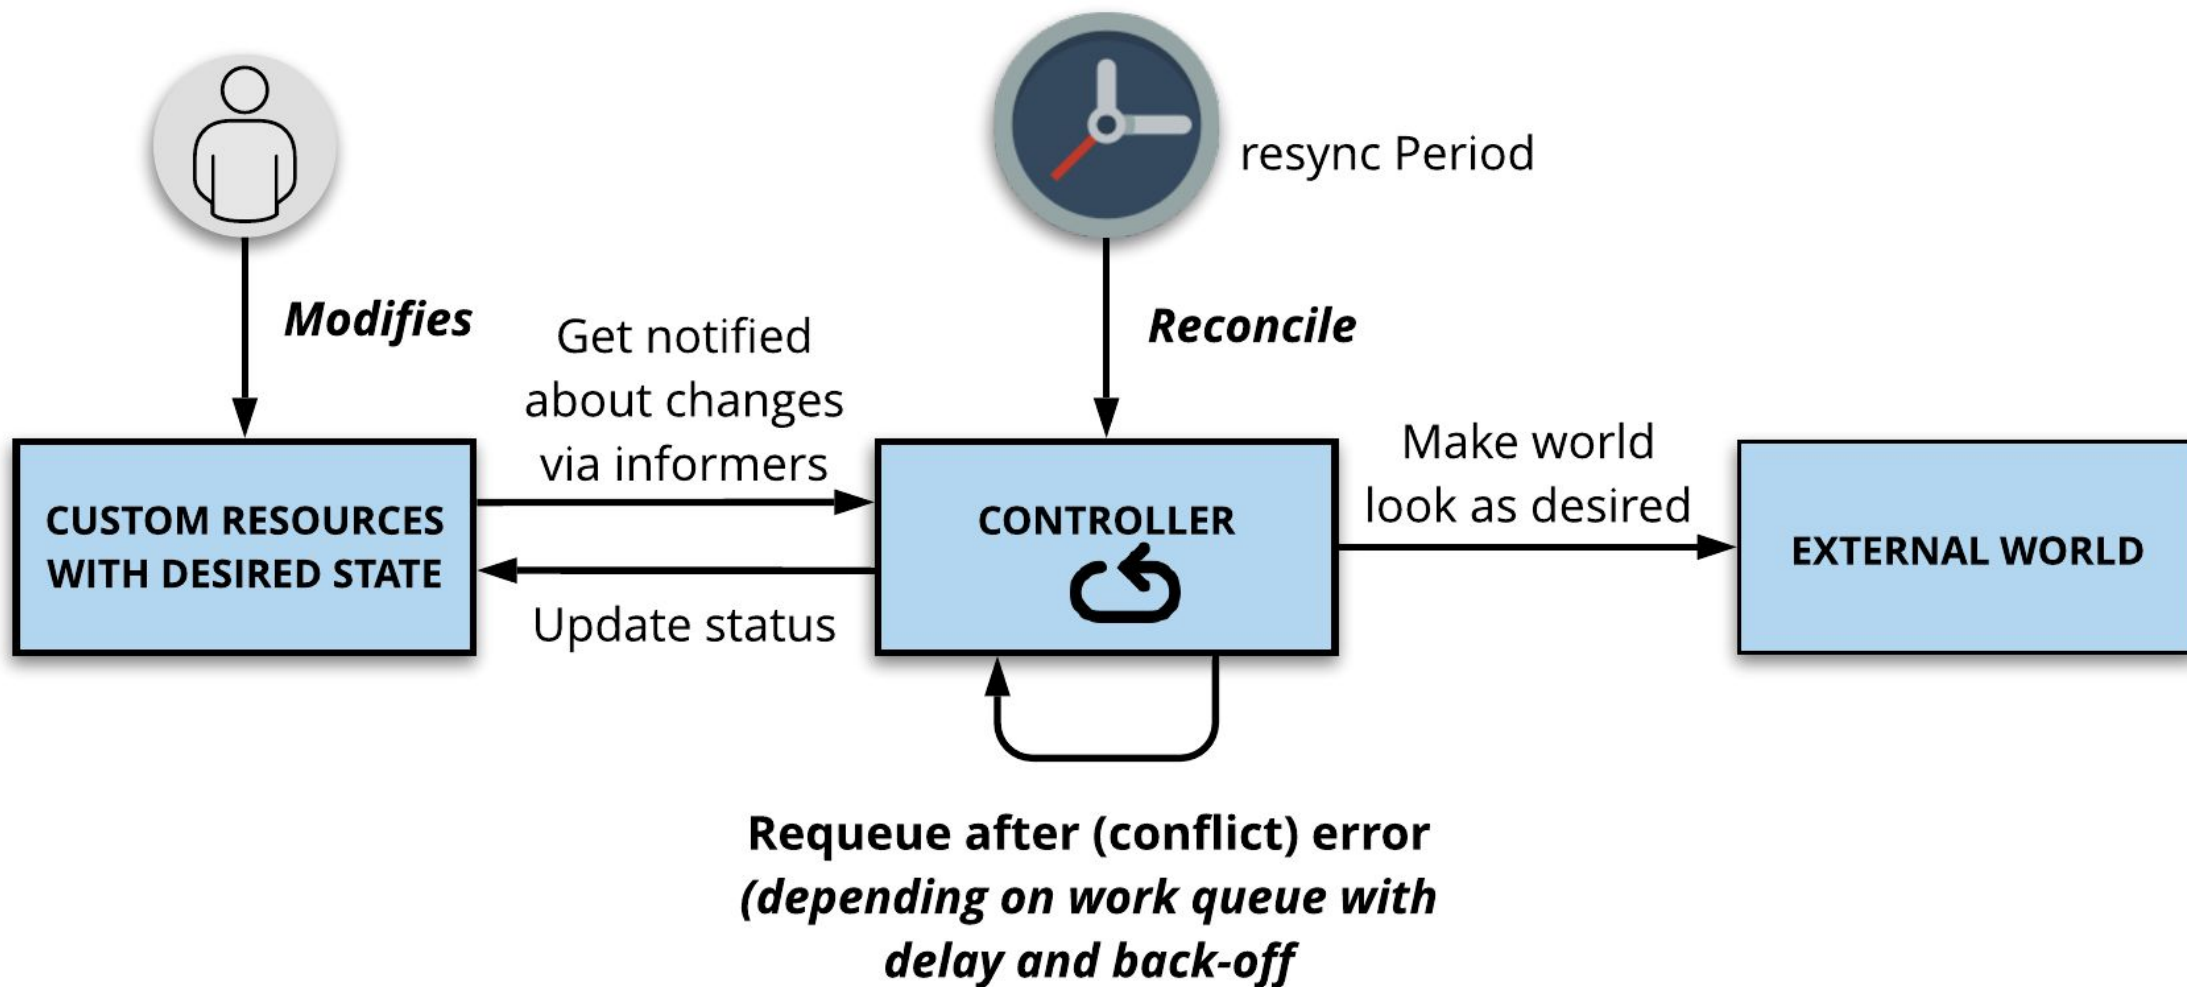
CRDs allow us to **EXTEND** the Kubernetes API

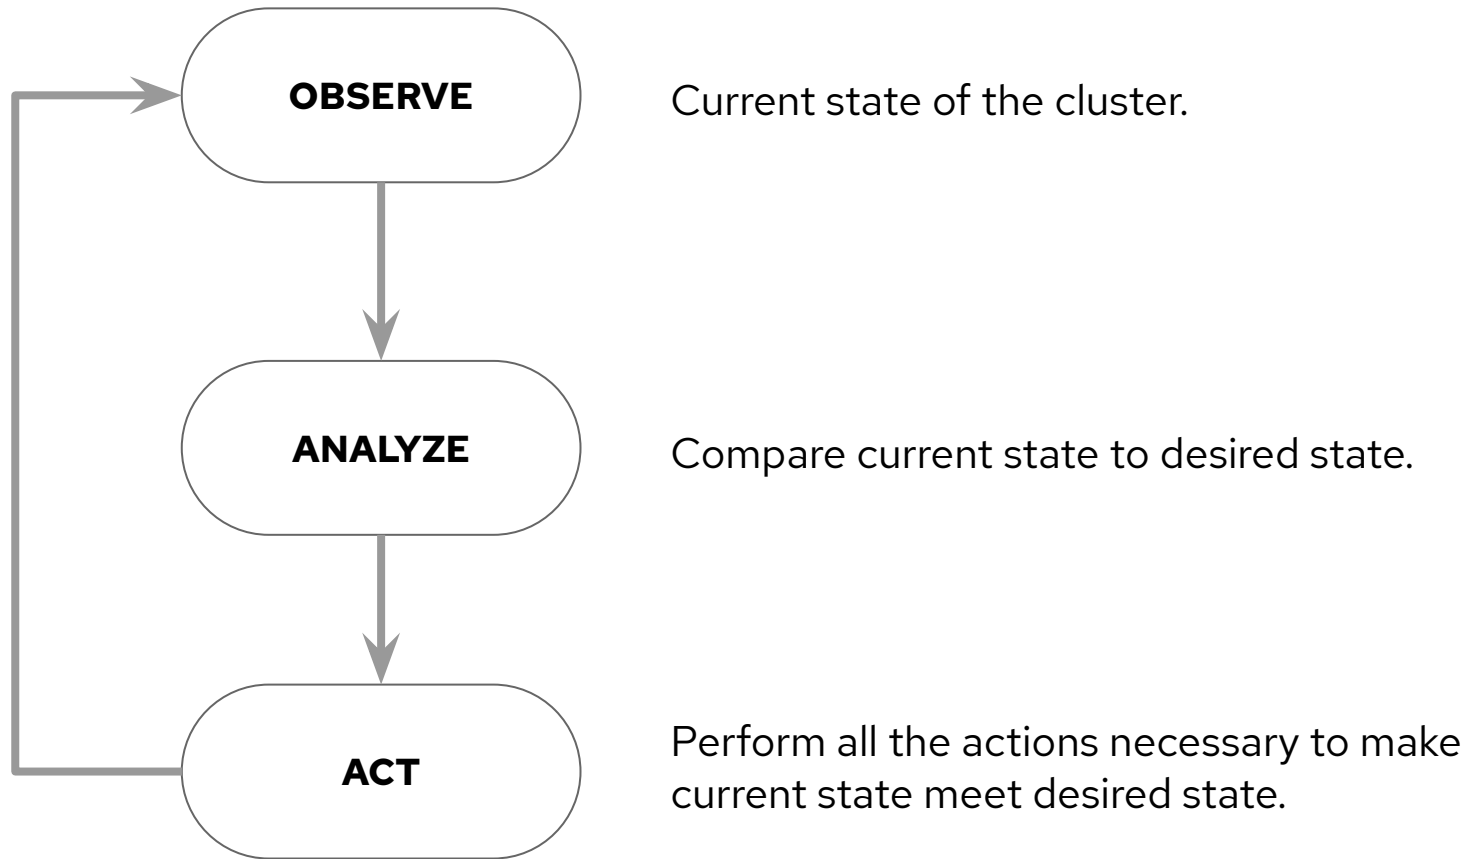
- ▶ modify the API without recompiling
- ▶ create our very own API resource/object
- ▶ resource/object exists but nothing acts on its presence and this is where controllers come in

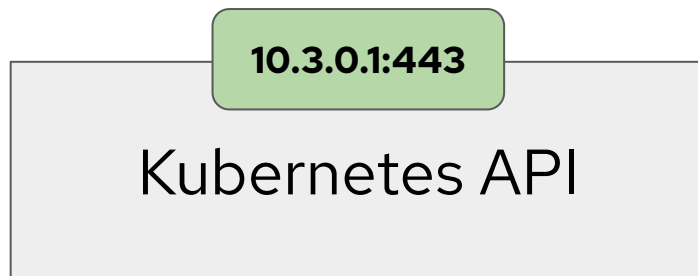
A Custom Resource
needs a controller
to **ACT**
upon its presence.

What do we mean by **ACT**?




- ▶ Create
- ▶ Read
- ▶ Update
- ▶ Delete



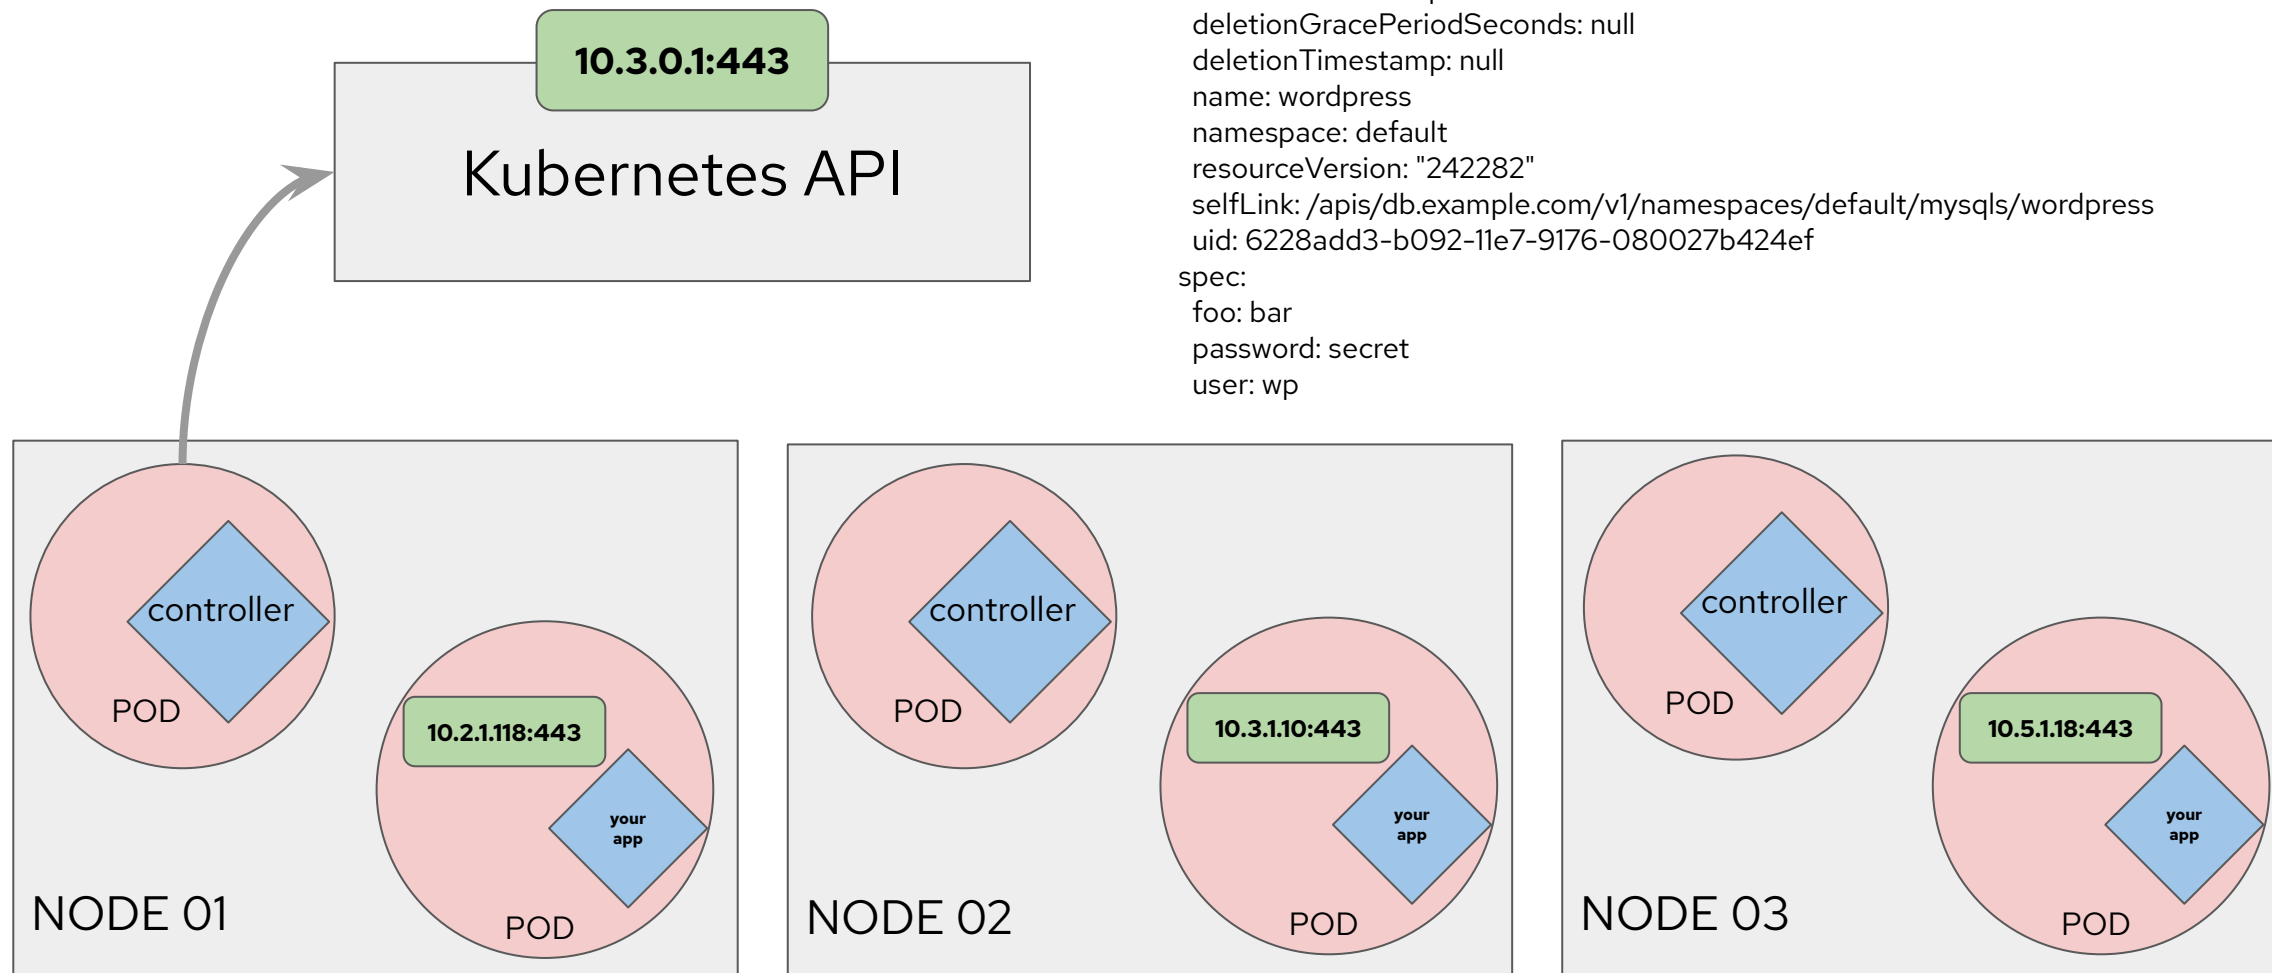




```
apiVersion: db.example.com/v1
kind: MySql
metadata:
  clusterName: ""
  creationTimestamp: 2017-10-14T03:47:21Z
  deletionGracePeriodSeconds: null
  deletionTimestamp: null
  name: wordpress
  namespace: default
  resourceVersion: "242282"
  selfLink: /apis/db.example.com/v1/namespaces/default/mysqls/wordpress
  uid: 6228add3-b092-11e7-9176-080027b424ef
spec:
  foo: bar
  password: secret
  user: wp
```



What is an Operator



What do we mean by **ACT**?

- ▶ Create
- ▶ Read
- ▶ Update
- ▶ Delete


Create, Read, Update, Delete...Probably Not Enough

- Server startup/shutdown
- Mastering the mysqladmin administrative client
- Using the mysql interactive client
- User account maintenance
- Log file maintenance
- Database backup/copying
- Hardware tuning
- Multiple server setups
- Software updates and upgrades
- File system security
- Server security
- Repair and maintenance
- Crash recovery
- Preventive maintenance
- Understanding the mysqld server daemon
- Performance analysis
- Choosing what else to install (e.g. Apache, Perl +modules, PHP)
- Which version of MySQL (stable, developer, source, binary)
- Creating a user account for the mysql user and group
- Download and unpack a distribution
- Compile source code and install (or rpm)
- Initialize the data directory and grant tables with mysql_install_db
- Starting the server
- Installing Perl DBI support
- Installing PHP
- Installing Apache
- Obtaining and installing the samp_db sample database

- Securing a new MySQL installation
- Running mysqld as an unprivileged user
 - Methods of starting the server
 - Invoking mysqld directly
 - Invoking safe_mysqld
 - Invoking mysql.server
 - Specifying startup options
 - Checking tables at startup
 - Shutting down the server
- Regaining control of the server if you can't connect
- Creating new users and granting privileges
- Determining who can connect from where
 - Who should have what privileges?
 - Administrator privileges
 - Revoking privileges
 - Removing users
 - deciding/finding the Data Directory's location
 - Structure of the Data Directory
 - How mysqld provides access to data
- Running multiple servers on a single Data Directory
 - Database representation
 - Table representation (form, data and index files)
 - OS constraints on DB and table names
- Data Directory structure and performance, resources, security
 - MySQL status files (.pid, .err, .log, etc)
 - Relocating Data Directory contents

- Creating new users and granting privileges
- Determining who can connect from where
 - Who should have what privileges?
 - Administrator privileges
 - Revoking privileges
 - Removing users
 - Methods: mysqldump vs. direct copying
 - Backup policies
 - Scheduled cycles
 - Update logging
 - Consistent and comprehensible file-naming
 - Backing up the backup files
 - Off-site / off-system backups
 - Backing up an entire database with mysqldump
 - Compressed backup files
 - Backing up individual tables
 - Using mysqldump to transfer databases to another server
 - mysqldump options (flush-logs, lock-tables, quick, opt)
 - Direct copying methods
 - Database replication (live and off-line copying)
 - Recovering an entire database
 - Recovering grant tables
 - Recovering from mysqldump vs. tar/cpio files
 - Using update logs to replay post-backup queries
 - Editing update logs to avoid replaying erroneous queries
 - Recovering individual tables
 - Default parameters

It builds upon the basic Kubernetes resource and controller concepts but includes domain or application-specific knowledge to automate common tasks.



Why do Operators matter to us at Red Hat?

Why Operators Matter to Red Hat

- ▶ Build an ecosystem of software on OpenShift that can be as easy, safe, and reliable to use and operate as a Cloud Service.
- ▶ Low-touch, remotely managed, one-click-updates.
- ▶ Super easy to deploy in an Operator in a Kubernetes environment.

OperatorHub

Operator Management

Workloads

Networking

Storage

Builds

Monitoring

Compute

Administration

All Items

AI/Machine Learning

Big Data

Database

Integration & Delivery

Logging & Tracing

Monitoring

Networking

OpenShift Optional

Security

Storage

Streaming & Messaging

Other

Filter by keyword...


INSTALL STATE

☐ Installed (0)

☐ Not Installed (34)

All Items


34 items



AMQ Streams

provided by Red Hat, Inc.

Red Hat AMQ Streams is a massively scalable, distributed, and high performance data streami




Community

Aqua Security Operator

provided by Aqua Security, Inc.

The Aqua Security Operator runs within a Openshift cluster and provides a means to deploy and manage Aqu




Community

Automation Broker Operator

provided by Red Hat, Inc.

Automation Broker is an implementation of the Open Service Broker API managi




Community

Camel-K Operator

provided by The Apache Software Foundation

Apache Camel K (a.k.a. Kamel) is a lightweight integration framework built from Apac




Community

Cluster Logging

provided by Red Hat, Inc

The Cluster Logging Operator for OKD provides a means for configuring and managing your aggregated logging




Community

CockroachDB

provided by Helm Community

CockroachDB Operator based on the CockroachDB helm chart

39



How do I create my very own Operator?

Life Before the Operator SDK

If only it were as simple as...

Resources

```
type MyCustomResourceDefinition struct {  
    // API obj kind & schema version  
    metav1.TypeMeta  
    // Standard object metadata (optional)  
    Metadata api.ObjectMeta  
    // Describe how the resource appears  
    Spec v1beta1.CustomResourceDefinitionSpec  
    // State of the CRD  
    Status CustomResourceDefinitionStatus  
}
```

Controllers

```
for {  
    current := getCurrentState()  
    desired := getDesiredState()  
    makeChanges(current, desired)  
}
```

Writing Operator from scratch is Challenging

- ▶ Research client-library.
- ▶ Repo organization.
- ▶ Write boiler-plate code.
- ▶ Use code generators.
- ▶ Knowledge of informers/shared informers and work queues for object cache and event handling.

We need an
easier way to
create
Operators

We need an
easier way to
manage
Operators



OPERATOR FRAMEWORK

Operator SDK

DEVELOP IN GO, ANSIBLE, OR HELM

GO

1. Create a new operator project using the SDK Command Line Interface (CLI)
2. Define new resource APIs by adding Custom Resource Definitions (CRD)
3. Define Controllers to watch and reconcile resources
4. Write the reconciling logic for your Controller using the SDK and controller-runtime APIs
5. Use the SDK CLI to build and generate the operator deployment manifests

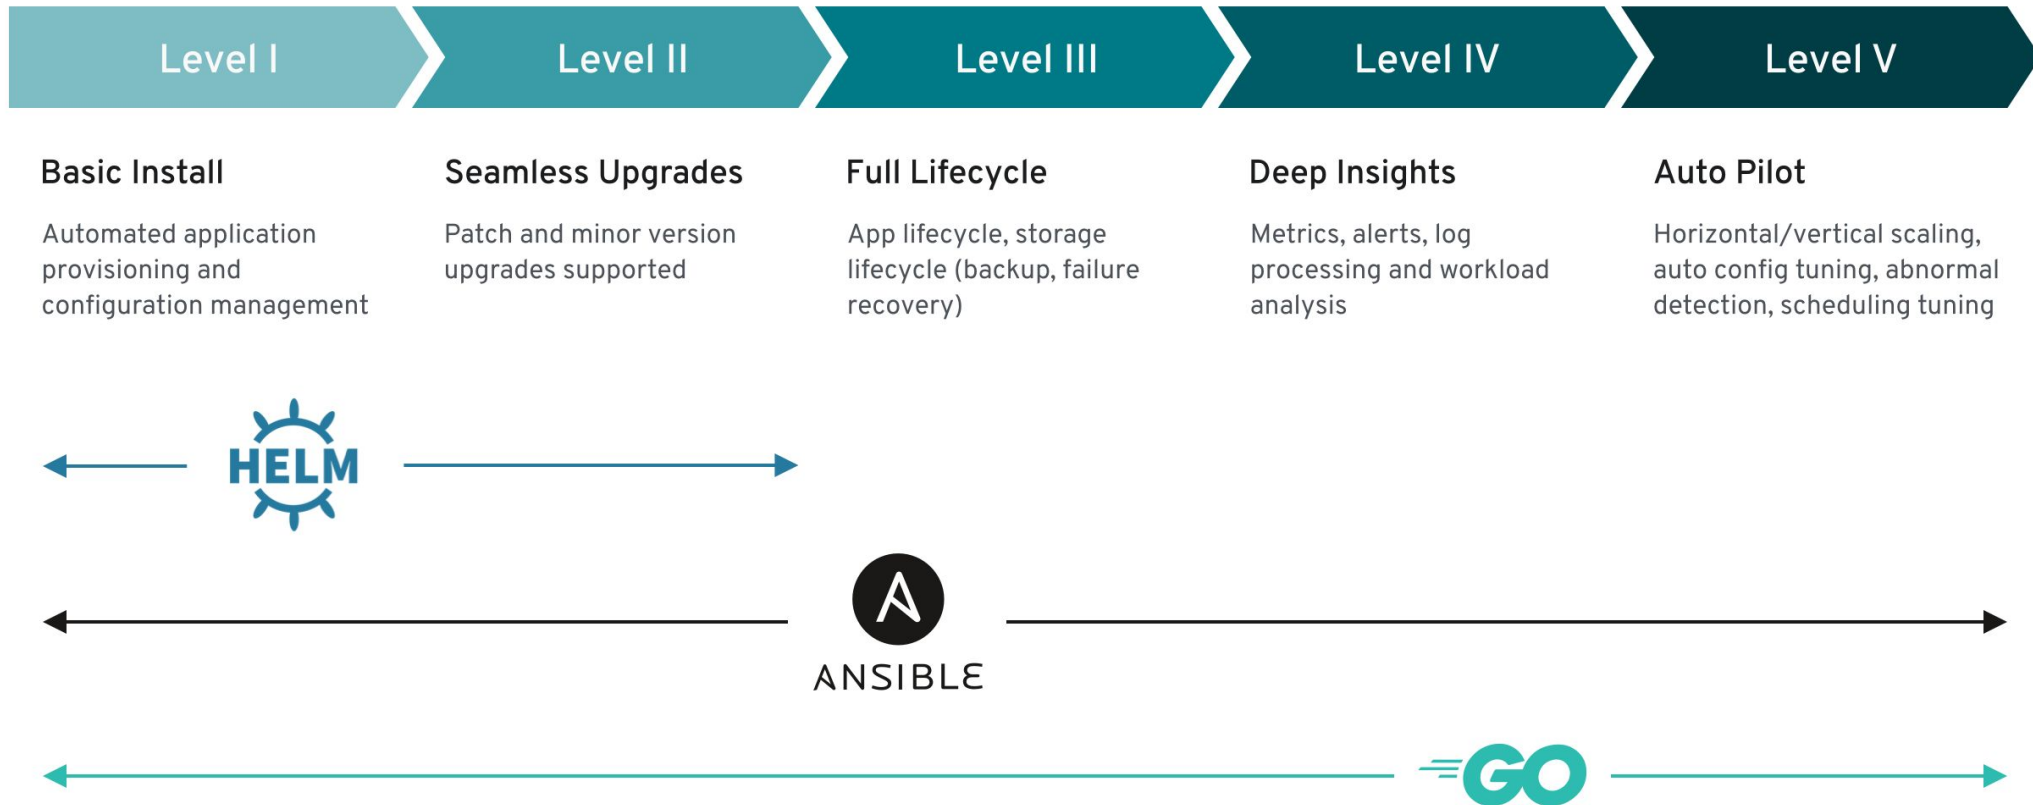
ANSIBLE

1. Create a new operator project using the SDK Command Line Interface (CLI)
2. Write the reconciling logic for your object using ansible playbooks and roles
3. Use the SDK CLI to build and generate the operator deployment manifests
4. Optionally add additional CRD's using the SDK CLI and repeat steps 2 and 3

HELM

1. Create a new operator project using the SDK Command Line Interface (CLI)
2. Create a new (or add your existing) Helm chart for use by the operator's reconciling logic
3. Use the SDK CLI to build and generate the operator deployment manifests
4. Optionally add additional CRD's using the SDK CLI and repeat steps 2 and 3

Operator SDK



Operator Lifecycle Manager

WHAT IS OPERATOR LIFECYCLE MANAGER?

This project is a component of the [Operator Framework](#), an open source toolkit to manage Kubernetes native applications, called Operators, in a streamlined and scalable way.

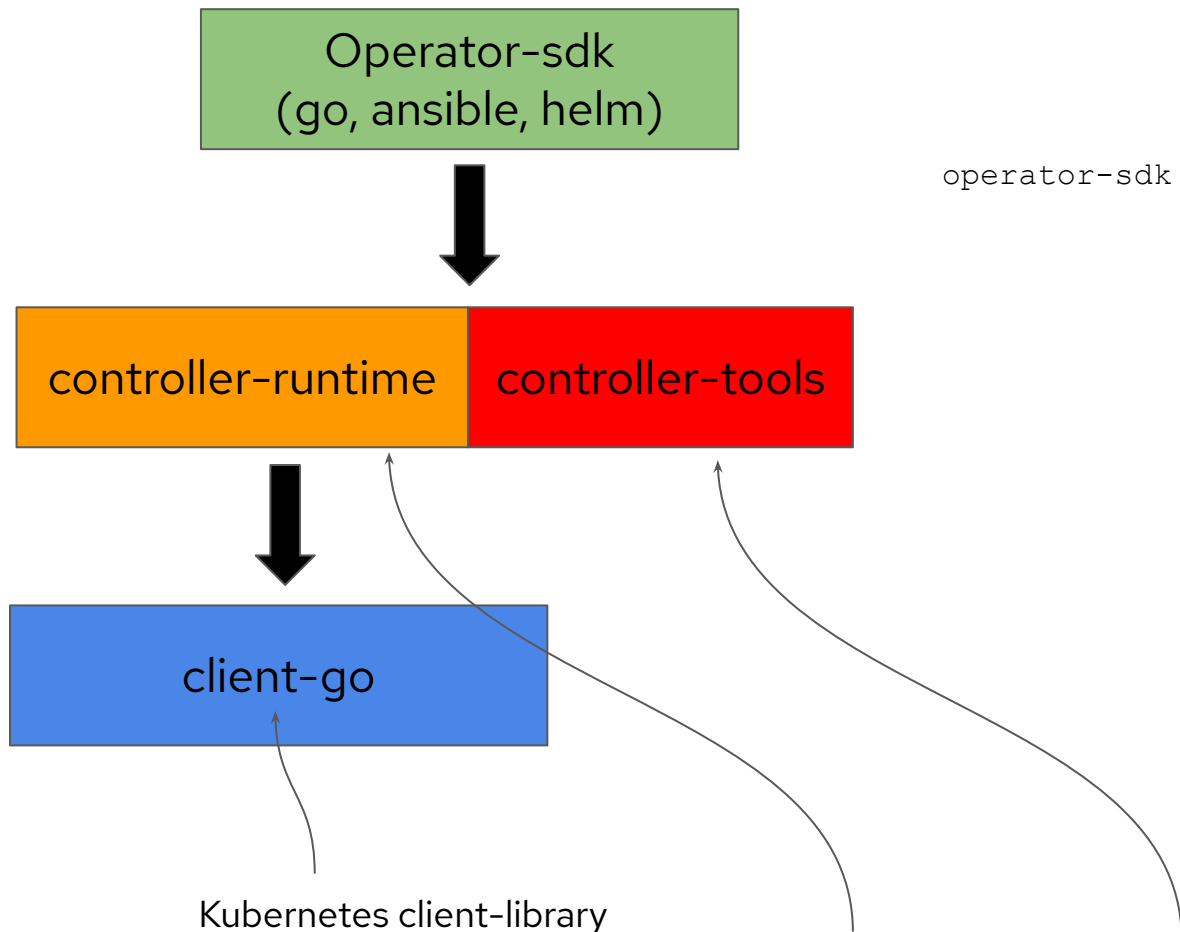
OLM FEATURES

OVER-THE-AIR UPDATES AND CATALOGS	DEPENDENCY MODEL	DISCOVERABILITY	CLUSTER STABILITY	DECLARATIVE UI CONTROLS
OLM provides rich update mechanisms to keep Kubernetes native applications up to date automatically.	With OLMs packaging format Operators can express dependencies on the platform and on other Operators.	OLM makes Operators and their services available for cluster users to select and install.	OLM will prevent conflicting Operators owning the same APIs being installed, ensuring cluster stability.	OLM enables Operators to behave like managed service providers through the APIs they expose.

About Operator-SDK

How things were before..

Operator-SDK (released in 2018 by RedHat)



Libraries for building the controller part of your operator

```
operator-sdk new create app-operator --type=go

operator-sdk add api --api-version=app.example.com/v1alpha1 --kind=App

operator-sdk generate k8s
operator-sdk generate crds

operator-sdk add controller --api-version=app.example.com/v1alpha1 --kind=App

operator-sdk run --local --kubeconfig=

operator-sdk build quay.io/example/operator:v0.0.1
podman push quay.io/example/operator:v0.0.1

operator-sdk olm install

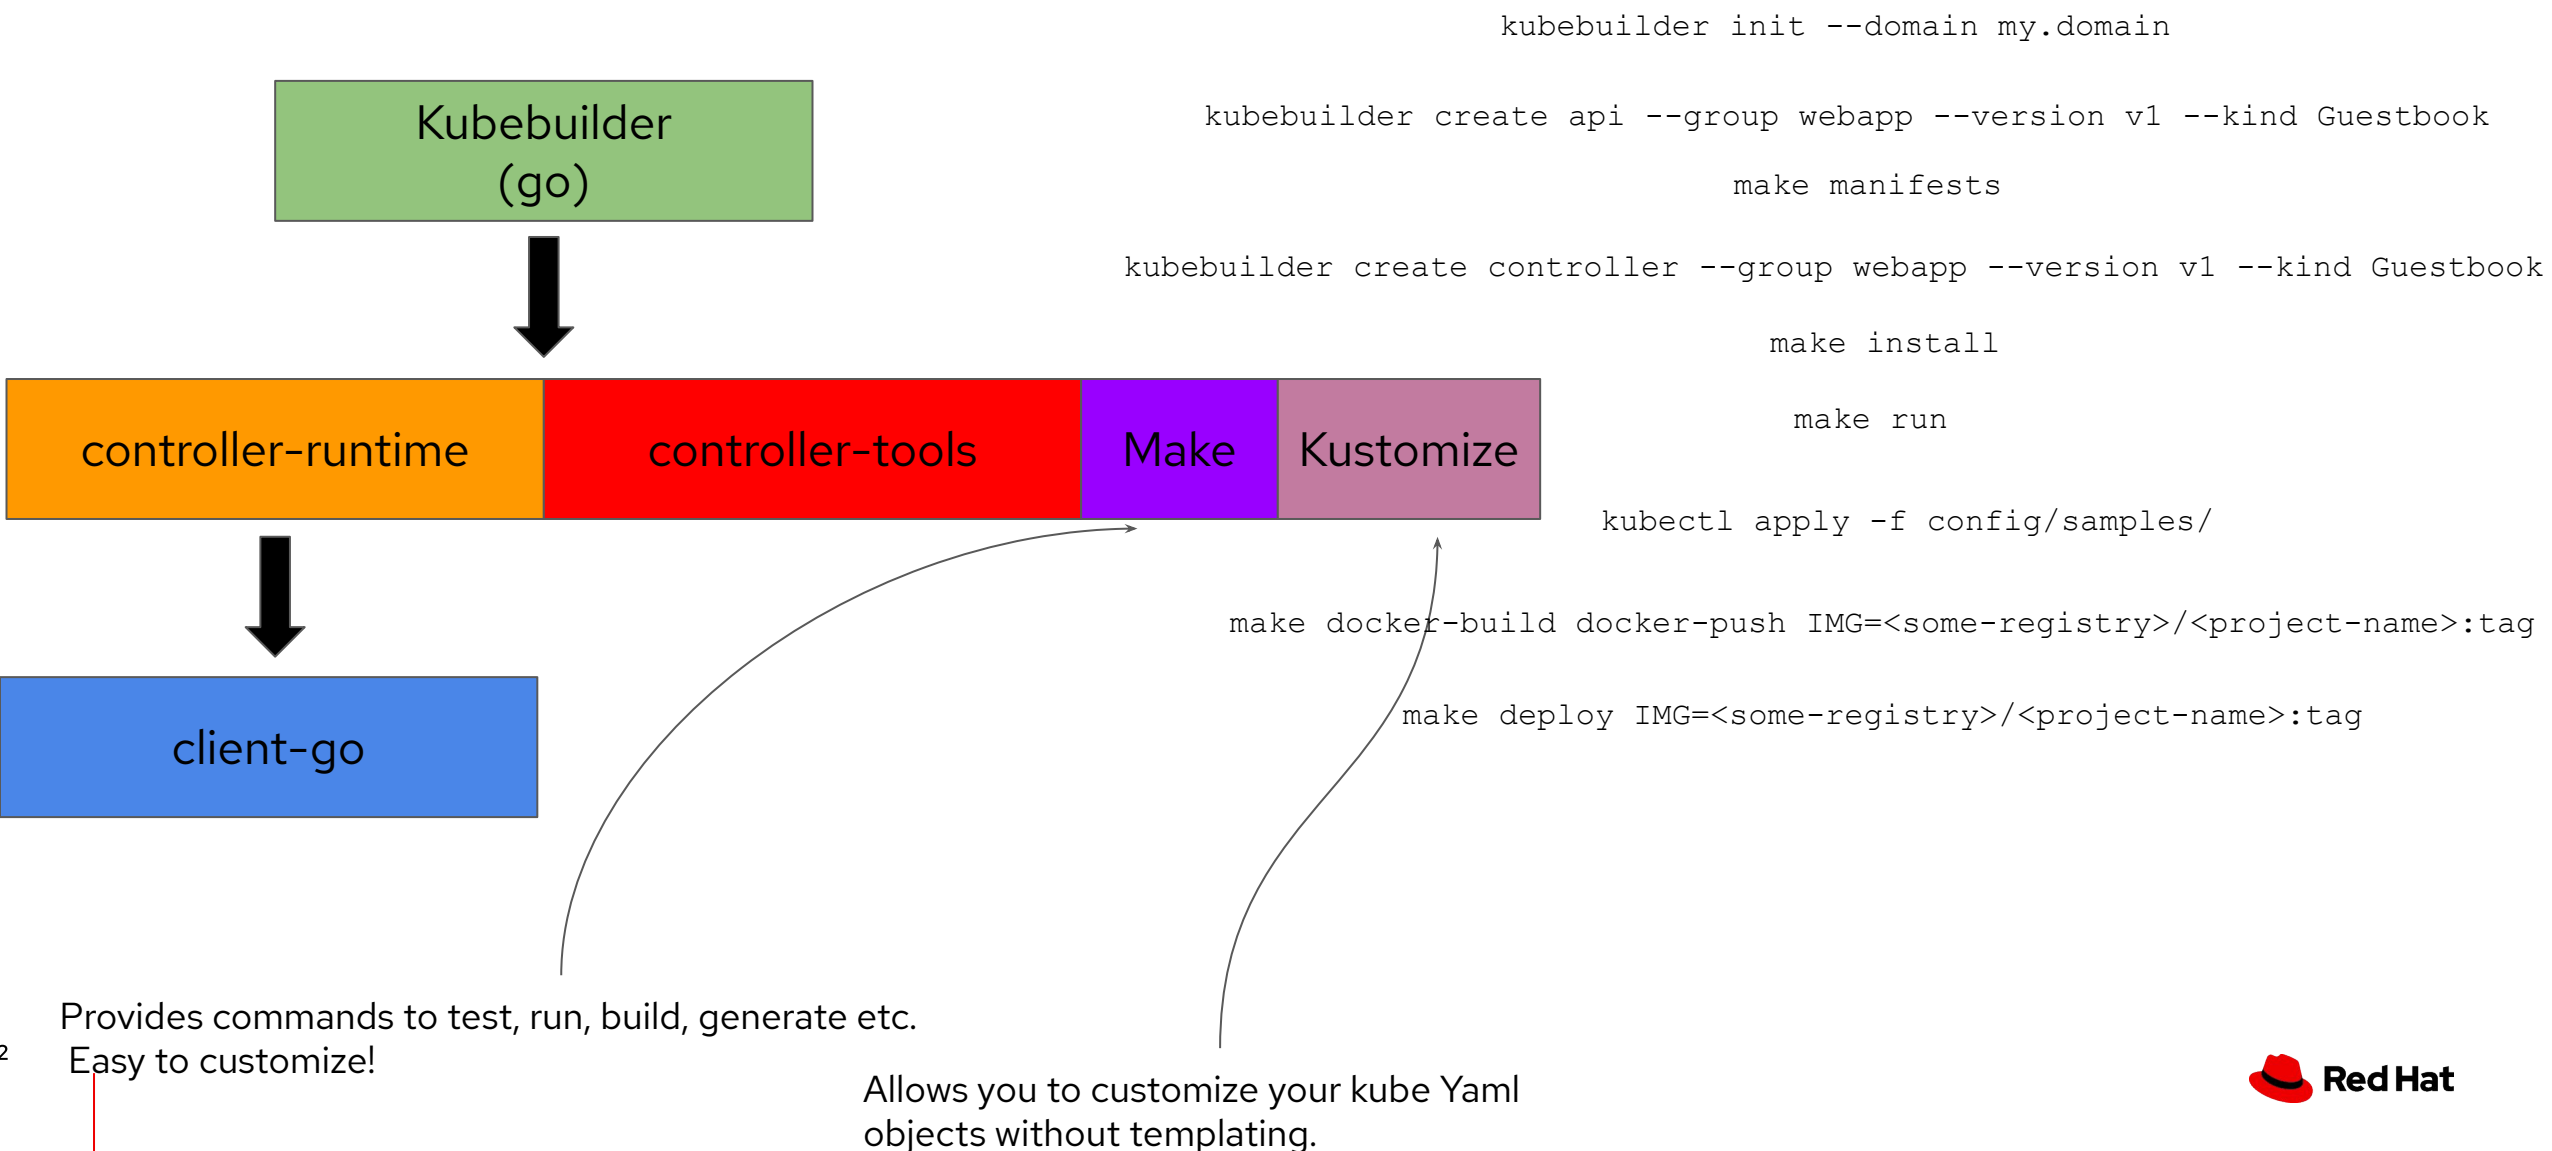
operator-sdk bundle create quay.io/example/operator:v0.0.1 \
  --directory ./deploy/olm-catalog/test-operator \
  --package test-operator \
  --channels stable,beta \
  --default-channel stable

podman build -t quay.io/example/operator-bundle:v0.0.1 -f
upstream-example.Dockerfile .

podman push quay.io/example/operator-bundle:v0.0.1
```

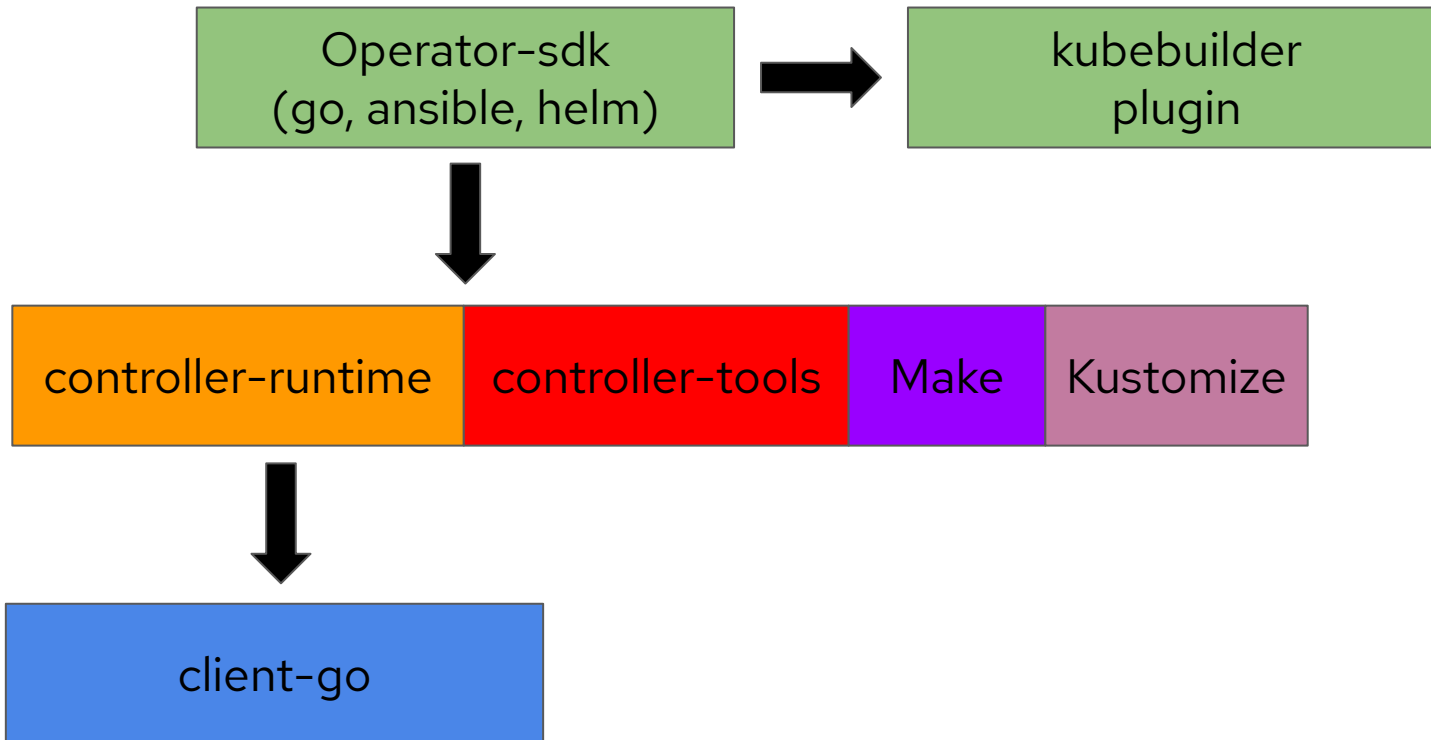
Tools for generating custom resource definitions, rbac artifacts, and more!

Kubebuilder (released in 2018 by API Machinery group)



Now with Operator-SDK 1.0.0...

Operator-SDK (1.0.0)



```
operator-sdk init --domain my.domain
```

```
operator-sdk create api --group webapp \  
--version v1 --kind Guestbook
```

```
make manifests
```

```
operator-sdk create controller --group webapp \  
--version v1 --kind Guestbook
```

```
make install
```

```
etc...
```

What else is new?

Separate binaries for Go, Ansible, and Helm

```
operator-sdk init --plugins=ansible --domain example.com
```

Support for Webhooks

```
operator-sdk create webhook --group batch --version v1 --kind CronJob --defaulting --programmatic-validation
```

Use Kustomize!

```
kustomize build config/manifests | operator-sdk generate bundle --overwrite --version 0.0.1
```

Other stuff...

Live Demo...but
first..some review..

Resource Schema Components

GVK aka TypeMeta

Metadata aka ObjectMeta

Spec


Status

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
```

```
metadata:
  name: my-first-replica-set
  namespace: myproject
```

```
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 5
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
```


```
status:
  availableReplicas: 1
  fullyLabeledReplicas: 1
  observedGeneration: 1
  readyReplicas: 1
  replicas: 1
```



Chapter 4

Designing Infrastructure Applications

*The **reconciler pattern** is a software pattern that can be used or expanded upon for managing cloud native infrastructure. The pattern enforces the idea of having two representations of the infrastructure—the first being the actual state of the infrastructure, and the second being the expected state of the infrastructure.*



*The **reconciler pattern** will force the engineer to have two independent avenues for getting either of these representations, as well as to implement a solution to reconcile the actual state into the expected state.*

ReplicaSets in Action!

```
kubectl create -f myfirstreplicaset.yaml
```

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: myfirstreplicaset
spec:
  selector:
    matchLabels:
      app: myfirstapp
  replicas: 3
  template:
    metadata:
      labels:
        app: myfirstapp
    spec:
      containers:
        - name: nodejs
          image: myimage
```

0 < spec.replicas?

1 < spec.replicas?

Selector: app=myfirstapp

ReplicaSet1

2 < spec.replicas?

3 < spec.replicas?

Pod

Pod

Pod

Label:
app=myfirstapp

Label: app=myfirstapp

Label:
app=myfirstapp

Kube-API

c.Watch(Replicaset)

ReplicaSetController

c.Watch(Pods, OwnerType: ReplicaSet)

ReplicaSet
Add Event

r.Client.List Pods by label: rs.metadata.label

r.Client.Create Pod 1

Pod 1
Add Event

r.Client.List Pods by label: rs.metadata.label

r.Client.Create Pod 2

Pod 2
Add Event

r.Client.List Pods by label: rs.metadata.label

r.Client.Create Pod 3

Pod 3
Add Event

r.Client.List Pods by label.metadata.label

ReplicaSets in Action!

```
kubectl create -f myfirstreplicaset.yaml
```

```
apiVersion: extensions/v1beta1
kind: ReplicaSet
metadata:
  name: myfirstreplicaset
spec:
  selector:
    matchLabels:
      app: myfirstapp
  replicas: 3
  template:
    metadata:
      labels:
        app: myfirstapp
    spec:
      containers:
        - name: nodejs
          image: myimage
```

2 < spec.replicas?

Selector: app=myfirstapp

ReplicaSet1

3 < spec.replicas?

Pod

Label: app=myfirstapp

Pod

Label: app=myfirstapp

Pod

Label: app=myfirstapp

c.Watch(Replicaset)

ReplicaSetController

c.Watch(Pods, OwnerType: ReplicaSet)

Pod 1
Delete Event

r.Client.List Pods by label: rs.metadata.label

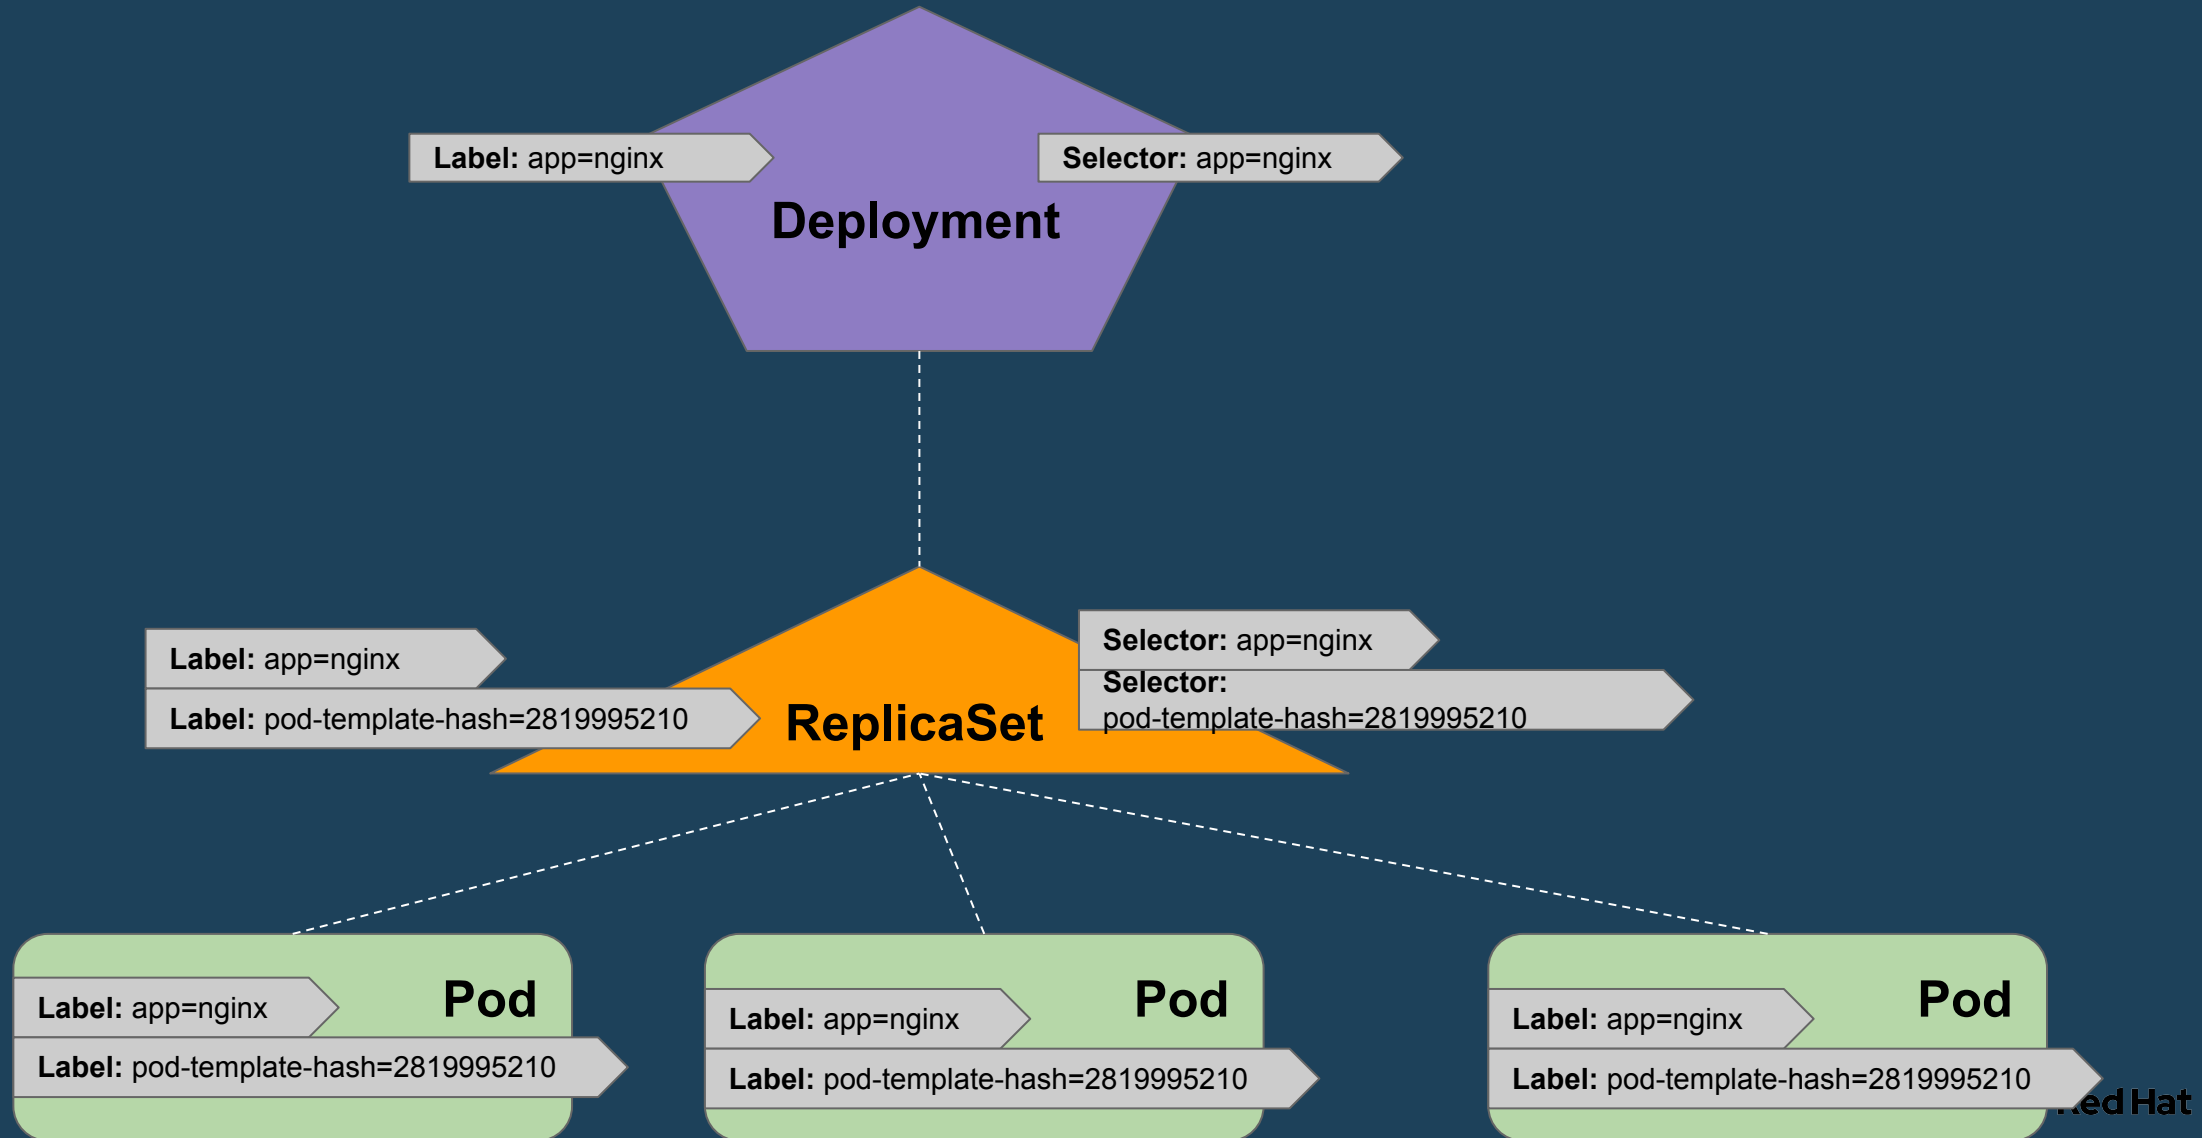
r.Client.Create Pod

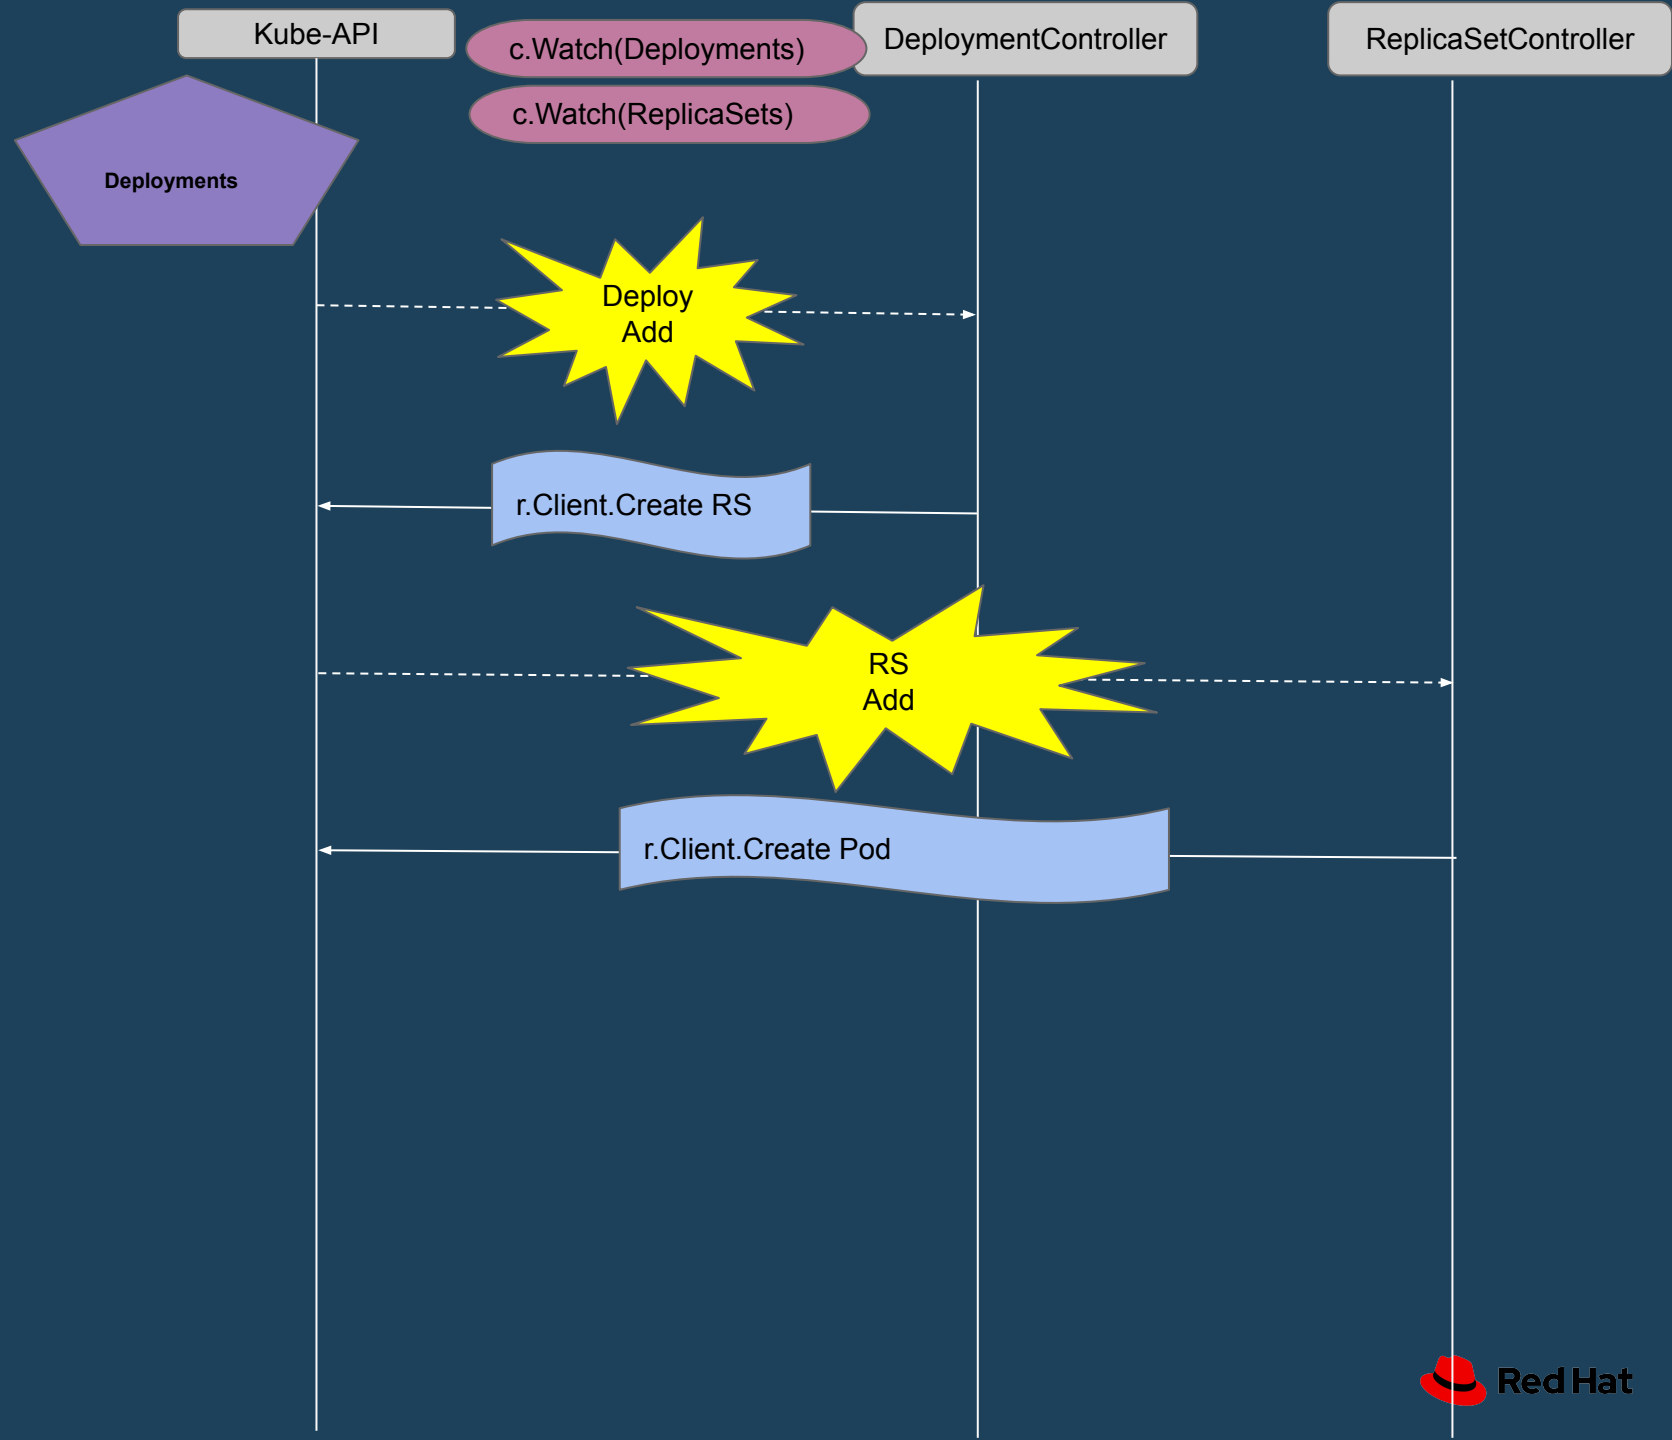
Pod 4
Add Event

r.Client.List Pods by label: rs.metadata.label



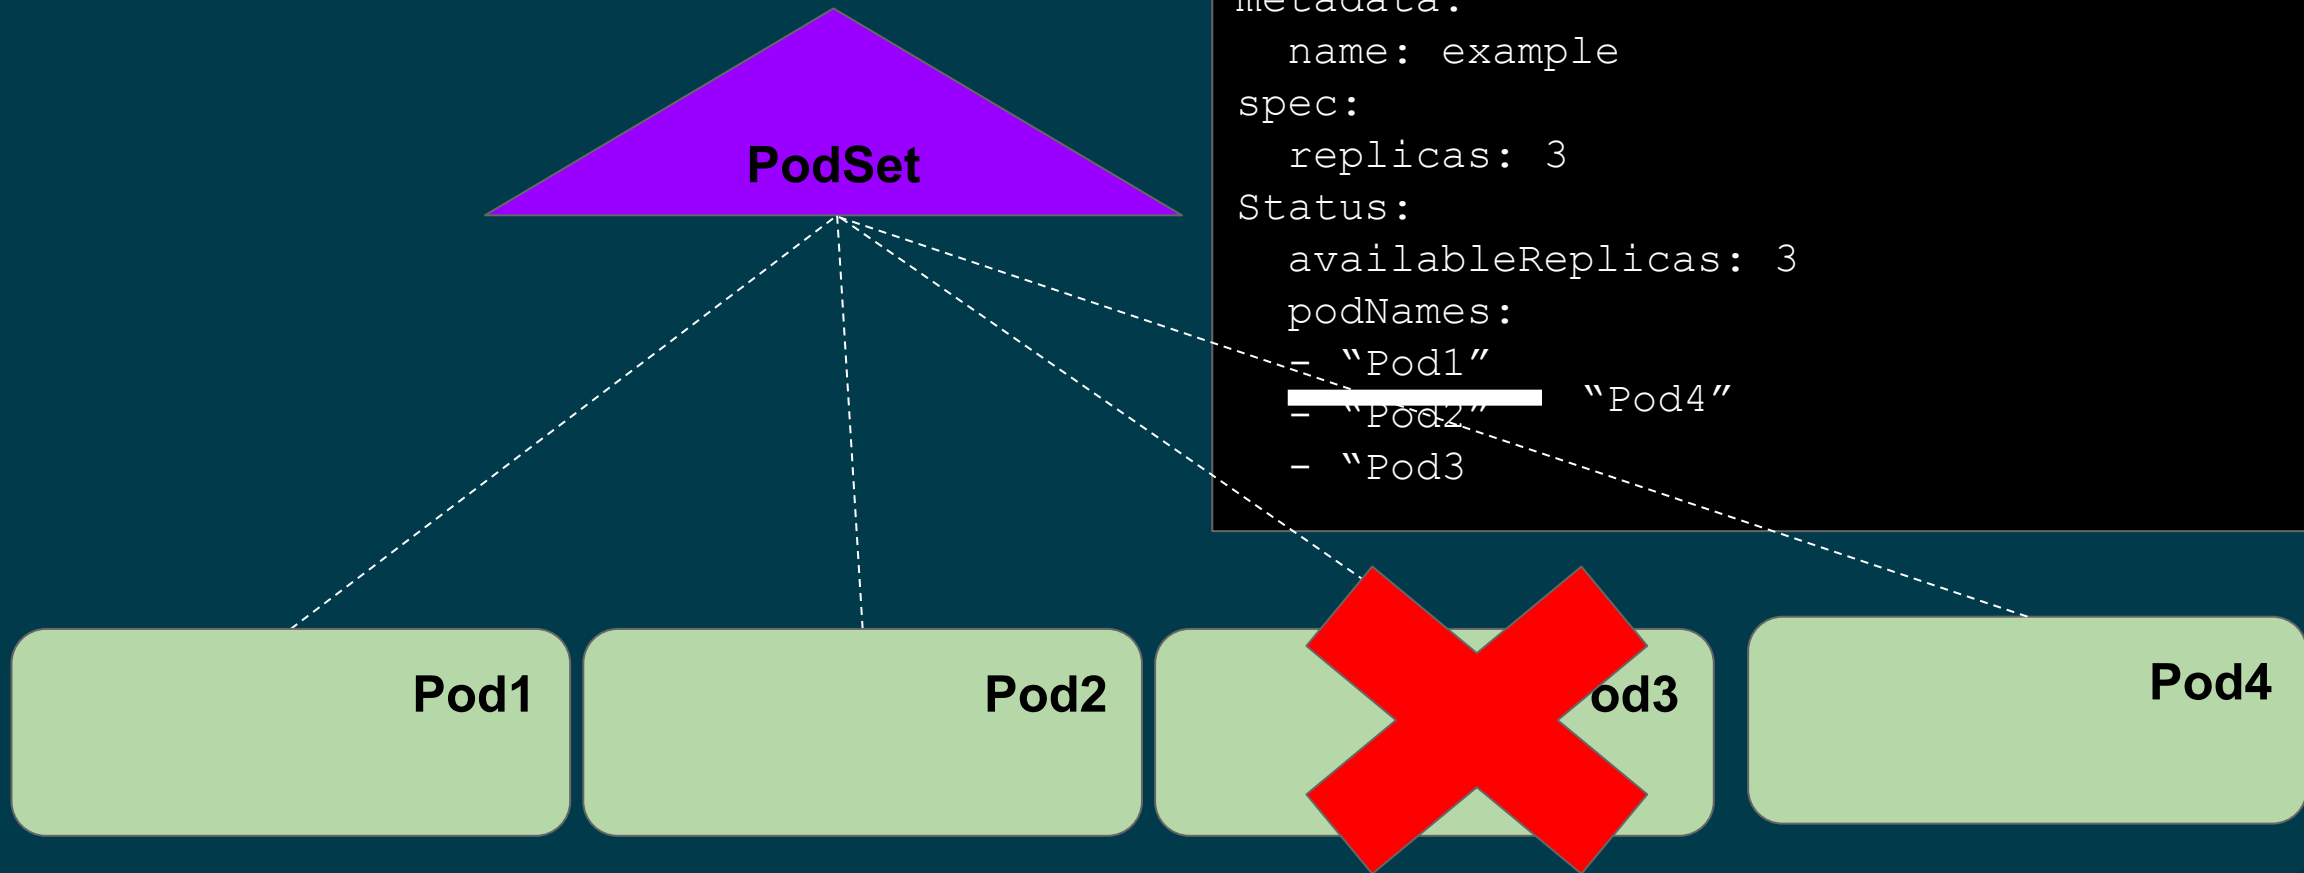
Deployments!





PodSet Operator

A Simple Controller that Manages Pods.



```
apiVersion: podset.redhat.com/v1alpha1
kind: PodSet
metadata:
  name: example
spec:
  replicas: 3
Status:
  availableReplicas: 3
  podNames:
    - "Pod1"
    - "Pod2"
    - "Pod3"
    - "Pod4"
```

A Pod Set Allows You to Scale Up/Down.

