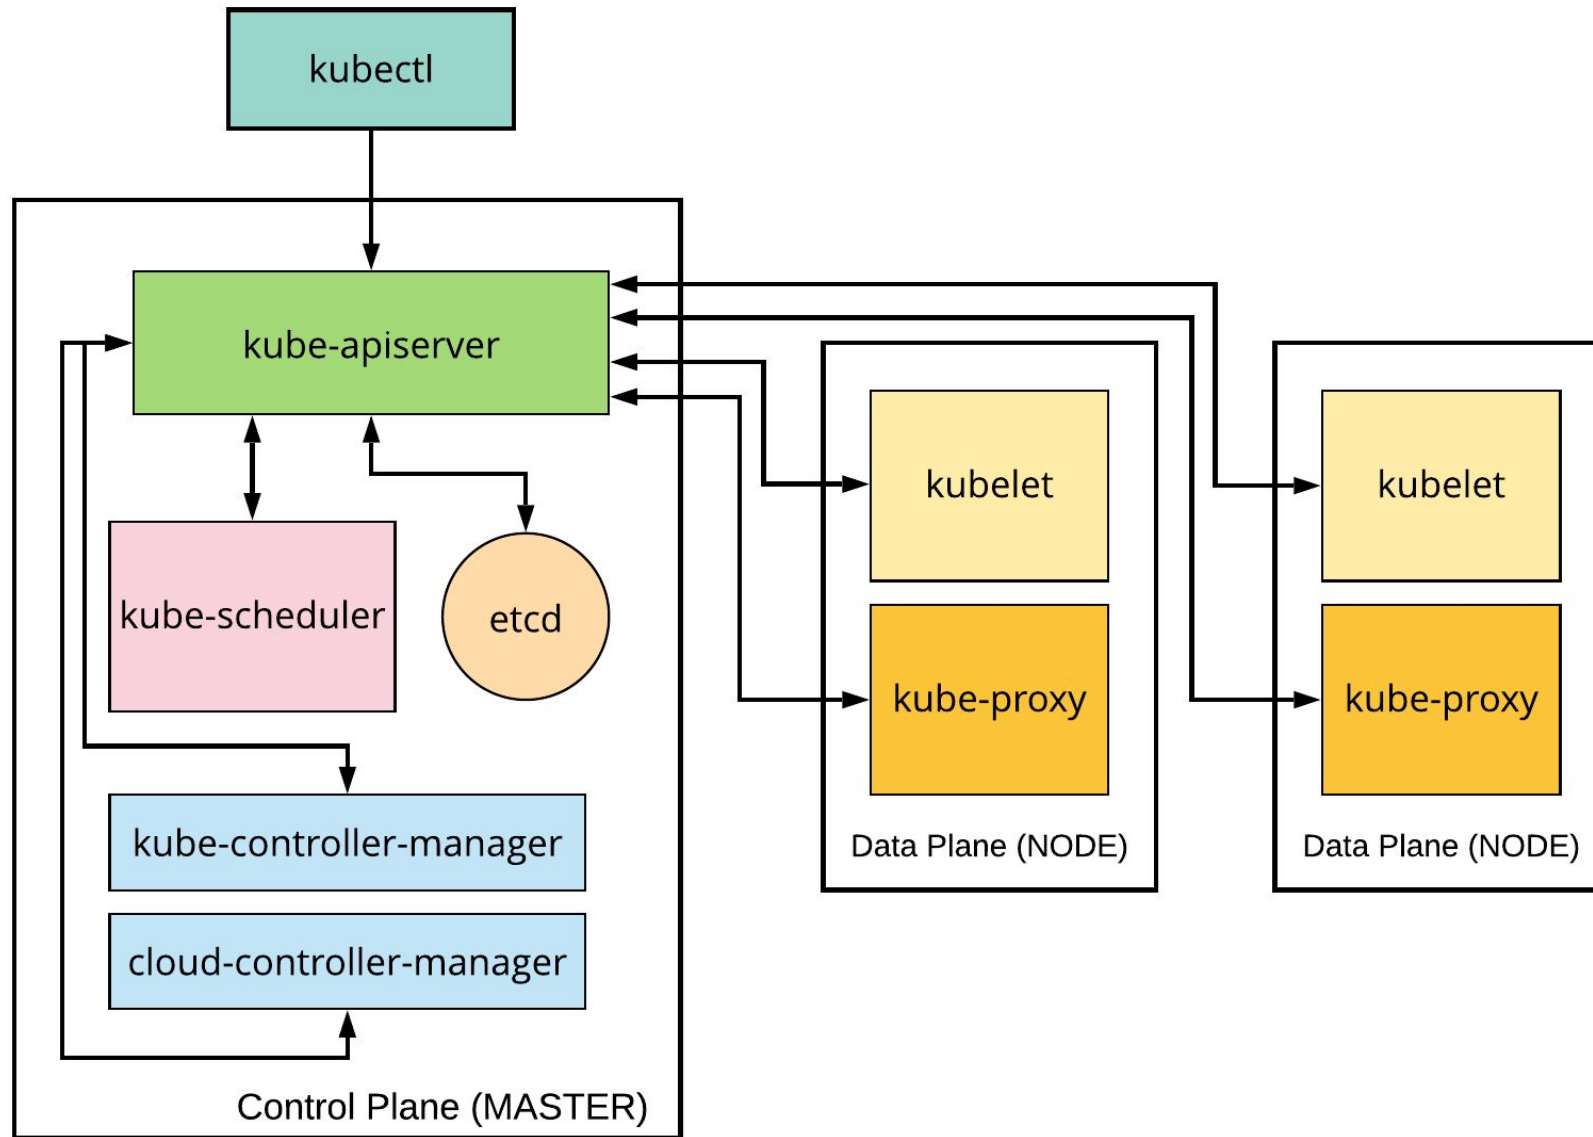


# What we'll discuss today

## Kubernetes API


- API Structure
- Optimistic Concurrency
- Versioning: Alpha, Beta, and Stable
- GroupVersionKind and GroupVersionResource
- Metadata, Spec, and Status
- Subcommands, API Actions, and HTTP Methods
- Deletions and Garbage Collection
- API Command Line Interaction

An Operator takes  
advantage of what  
Kubernetes does best



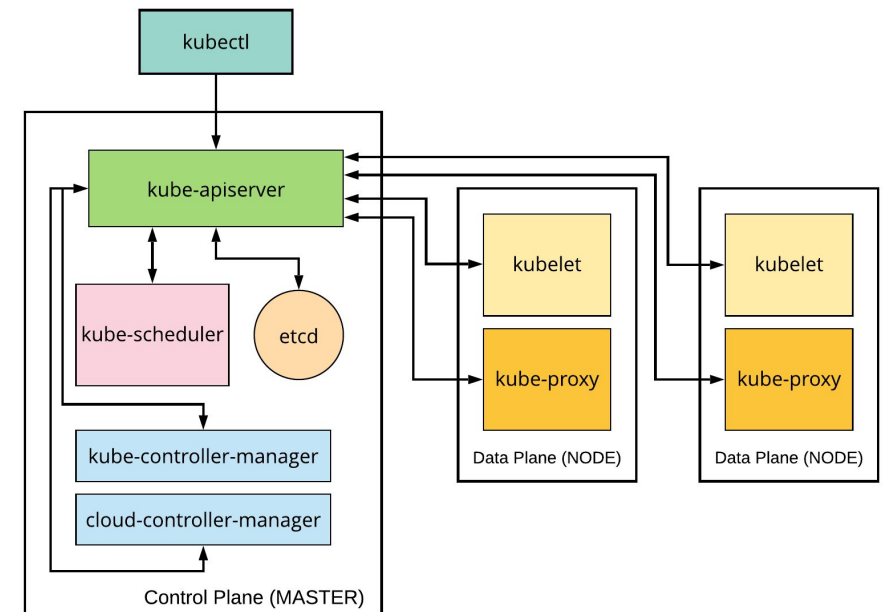
# Control Plane (Master)

- ▶ kube-apiserver
  - the only component that all other master and worker components directly communicate with.
  - validates and configures data for the api objects which include pods, services, deployments, and others.
- ▶ kube-scheduler
  - responsible for managing the scheduling of pods.
- ▶ kube-controller-manager
  - embeds the core control loops shipped with Kubernetes.
  - performs cluster-level functions like keeping track of workers and handling node failures.
- ▶ cloud-controller-manager
  - embeds the cloud specific control loops shipped with Kubernetes.
- ▶ etcd
  - distributed data store that persistently stores the cluster configuration.



# Data Plane (Worker)

- ▶ kube-proxy
  - network proxy that runs on each node.
  - can do simple TCP, UDP, and SCTP stream forwarding or round robin TCP, UDP, and SCTP forwarding across a set of backends.
  - reflects services as defined in the Kubernetes API on each node.
- ▶ kubelet
  - primary agent running on the node and registers the node with the API server.
  - ensures containers described in PodSpecs are running.
  - able to receive container manifests via File, HTTP endpoint, HTTP server
- ▶ container runtime
  - cri-o
  - rkt
  - containerd



## Command Line Interaction

```
oc proxy --port=8080
```

```
curl http://127.0.0.1:8080/apis/batch/v1
```

```
oc get --raw /apis/batch/v1
```


```
oc api-resources
```

```
oc api-versions
```


\* kubectl and oc are essentially the same (oc utilizes the packages of kubectl)

# Concurrency


## Monolithic



## Two-level



## Shared state




### Pessimistic Concurrency

- conflicts will happen often
- uses locks

### Optimistic Concurrency

- we believe conflicts will not happen
- if they happen react in some way

# API Versioning





# alpha



The version names contain alpha (e.g. v1alpha1).

May be buggy. Enabling the feature may expose bugs.  
Disabled by default.

Support for feature may be dropped at any time **without notice**.

The API may change in incompatible ways in a later software release **without notice**.

Recommended for use only in short-lived testing clusters, due to the increased risk of bugs and lack of long-term support.

# beta



The version names contain beta (e.g. v2beta3).

Code is well tested. Enabling the feature is considered safe.

**Enabled by default.**

Support for the overall feature will not be dropped, though details may change.

The schema and/or semantics of objects may change in incompatible ways in a subsequent beta or stable release.

Recommended for only **non-business-critical uses**.

# stable



The version name is vX where X is an integer.

Stable versions of features will appear in released software for many subsequent versions.

# GroupVersionKind or GVK

apiVersion: batch/v1

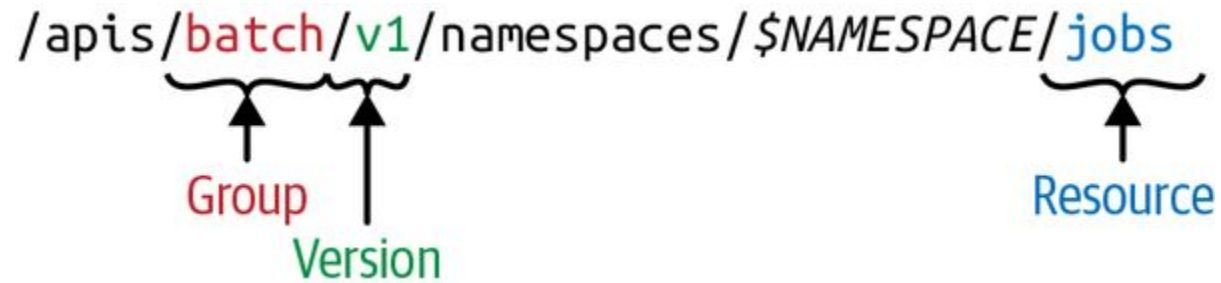
kind: Job

The entity **Group** is similar to package in a language. It disambiguates different APIs that may happen to have identically named Kinds. Groups often contain a domain name, such as redhat.com.

The entity **Version** defines the stability of the API and backward compatibility guarantees - such as v1beta1 or v1.

The entity **Kind** is the name of the API - such as Deployment or Service.

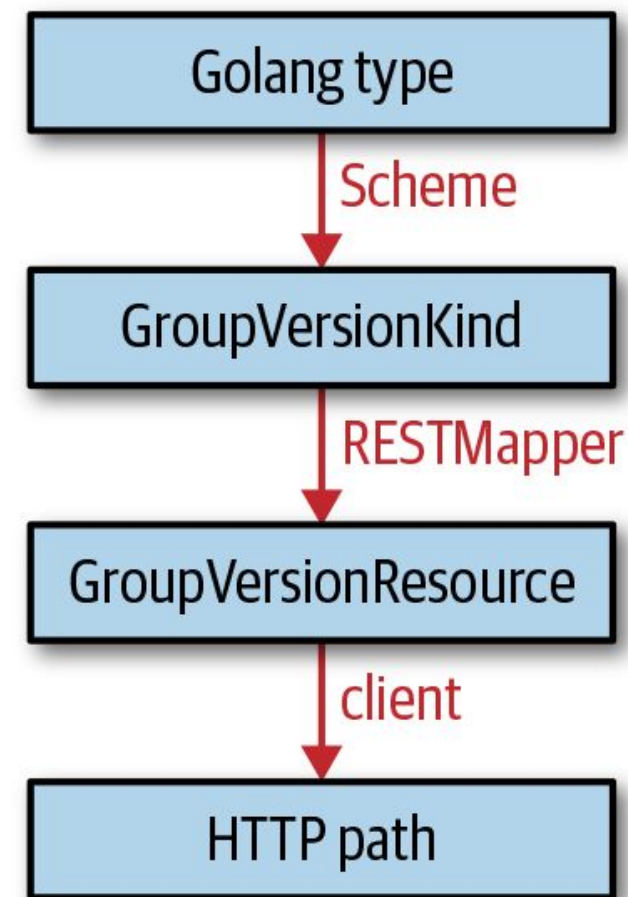
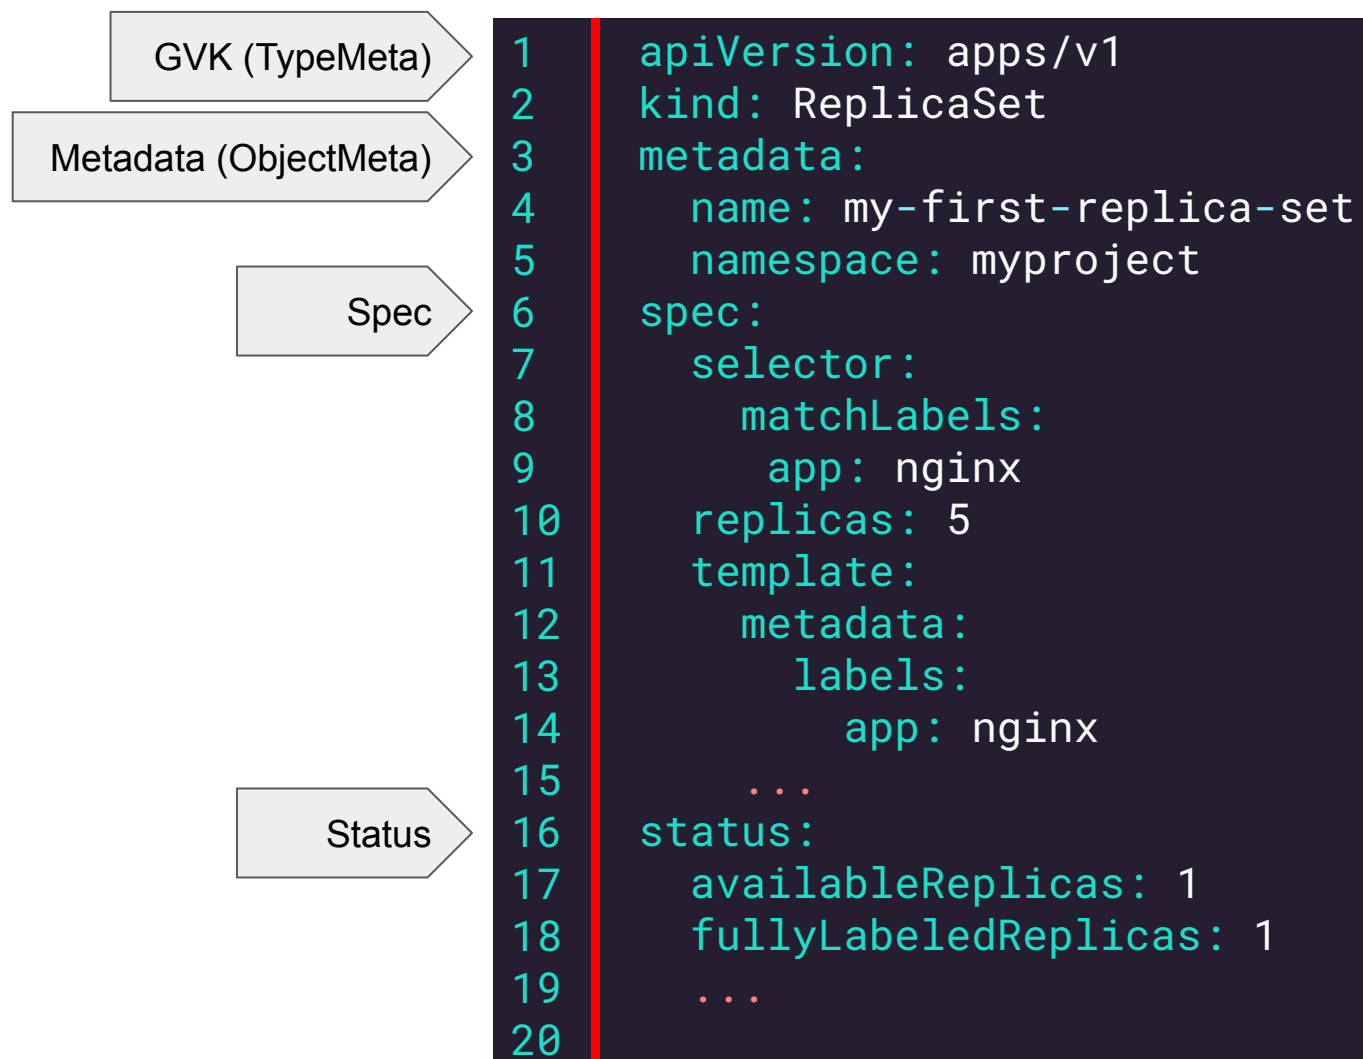
# GroupVersionResource or GVR



The resource **Group** is similar to package in a language. It disambiguates different APIs that may happen to have identically named Kinds.

The resource **Version** defines the stability of the API and backward compatibility guarantees – such as `v1beta1` or `v1`.

The **Resource** is the name of the path





# Kubernetes API Actions and HTTP Methods

Verb (API)

HTTP Method

get	GET
list	GET
watch	GET
create	POST
update	PUT
patch	PATCH
delete	DELETE
deletecollection	DELETE

## Kubernetes Subcommand and HTTP Methods

Subcommand	Object Does Not Exist	Object Exists
------------	-----------------------	---------------

apply	POST	PATCH   DELETE
create	POST	error!
replace	error!	PUT
delete	error!	DELETE
patch	error!	PATCH

**Garbage Collection**  
assists in deleting objects  
that have an **owner** that  
no longer exists.

# ownerReferences

Parent/child or owner/dependents Relationship

GroupVersion of  
owner object  
(required)

Kind of owner  
object (required)

Name of owner  
object (required)

uid of owner  
object (required)

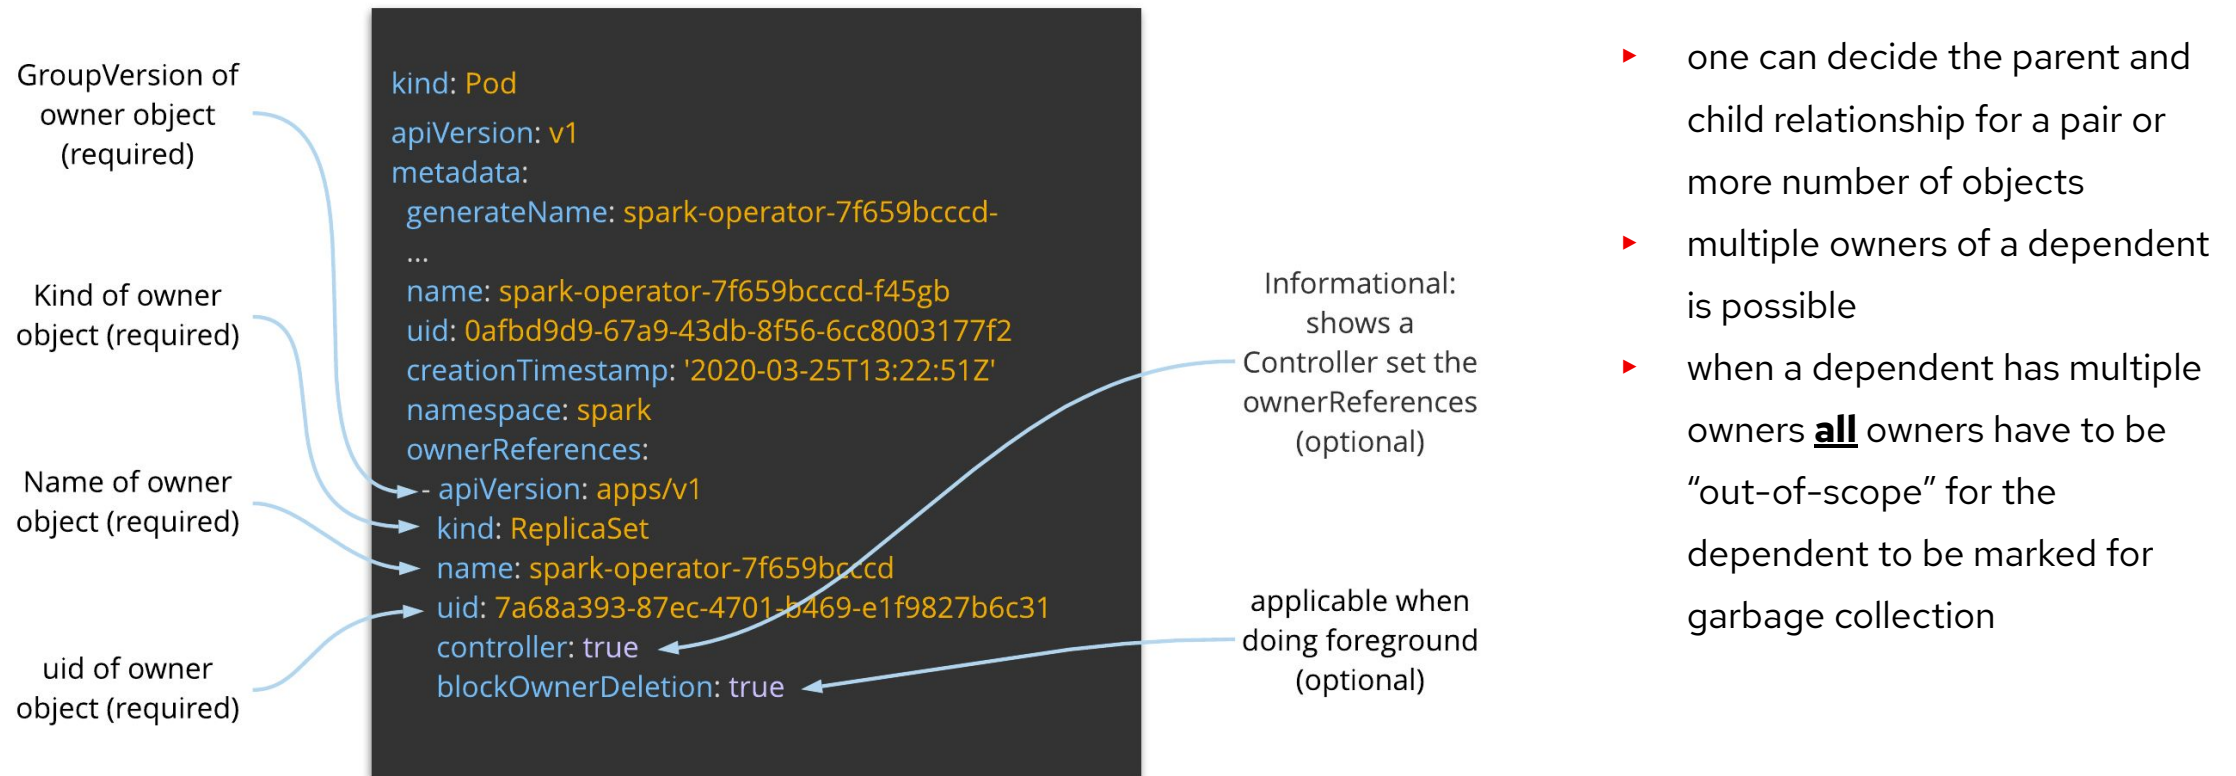
```
kind: Pod
apiVersion: v1
metadata:
  generateName: spark-operator-7f659bcccd-
  ...
  name: spark-operator-7f659bcccd-f45gb
  uid: 0afb9d9d-67a9-43db-8f56-6cc8003177f2
  creationTimestamp: '2020-03-25T13:22:51Z'
  namespace: spark
  ownerReferences:
  - apiVersion: apps/v1
    kind: ReplicaSet
    name: spark-operator-7f659bcccd
    uid: 7a68a393-87ec-4701-b469-e1f9827b6c31
    controller: true
    blockOwnerDeletion: true
```

Informational:  
shows a  
Controller set the  
ownerReferences  
(optional)

applicable when  
doing foreground  
(optional)

# ownerReferences

## Parent/child or owner/dependents Relationship



# Background, Foreground, and Finalizers

## BACKGROUND

- ▶ Kubernetes deletes the **owner** object immediately and the garbage collector then deletes the dependents in the background

## FOREGROUND

- ▶ object is still visible via the REST API
- ▶ the root object first enters a “deletion in progress” state
- ▶ once the garbage collector has deleted all “blocking” dependents it deletes the **owner** object

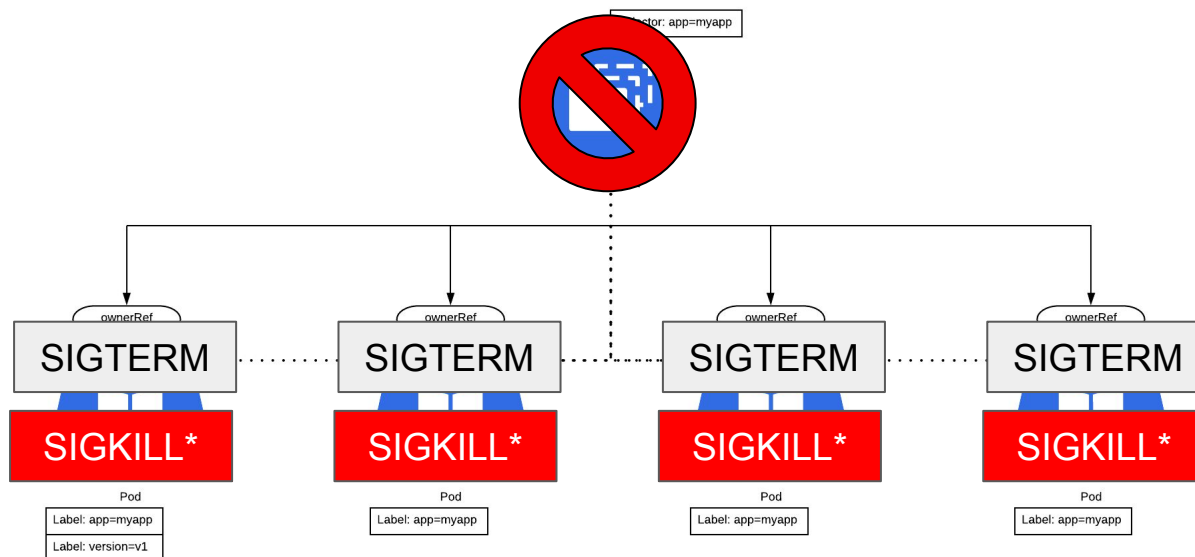
## FINALIZERS

- ▶ allow controllers to implement asynchronous pre-delete hooks
- ▶ arbitrary string values, that when present ensure that a hard delete of a resource is not possible while they exist
- ▶ Kubernetes only finally deletes the object if the list of finalizers is empty, meaning all finalizers have been executed

# Background Deletion

```
oc delete -f spark-rs.yaml
```

```
curl -x DELETE
http://localhost:8080/apis/apps/v1/namespaces/spark/replicasets/spark-operator-7f659bcccd \
-d '{"kind":"DeleteOptions","apiVersion":"v1","propagationPolicy":"Background"}' \
-H "Content-Type: application/json"
```



ownerReferences:

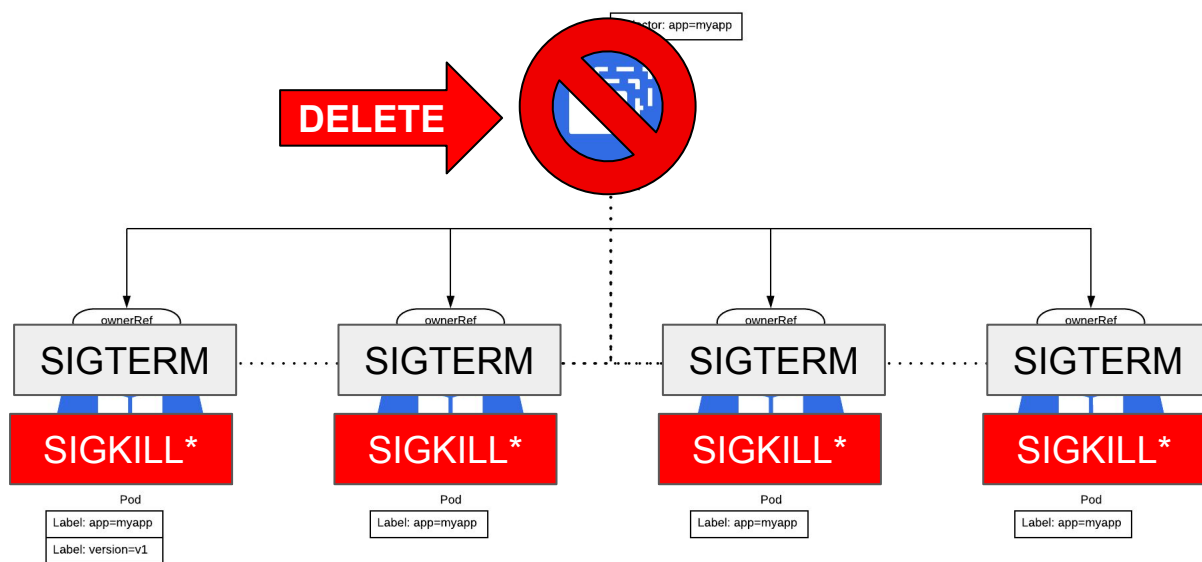
- apiVersion: apps/v1
- kind: ReplicaSet
- name: spark-operator-7f659bcccd
- uid: 7a68a393-87ec-4701-b469-e1f9827b6c31
- controller: true
- blockOwnerDeletion: true

\*wait 30s after SIGTERM then SIGKILL

# Foreground Deletion

```
oc delete -f spark-rs.yaml
```

```
curl -x DELETE
http://localhost:8080/apis/apps/v1/namespaces/spark/replicasets/spark-operator-7f659bcccd \
-d '{"kind":"DeleteOptions","apiVersion":"v1","propagationPolicy":"Foreground"}' \
-H "Content-Type: application/json"
```



ownerReferences:

- apiVersion: apps/v1
- kind: ReplicaSet
- name: spark-operator-7f659bcccd
- uid: 7a68a393-87ec-4701-b469-e1f9827b6c31
- controller: true
- blockOwnerDeletion: true

\*wait 30s after SIGTERM then SIGKILL



# Foreground Deletion

metadata:

deletionTimestamp: 2019-10-20T01:16:04Z

Finalizers: "foregroundDeletion"

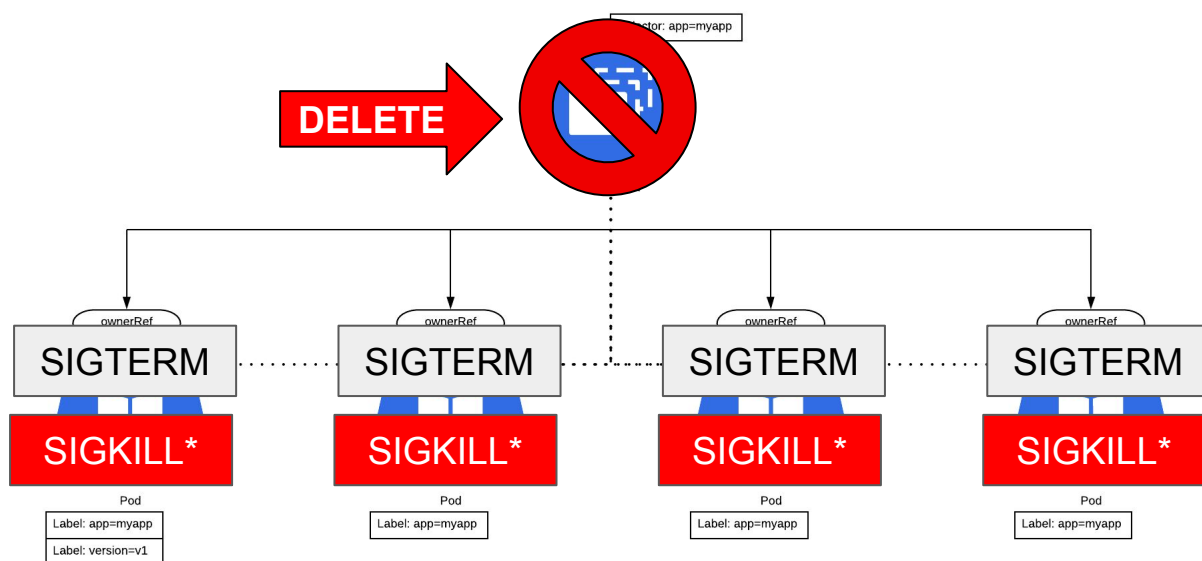
`oc delete -f spark-rs.yaml`

`curl -x DELETE`

`http://localhost:8080/apis/apps/v1/namespaces/spark/replicasets/spark-operator-7f659bcccd \`  
`-d '{"kind":"DeleteOptions","apiVersion":"v1","propagationPolicy":"Foreground"}' \`  
`-H "Content-Type: application/json"`

ownerReferences:

- apiVersion: apps/v1
- kind: ReplicaSet
- name: spark-operator-7f659bcccd
- uid: 7a68a393-87ec-4701-b469-e1f9827b6c31
- controller: true
- blockOwnerDeletion: true



\*wait 30s after SIGTERM then SIGKILL

# Foreground Deletion

metadata:

deletionTimestamp: 2019-10-20T01:16:04Z

Finalizers: "foregroundDeletion"

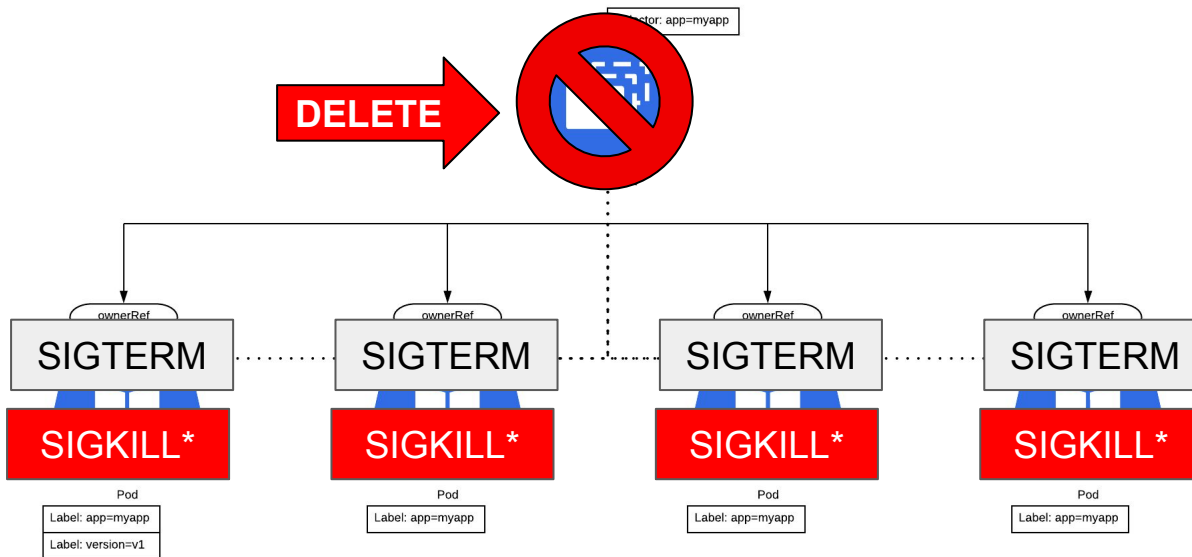
`oc delete -f spark-rs.yaml`

`curl -x DELETE`

`http://localhost:8080/apis/apps/v1/namespaces/spark/replicasets/spark-operator-7f659bcccd \`  
`-d '{"kind":"DeleteOptions","apiVersion":"v1","propagationPolicy":"Foreground"}' \`  
`-H "Content-Type: application/json"`

ownerReferences:

- apiVersion: apps/v1
- kind: ReplicaSet
- name: spark-operator-7f659bcccd
- uid: 7a68a393-87ec-4701-b469-e1f9827b6c31
- controller: true
- blockOwnerDeletion: true



\*wait 30s after SIGTERM then SIGKILL

# Finalizers

metadata:

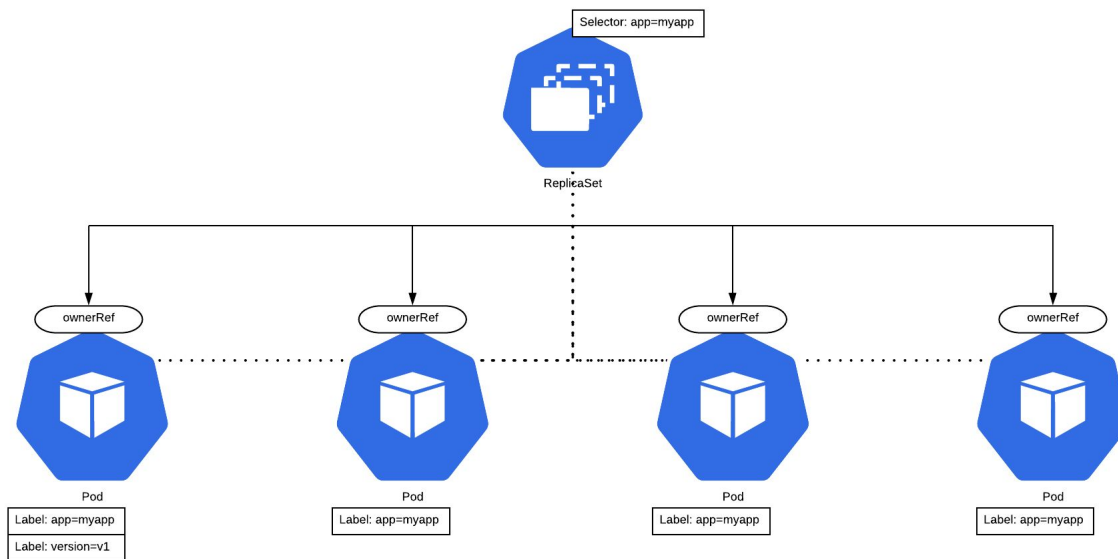
deletionTimestamp: 2019-10-20T01:16:04Z

Finalizers: "foregroundDeletion"

`oc delete -f spark-rs.yaml`

`curl -x DELETE`

`http://127.0.0.1:8080/apis/apps/v1/namespaces/spark/replicasets/spark-operator-7f659bcccd \`  
`-d '{"kind":"DeleteOptions","apiVersion":"v1","propagationPolicy":"Background"}' \`  
`-H "Content-Type: application/json"`



ownerReferences:

- apiVersion: apps/v1
- kind: ReplicaSet
- name: spark-operator-7f659bcccd
- uid: 7a68a393-87ec-4701-b469-e1f9827b6c31
- controller: true
- blockOwnerDeletion: true

metadata:

deletionTimestamp: 2019-10-20T01:16:04Z

Finalizers: "foregroundDeletion"